

Lecture 4:

3D and AI

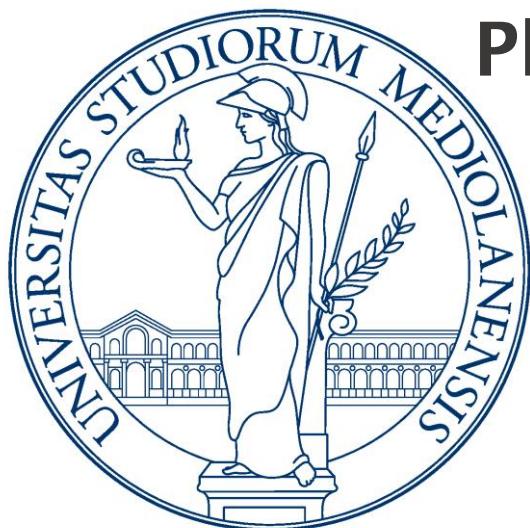
Simone Melzi, Marco Tarini

Milano, 28/07/2021

PhD School – Learning on 3D geometries

LA STATALE Università degli Studi di Milano

SAPIENZA Università di Roma



Why Geometric deep Learning?



Pixels (Euclidean)

Why Geometric deep Learning?

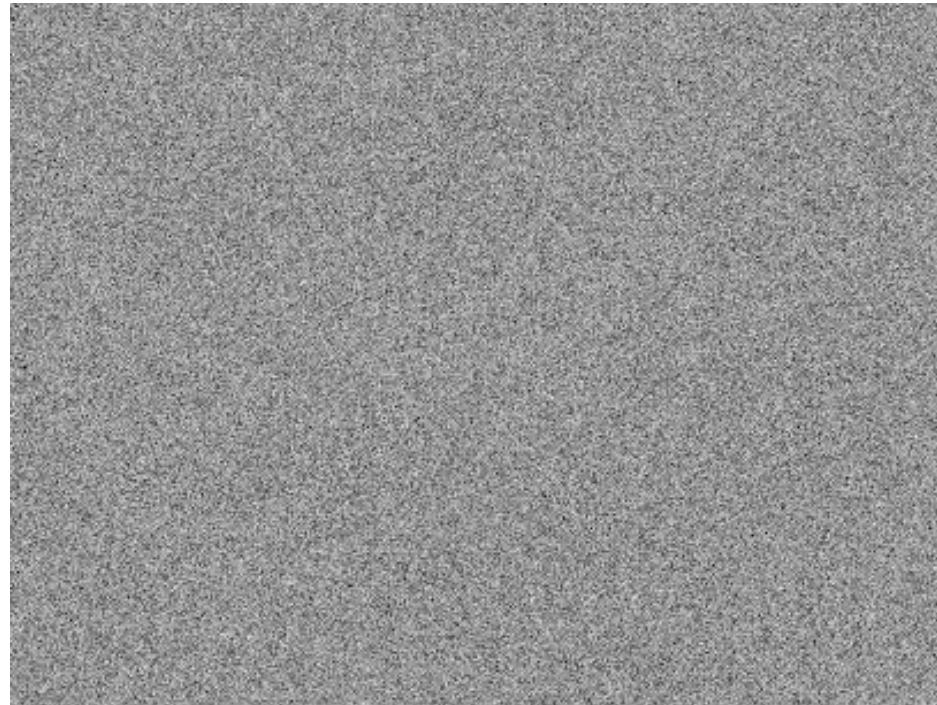


Geometry (Non-Euclidean)

Pixels (Euclidean)

Need for priors:

Data often owns structural priors in terms of repeating patterns, compositionality, ...



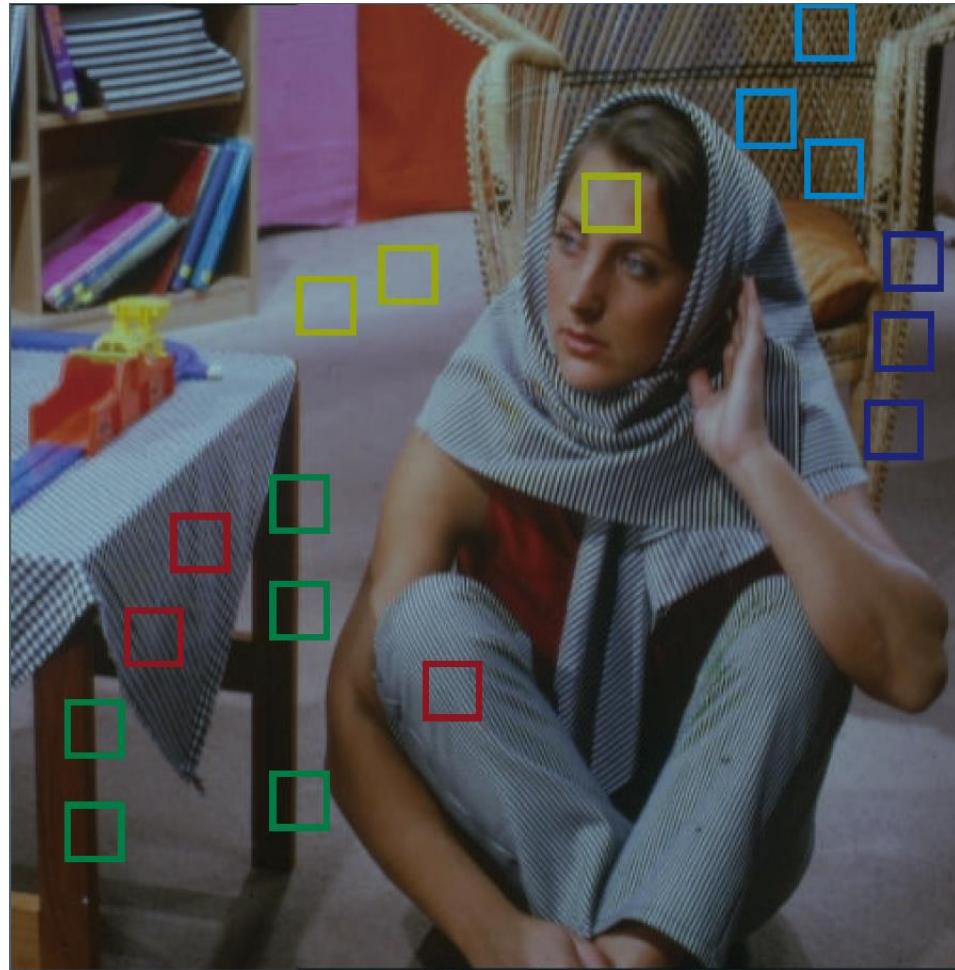
Need for priors:

Data often owns structural priors in terms of repeating patterns, compositionality, ...



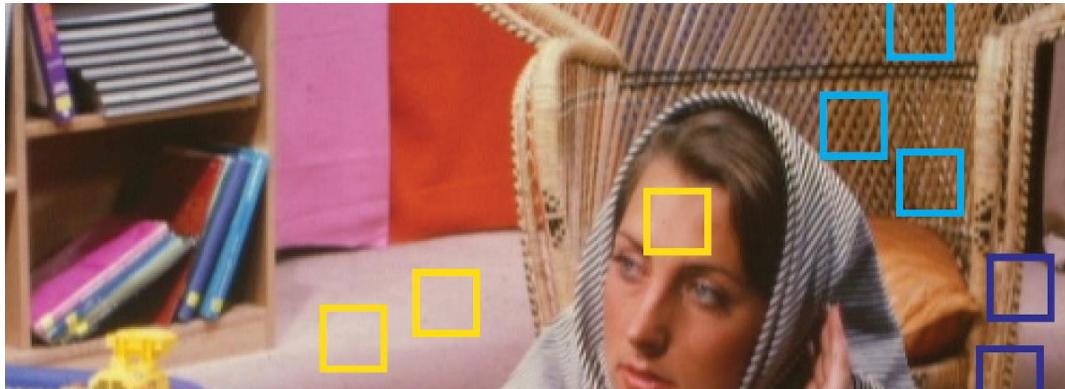
Self similarity:

Data tends to be **self-similar** across the domain



Hierarchy:

Translation invariance is desirable across multiple scales



We expect local features to be invariant to their location in the image:

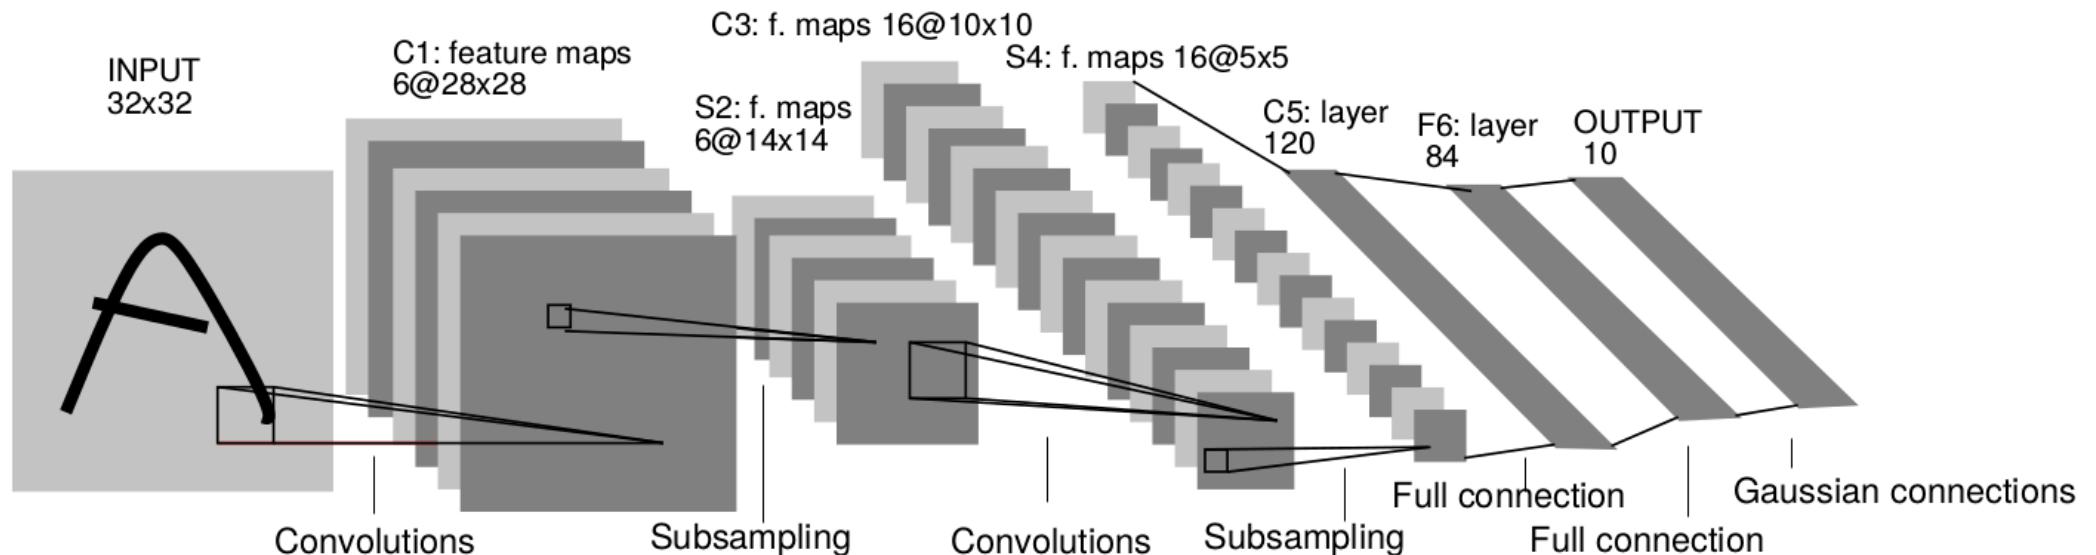
$$z(T_v p) = z(p), \quad \forall p, T_v$$

Where p are image patches of variable size.

CNN:

Data are often composed of hierarchical, local, shift-invariant patterns.

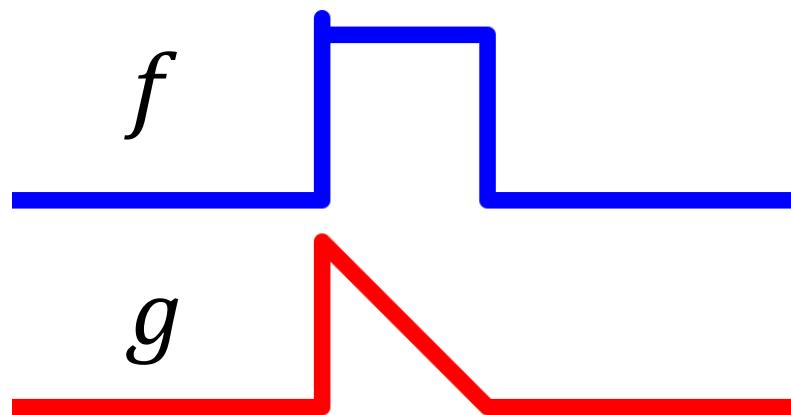
CNNs directly exploit this fact as a prior.



Convolution:

Given two functions $f, g: D = [0, 2\pi] \rightarrow \mathbb{R}$
then their **convolution** is the function:

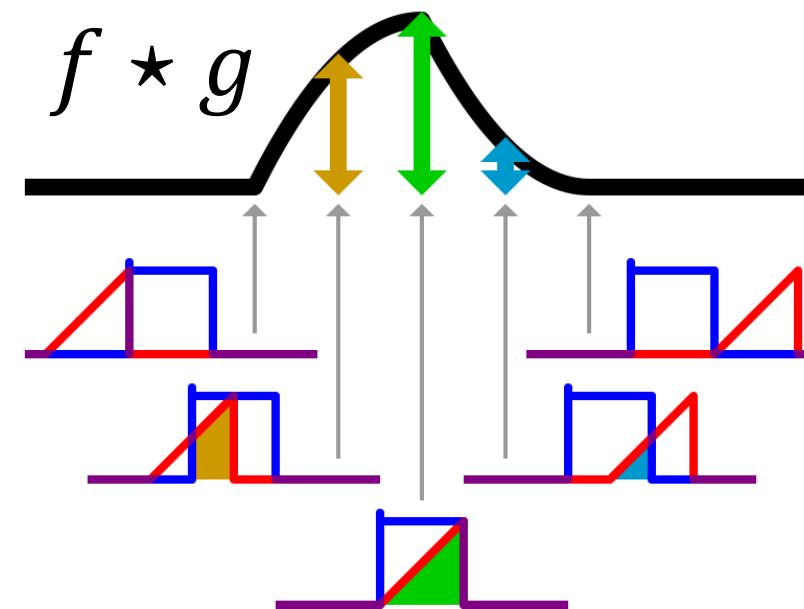
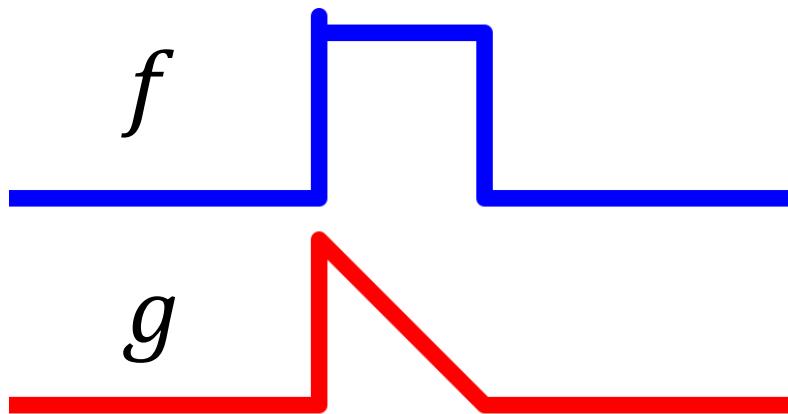
$$(f * g)(x) = \int_0^{2\pi} f(t)g(x - t) dt$$



Convolution:

Given two functions $f, g: D = [0, 2\pi] \rightarrow \mathbb{R}$
then their **convolution** is the function:

$$(f * g)(x) = \int_0^{2\pi} f(t)g(x - t) dt$$

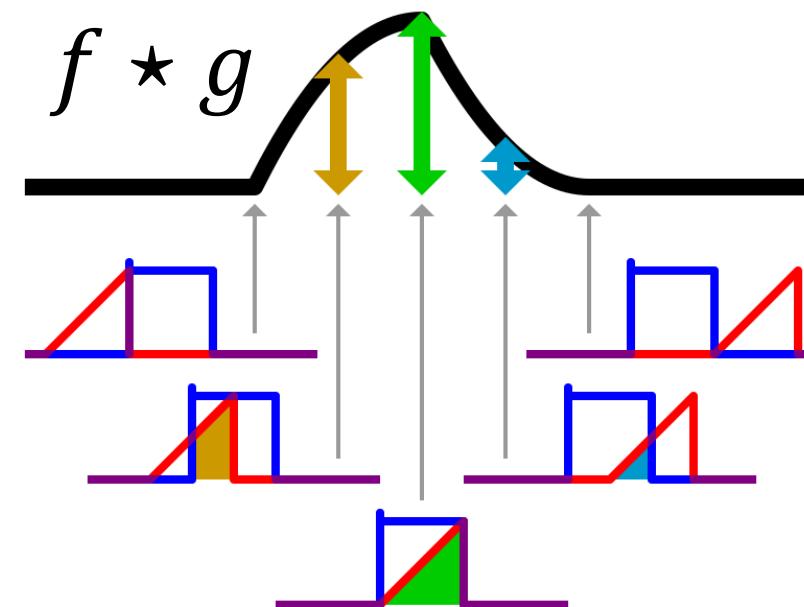
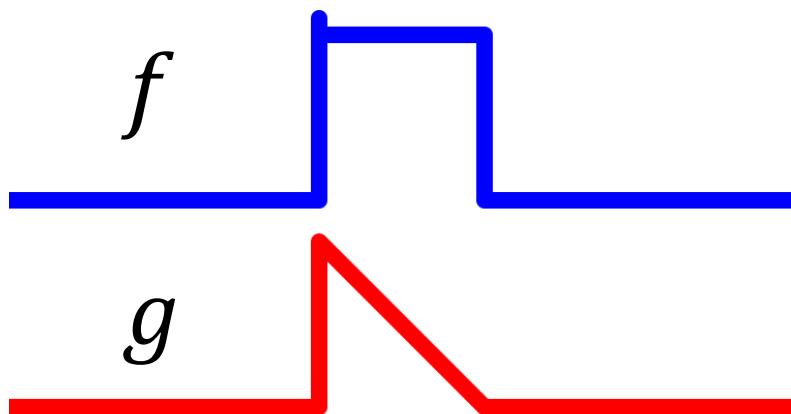


Convolution:

Given two functions $f, g: D = [0, 2\pi] \rightarrow \mathbb{R}$
then their **convolution** is the function:

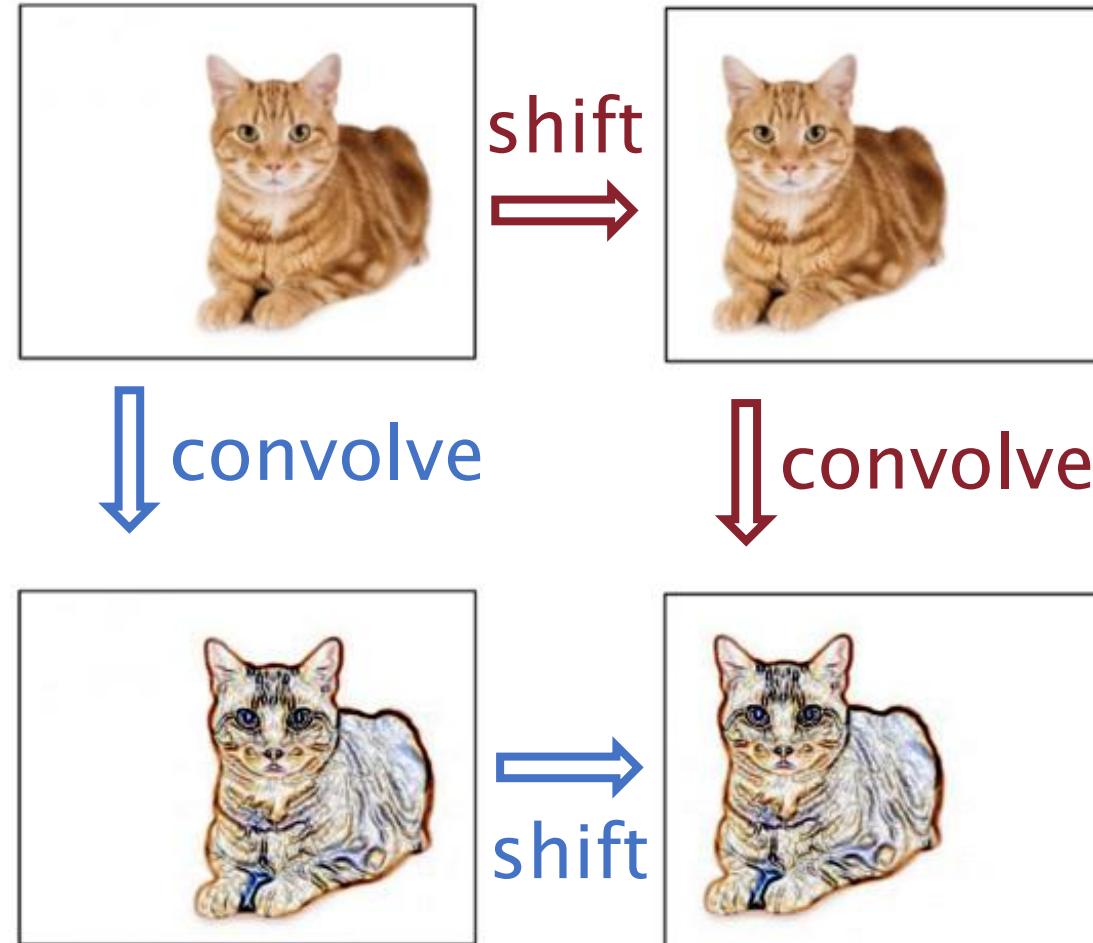
$$(f * g)(x) = \int_0^{2\pi} f(t)g(x - t) dt$$

Feature map Kernel



Convolution: Shift-equivariance

Convolution is shift-equivariant



Convolution: Linearity

Convolution could be seen as the application of a linear operator \mathcal{G}

$$\mathcal{G}f(x) = (f \star g)(x) = \int_0^{2\pi} f(t)g(x - t) dt$$

$$\mathcal{G} \text{ is linear: } \mathcal{G}(\alpha f(x)) = \alpha \mathcal{G}(f(x))$$

$$\mathcal{G}(f + h)(x) = \mathcal{G}(f(x)) + \mathcal{G}(h(x))$$

The shift-equivariance can be written as:

$$\mathcal{G}(\mathcal{T}f) = \mathcal{T}(\mathcal{G}f) \Leftrightarrow \mathcal{G} \text{ and } \mathcal{T} \text{ commute}$$

Discrete convolution:

In the discrete setting we deal with vectors f and g and we define the convolution as:

$$(f \star g)[x_i] = \sum_{j=1}^n f[x_j] g[x_i - x_j]$$

Assuming cyclic boundary conditions the convolution operator can be encoded as a matrix (Toeplitz matrix)

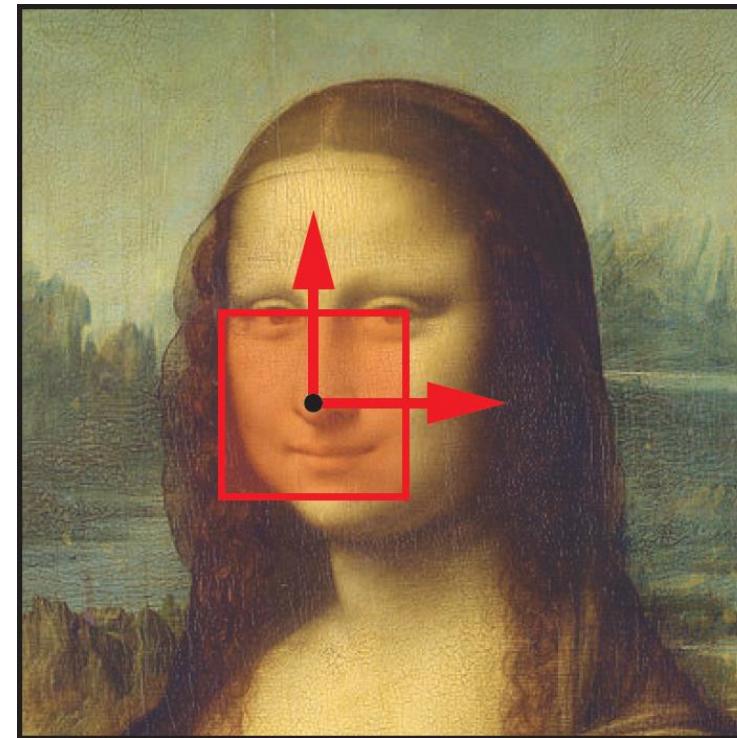
$$f \star g = \begin{bmatrix} g_1 & g_2 & \cdots & \cdots & g_n \\ g_n & g_1 & g_2 & \cdots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_2 & g_3 & \cdots & \cdots & g_1 \end{bmatrix} = \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Discrete convolution:

On 2D domains (RGB images $f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$) for each channel

$$(f \star g)[x_i, y_l] = \sum_{j=1}^{nr} \sum_{k=1}^{nc} f[x_j, y_k] g[x_i - x_j, y_l - y_k]$$

We get the classical interpretation of the moving window

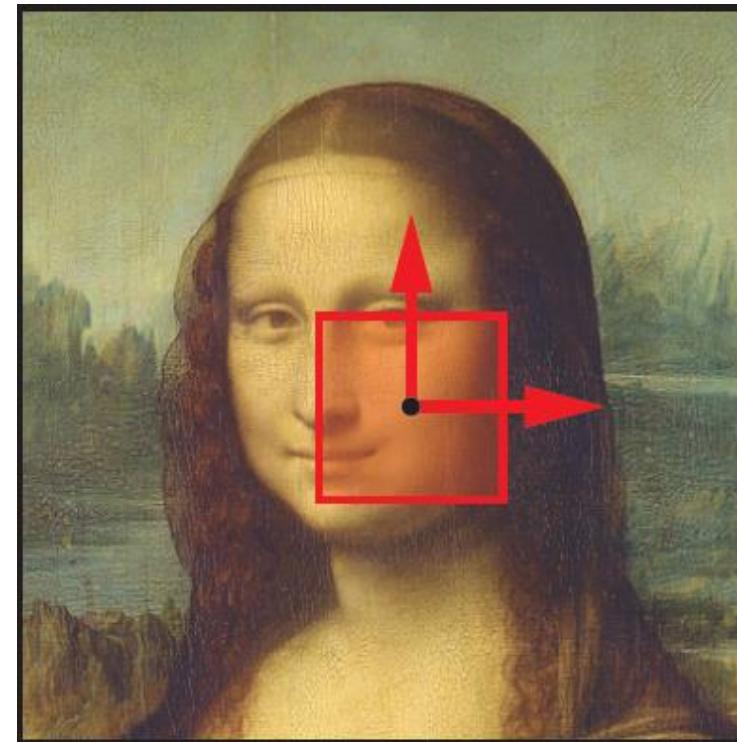


Discrete convolution:

On 2D domains (RGB images $f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$) for each channel

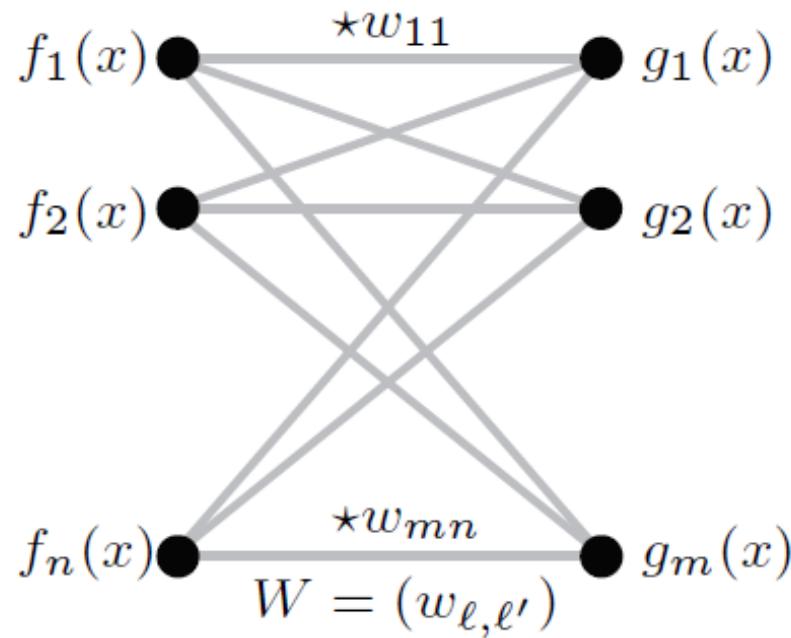
$$(f * g)[x_i, y_l] = \sum_{j=1}^{nr} \sum_{k=1}^{nc} f[x_j, y_k] g[x_i - x_j, y_l - y_k]$$

We get the classical interpretation of the moving window



Convolutional Neural Network

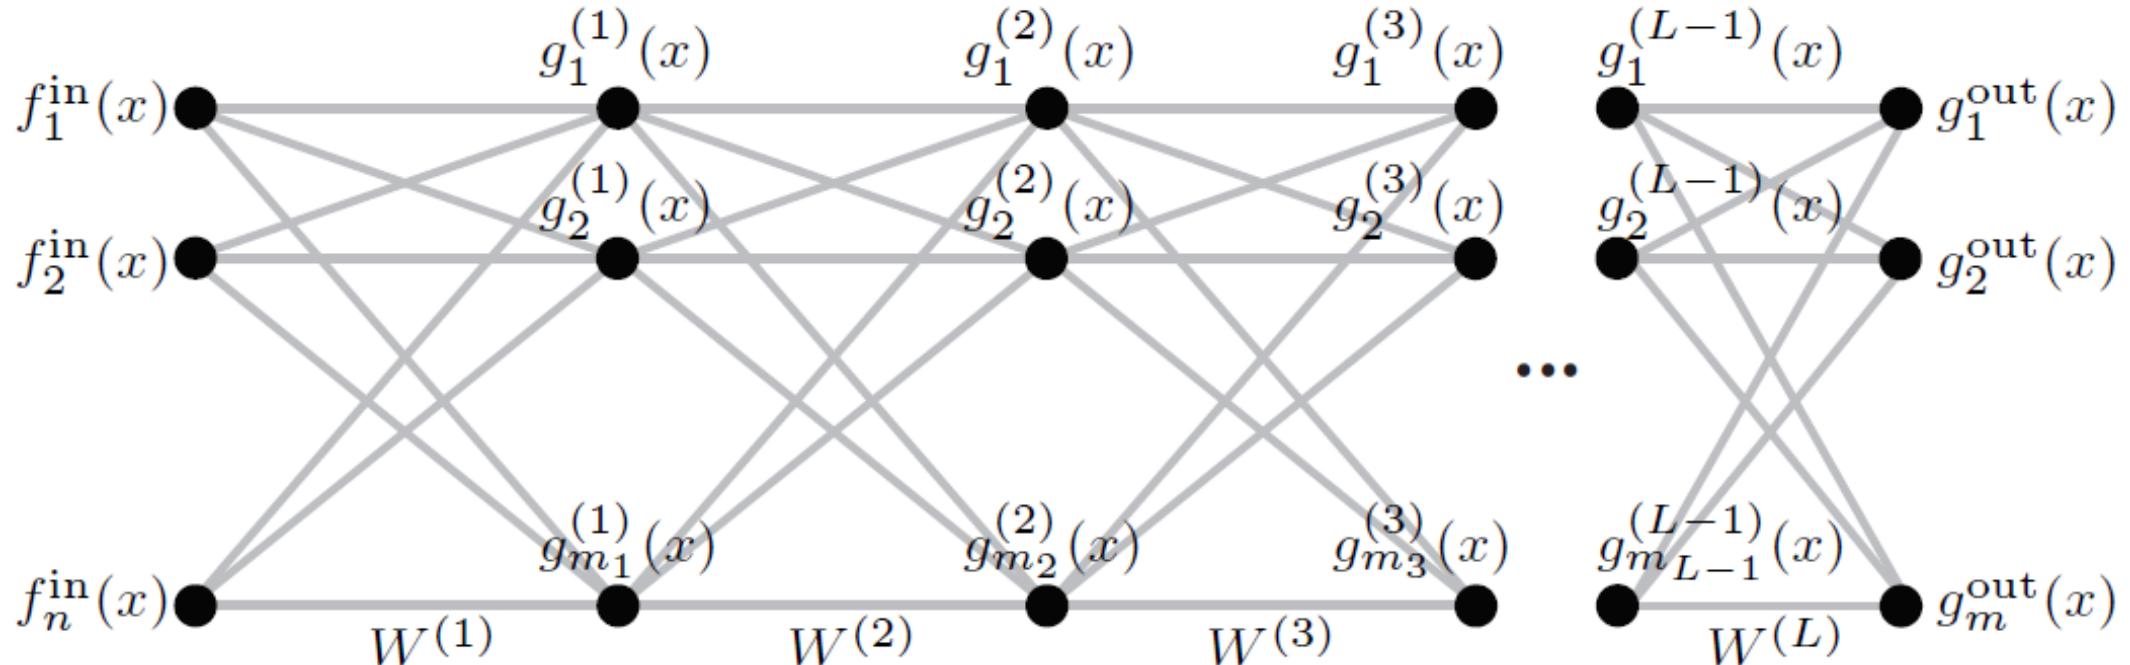
Main Idea: Compose equivariant layers (based on convolutions)



Single convolutional layer

Convolutional Neural Network

Main Idea: Compose equivariant layers (based on convolutions)



CNN consisting of L convolutional layers

Convolutional Neural Network

Main Idea: Compose equivariant layers (based on convolutions)

Conv. Layer:

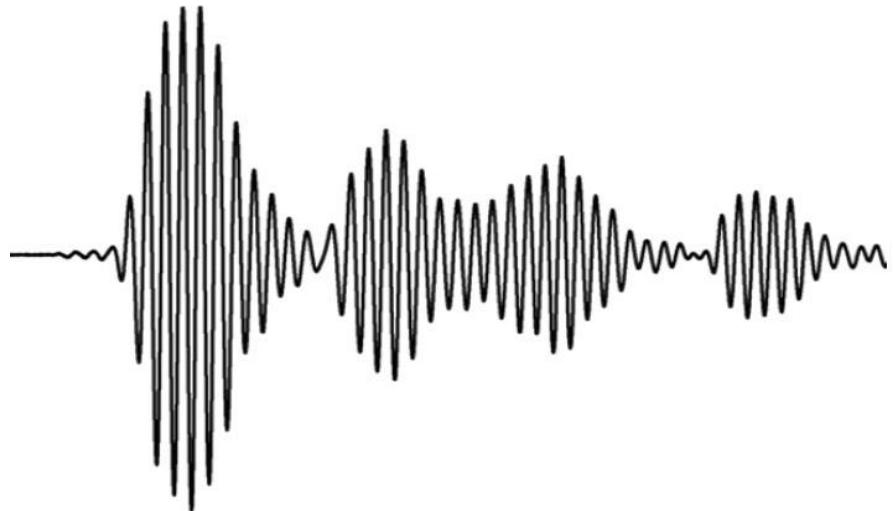
$$g_l^{(k)}(x) = \sigma \left(\sum_{l'=1}^{m_{k-1}} (g_{l'}^{(k-1)} \star w_{l,l'}^{(k)})(x) \right) \quad \begin{matrix} l = 1, \dots, m_k \\ l' = 1, \dots, m_{k-1} \end{matrix}$$

Activation: $\sigma(x) = \max(\{0, x\})$ rectified linear unit (ReLU)

Parameters: Filters of all layers $W^{(1)}, \dots, W^{(L)}$

Usual data for deep learning:

Audio signals and images



Different signals:



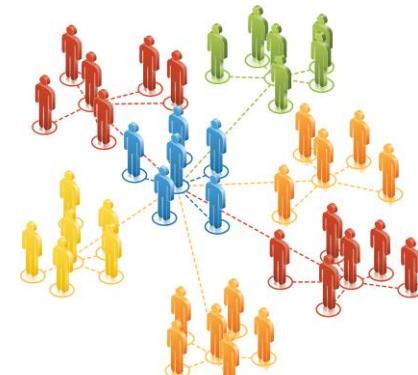
Audio



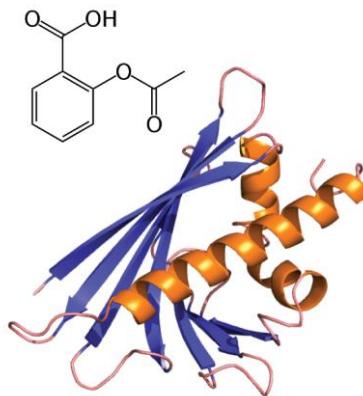
Image



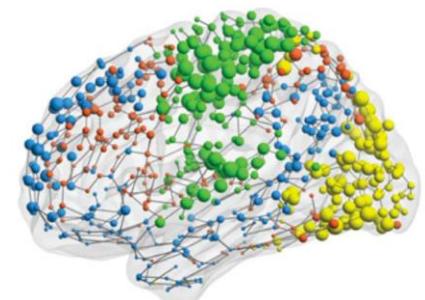
Road maps



Social
networks



Molecules



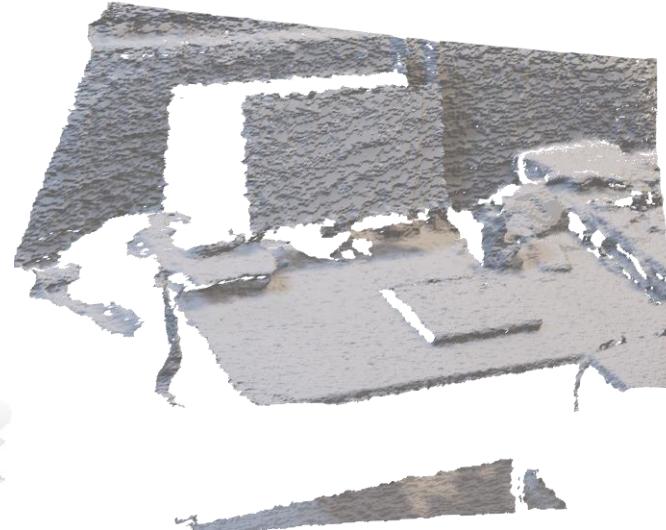
Functional
networks

Different domains:

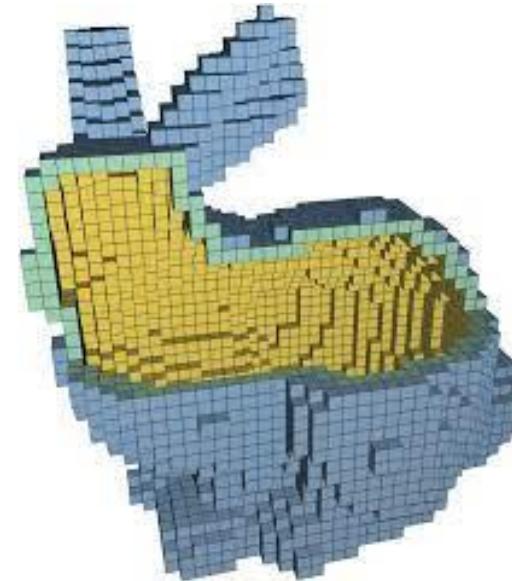
We want to consider non-Euclidean domains



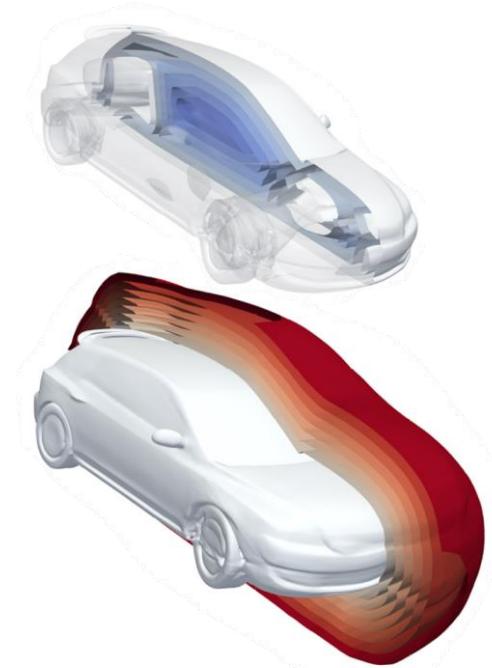
Point clouds



Range map



Volumetric

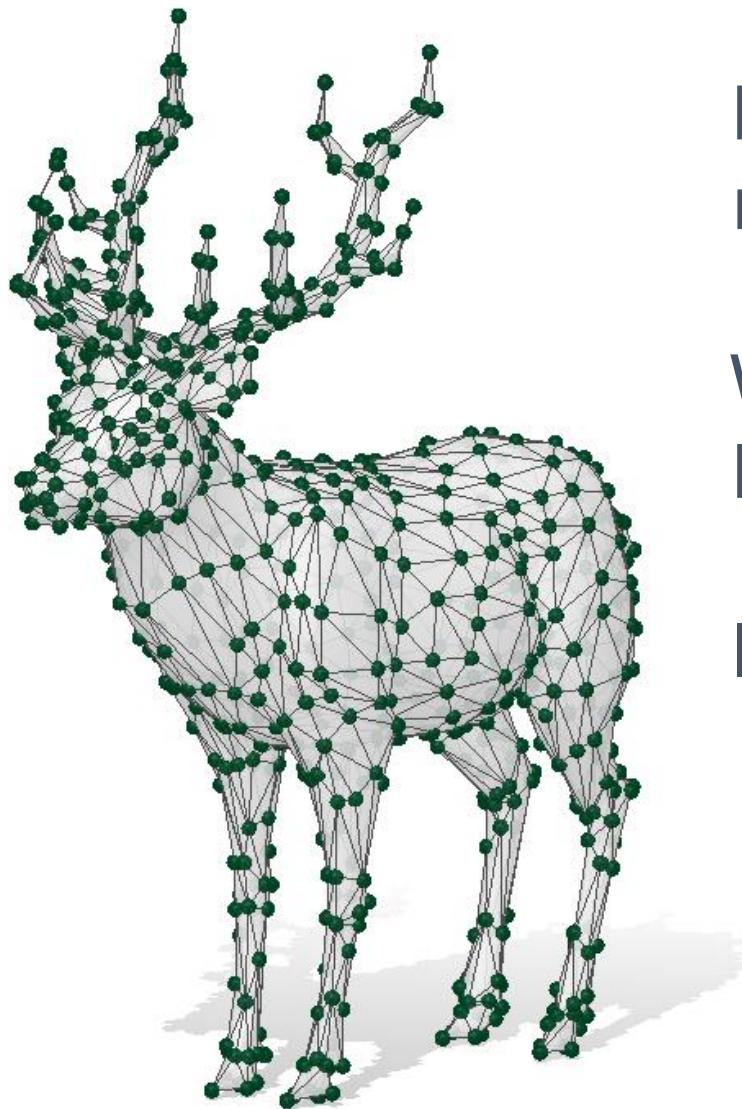


Implicit

["Fast Parallel Surface and Solid Voxelization on GPUs", M. Schwarz et al., 2010](#)

["Implicit Geometric Regularization for Learning Shapes", A. Gropp et al., 2020](#)

... and Meshes:



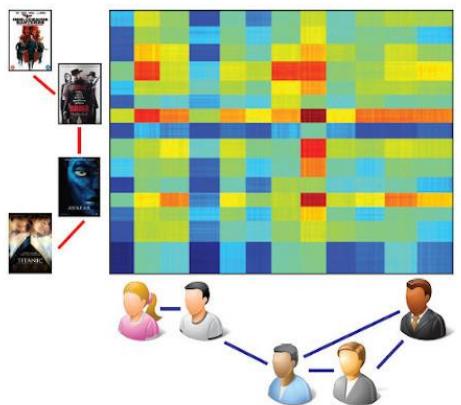
Piecewise linear approximation of manifolds.

We want to inject geometry in the Deep learning processing.

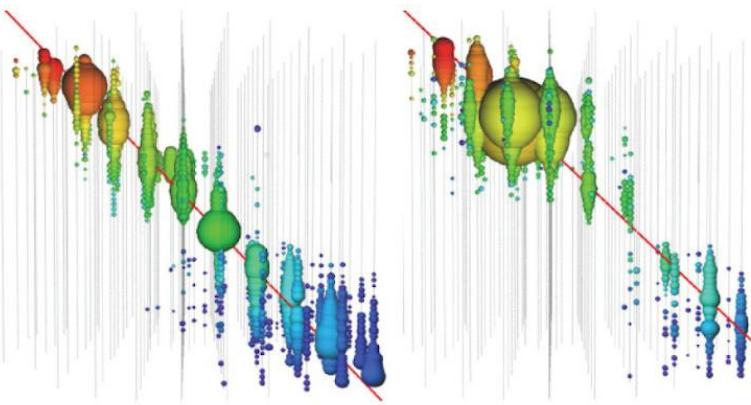
Is this new?

We want to do learning on the geometry of the domain!

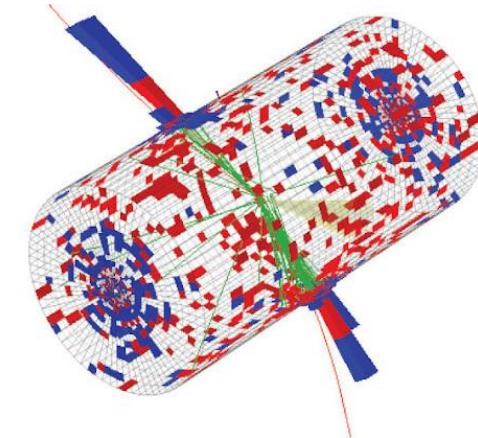
Applications:



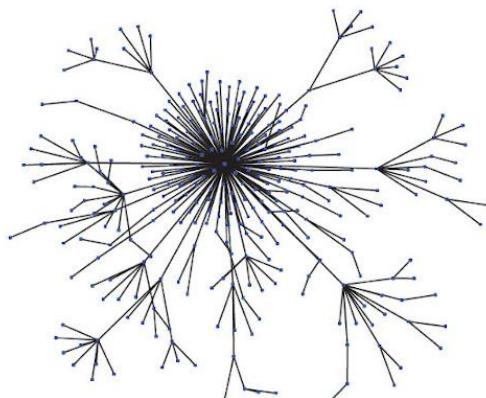
Recommender system



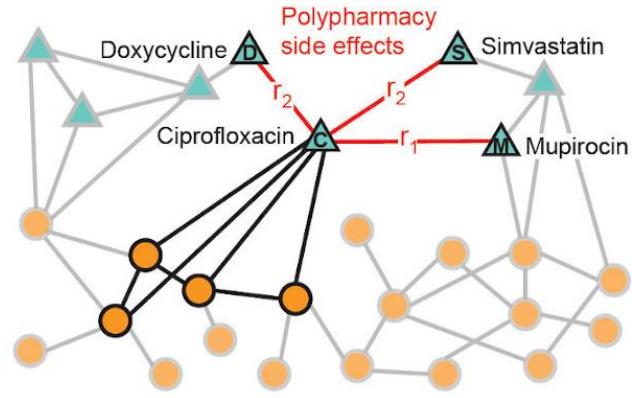
Neutrino detection



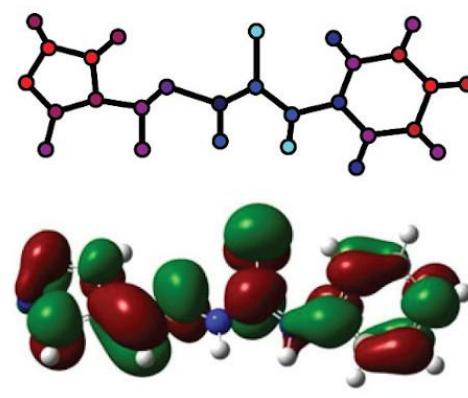
LHC



Fake news detection



Drug repurposing

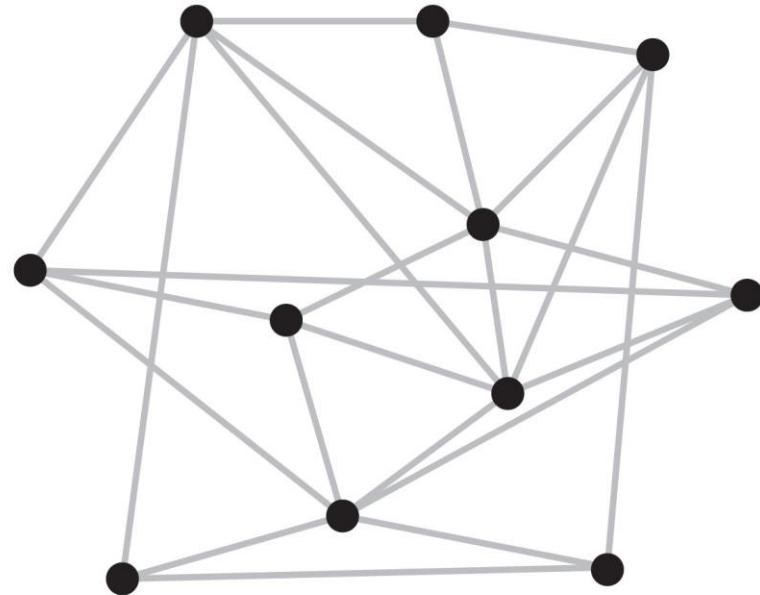


Chemistry

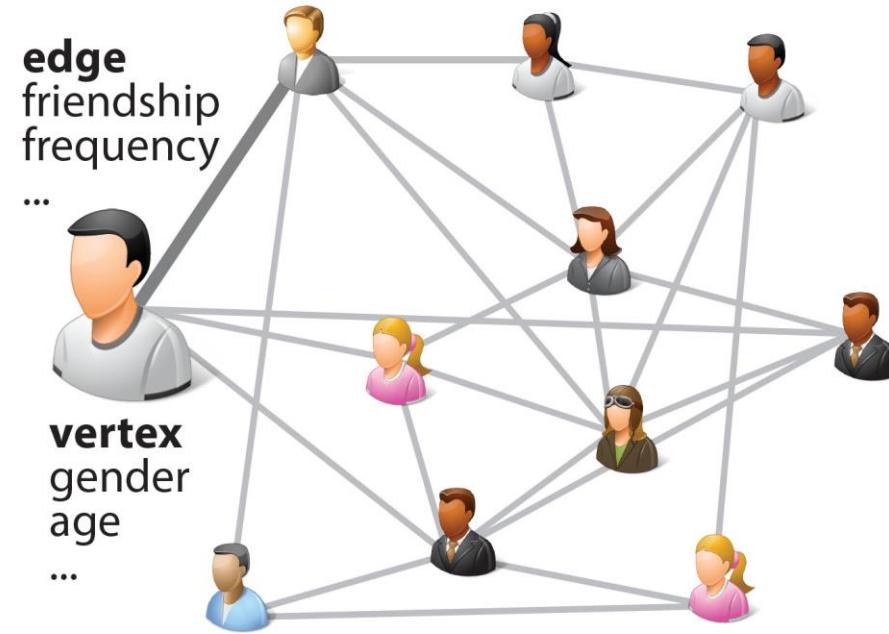
Domain structure:



Signals on a domain:



Domain structure



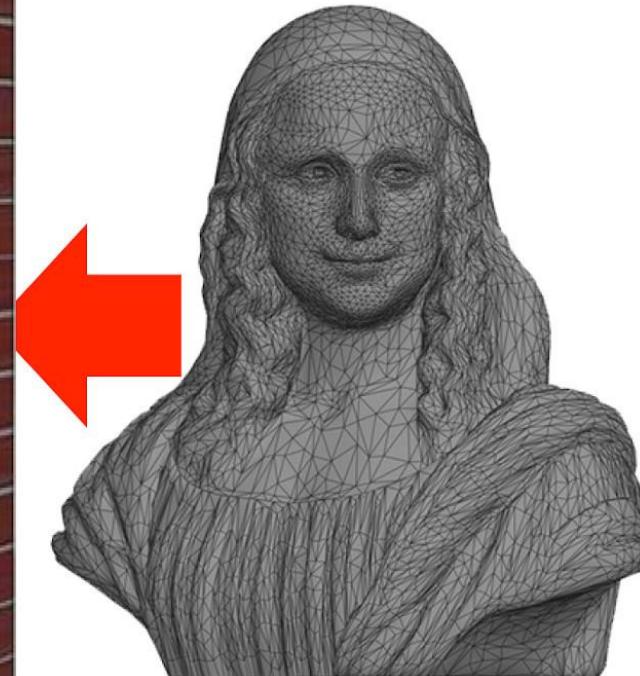
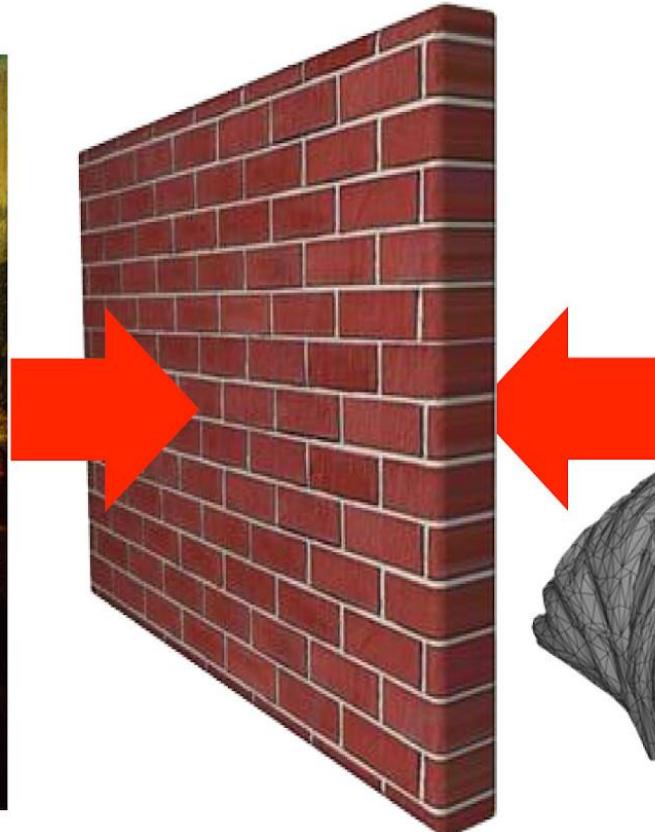
Signals on a domain

Signals vs domain:



2D

Only signal
(fixed domain/template)



3D

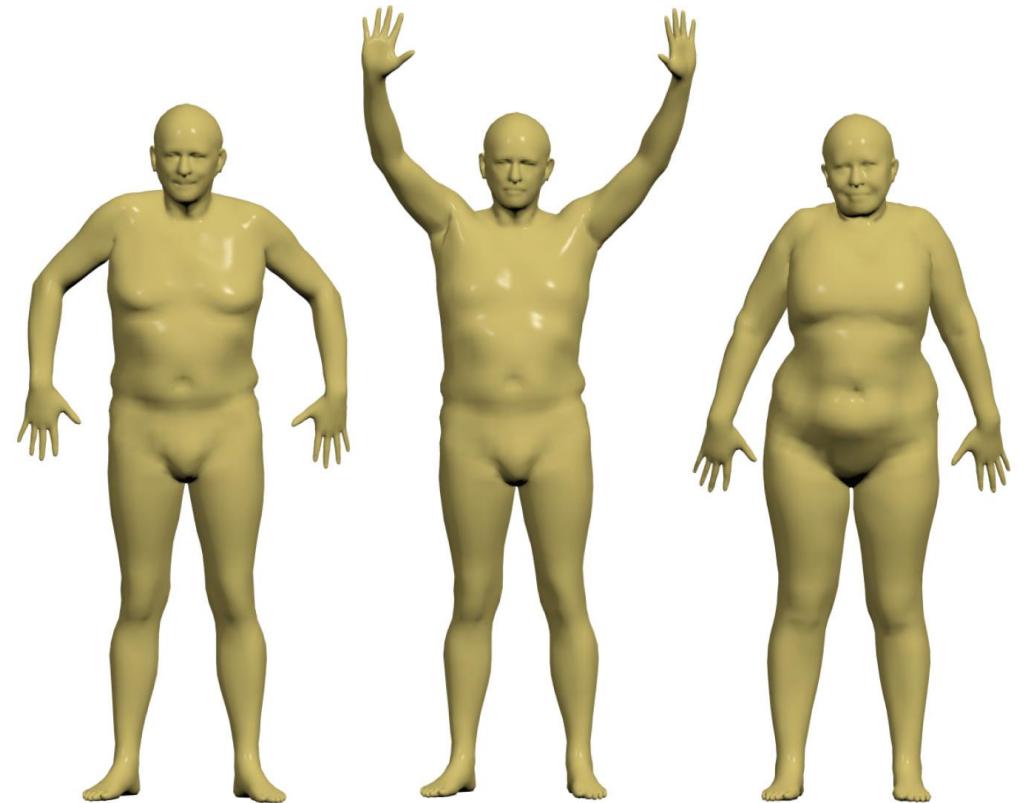
Only domain
(constant/no signal)

Slide credit E. Rodolà

Fixed or non fixed domains:



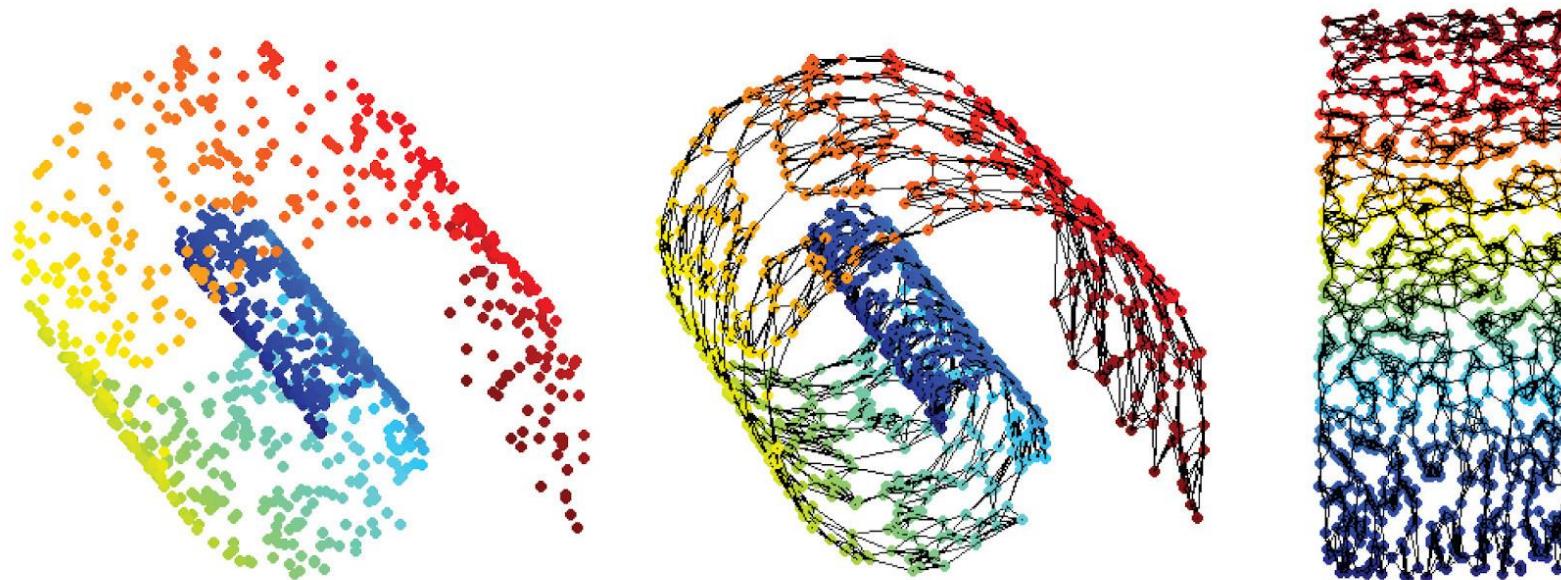
Fixed domain



Deformable domain

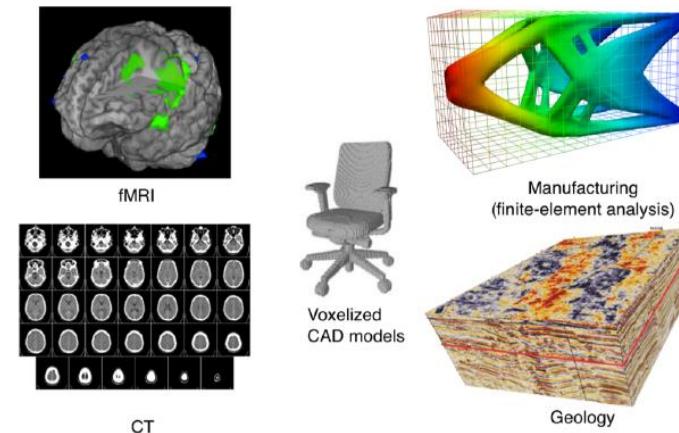
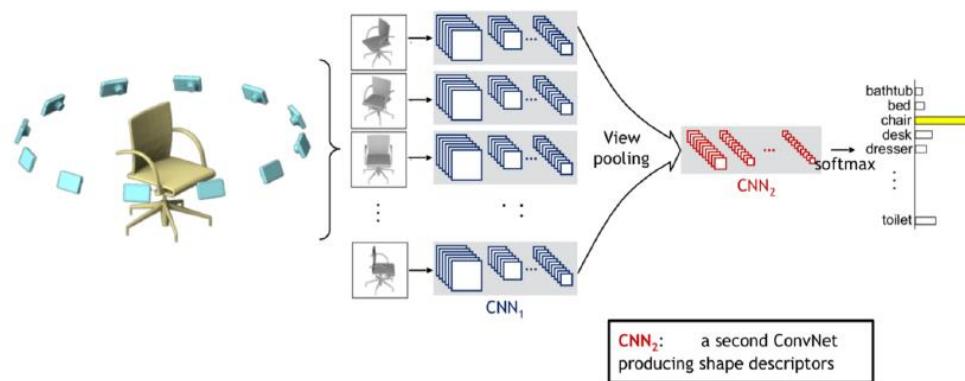
Geometric \neq metric learning:

In manifold learning we look for a (high dimensional) manifold that justifies a given set of data samples

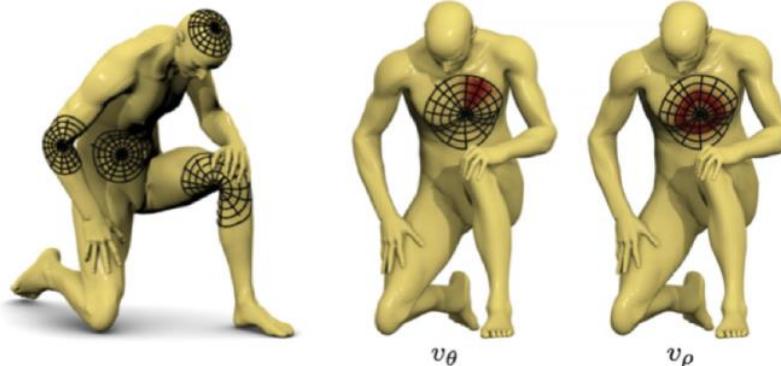


In geometric deep learning the data has a well-known geometric structure

Geometric deep learning

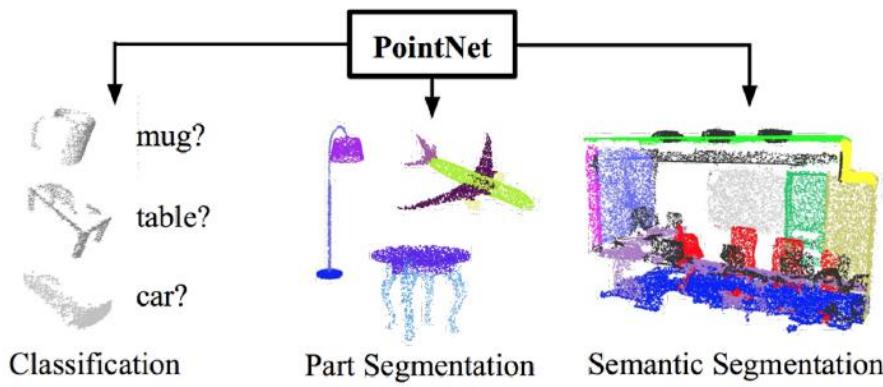


View-based



Intrinsic (surface-based)

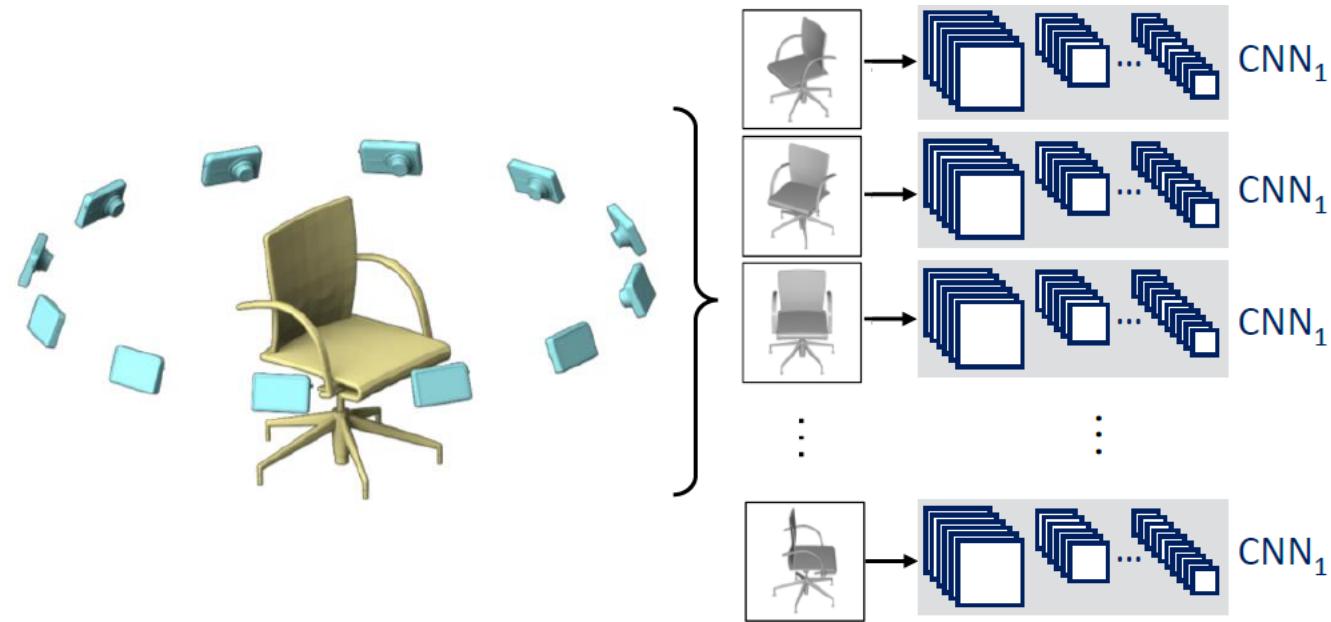
Volumetric



Point-based

Multi-view CNN:

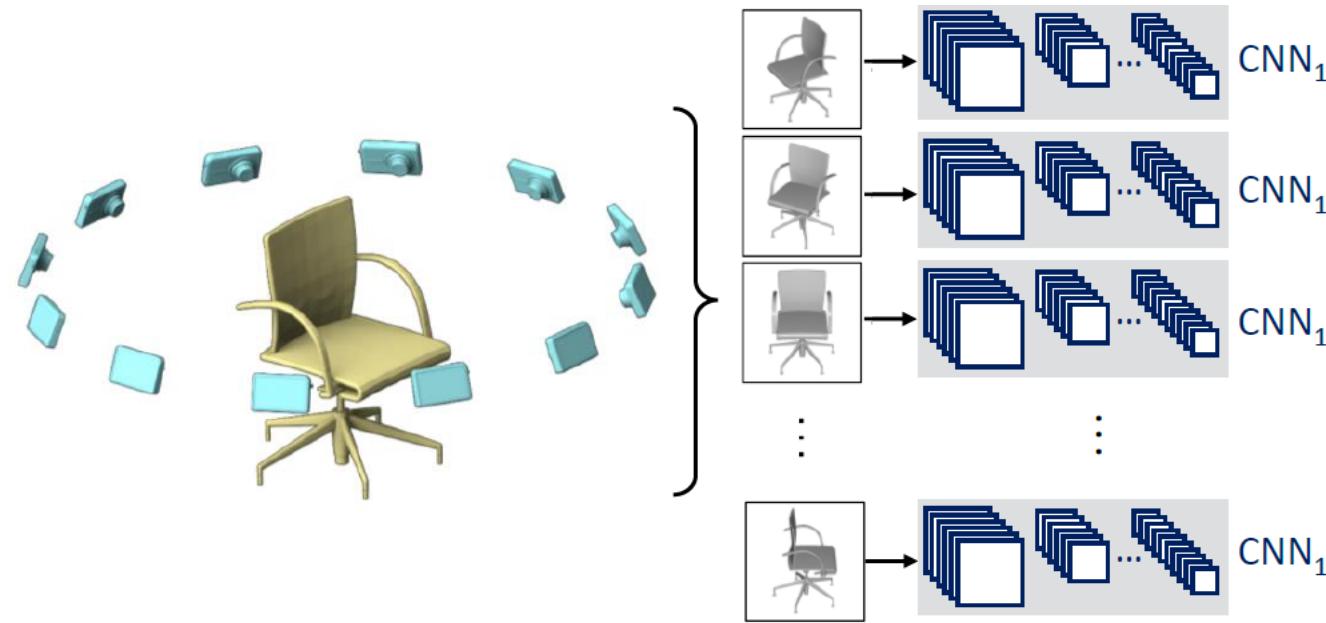
Slide credit E. Rodolà



- Represent 3D object as a collection of range images.

Multi-view CNN:

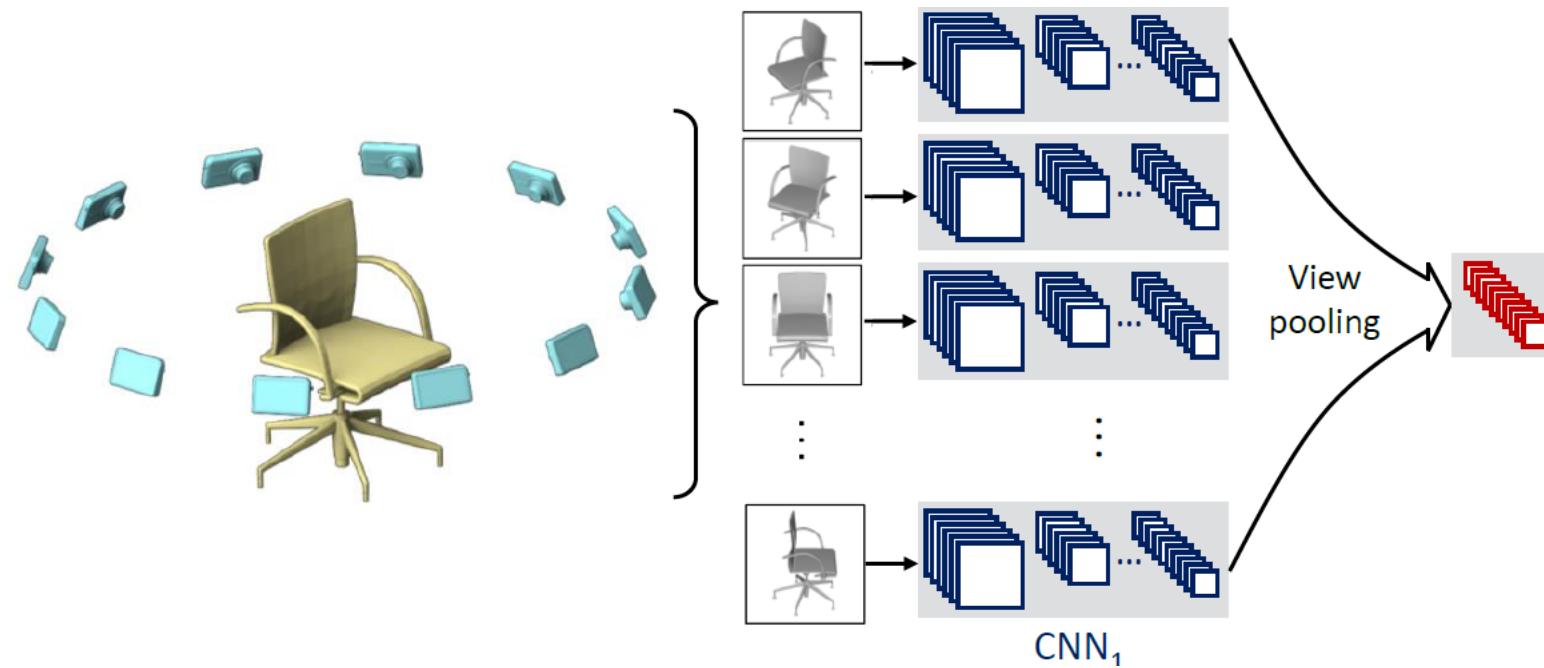
Slide credit E. Rodolà



- Represent 3D object as a collection of range images.
- CNN_1 : extracts image features (views share the parameters)

Multi-view CNN:

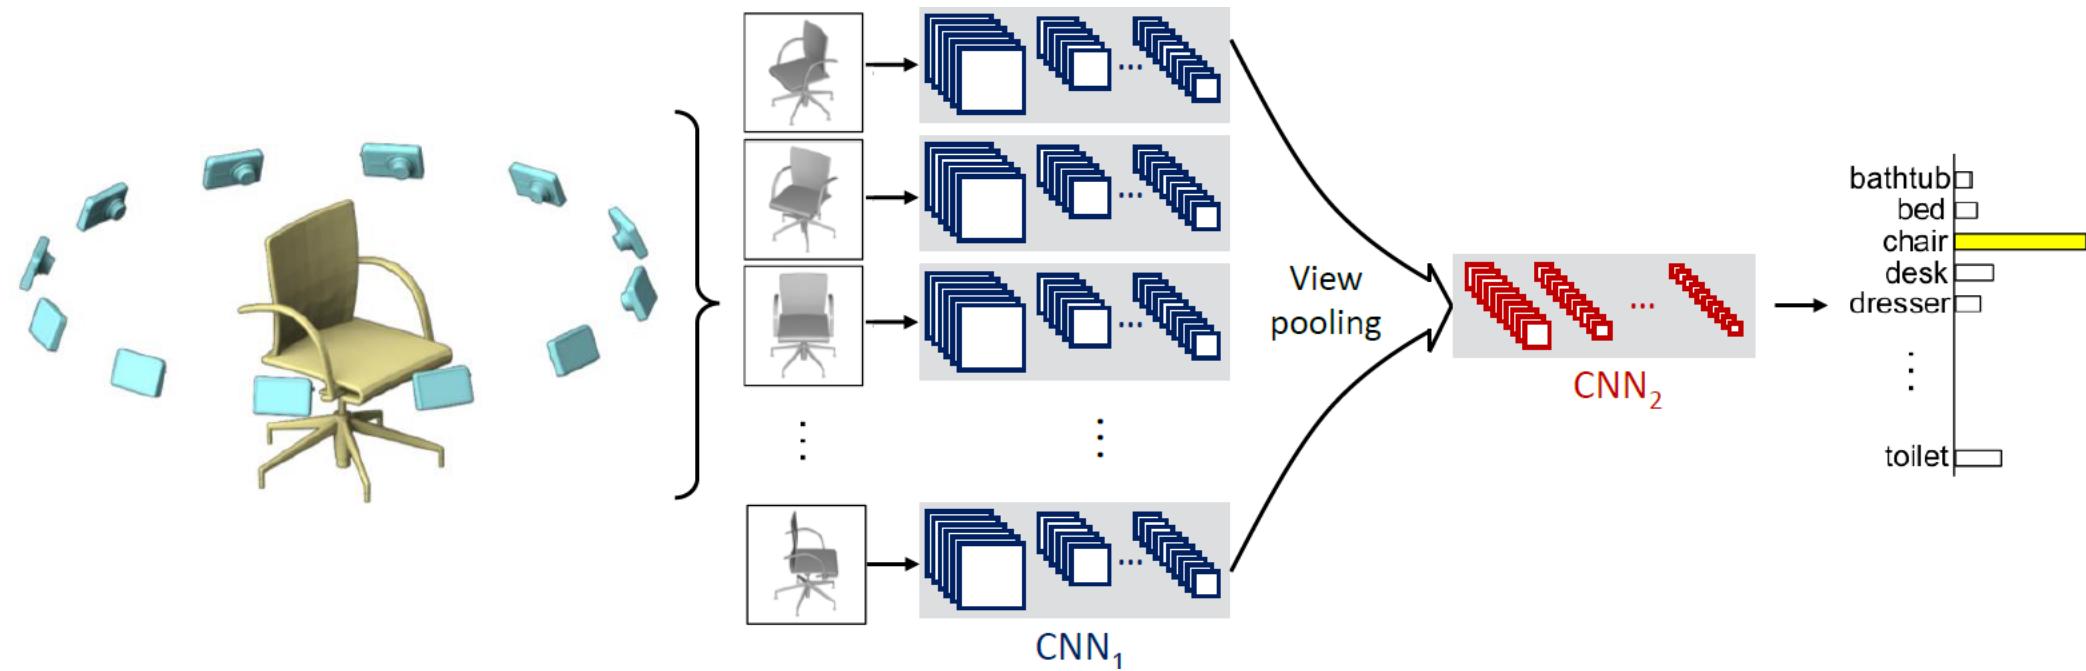
Slide credit E. Rodolà



- Represent 3D object as a collection of range images.
- CNN_1 : extracts image features (views share the parameters)
- Element-wise max pooling across all views

Multi-view CNN:

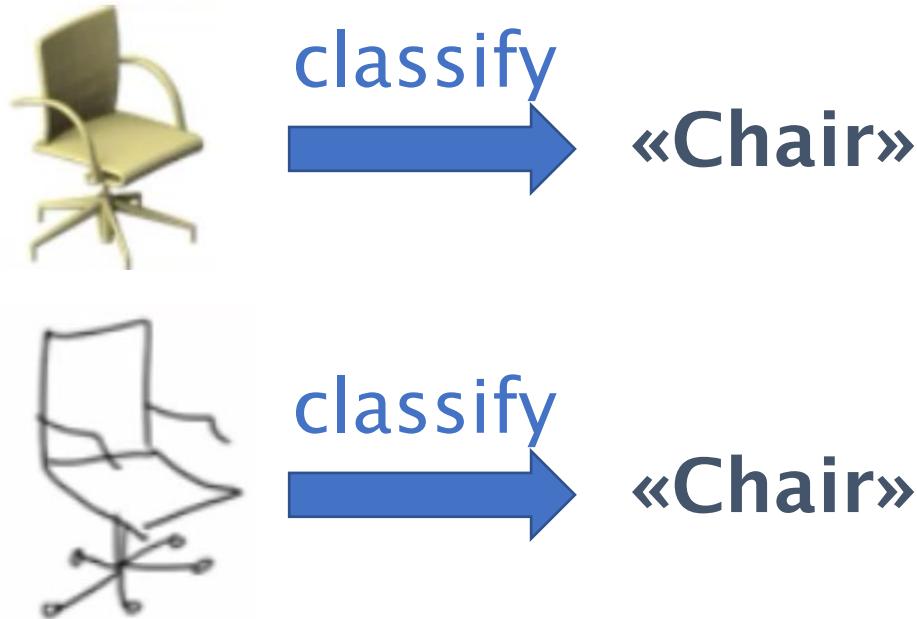
Slide credit E. Rodolà



- Represent 3D object as a collection of range images.
- CNN_1 : extracts image features (views share the parameters)
- Element-wise max pooling across all views
- CNN_2 : produces shape descriptors + final prediction

Multi-view CNN applications

Slide credit E. Rodolà

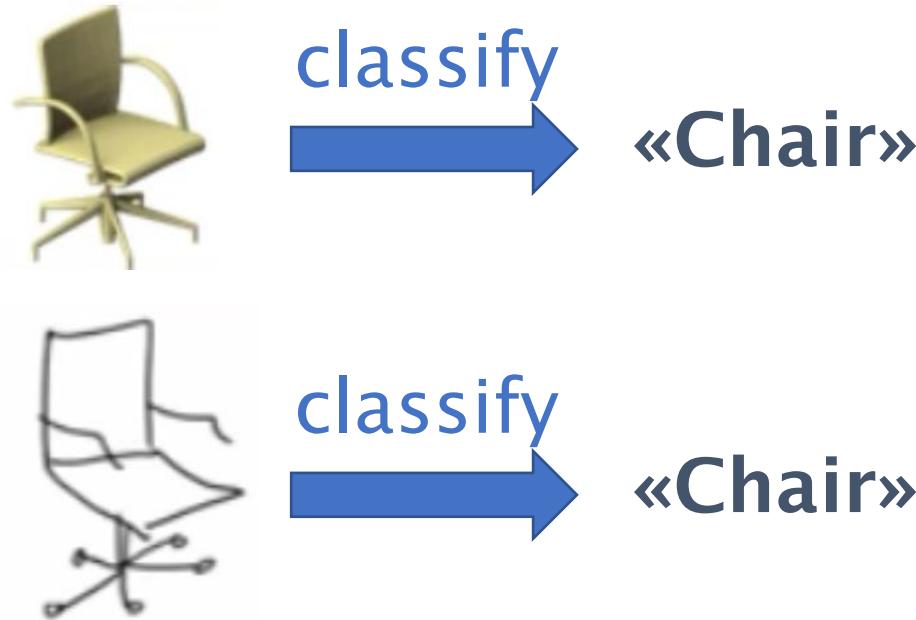


3D shape/sketch classification

- Pretrained on ImageNet
- Fine-tuned on 2D views

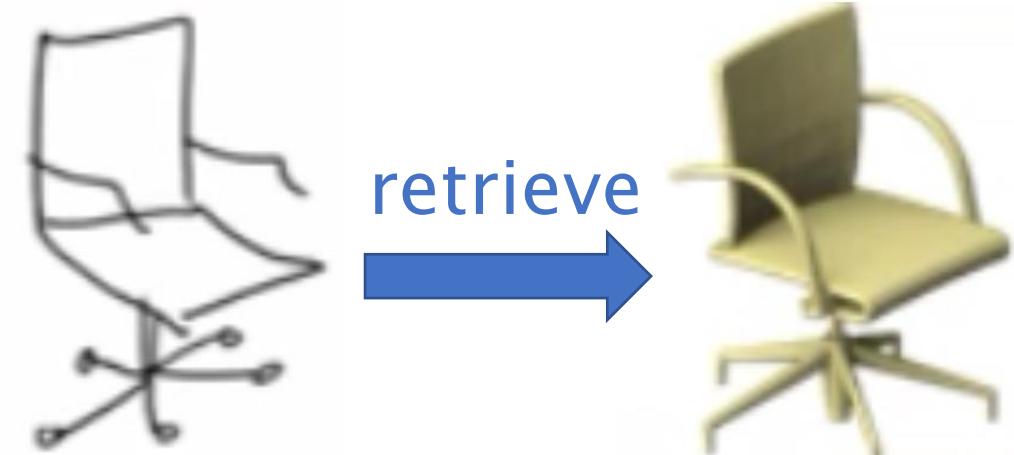
Multi-view CNN applications

Slide credit E. Rodolà



3D shape/sketch classification

- Pretrained on ImageNet
- Fine-tuned on 2D views



Sketch-based shape retrieval

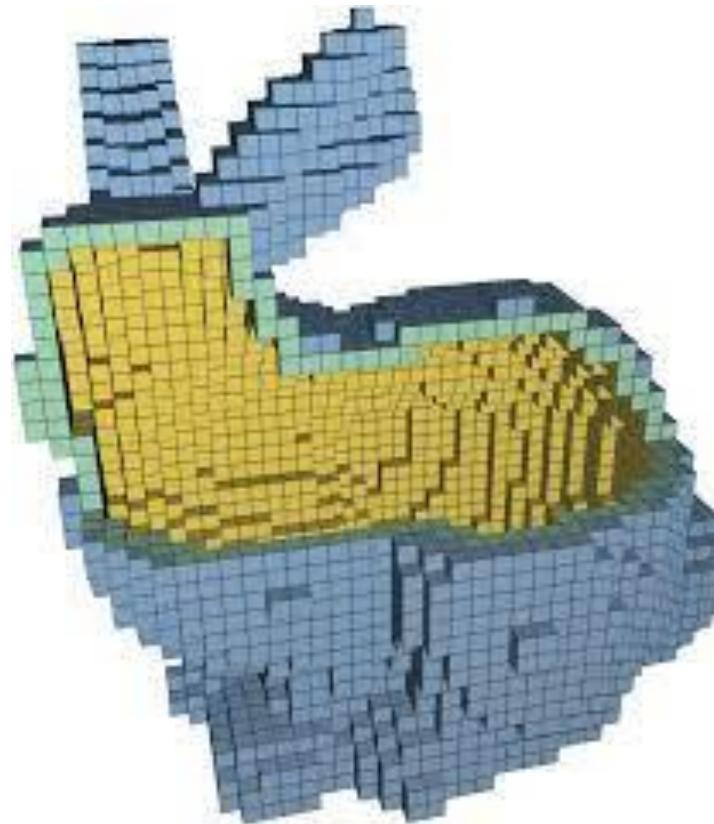
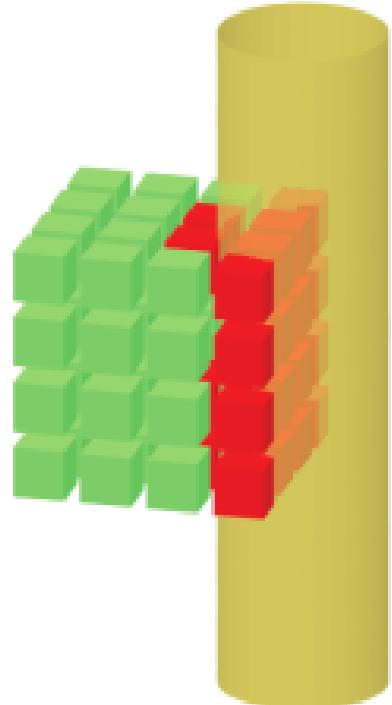
- Render views with hand-drawn style (edge maps)

3D Shape-Net:

Slide credit
E. Rodolà

Volumetric representation
(binary voxels on 3D grid)

3D convolutional network

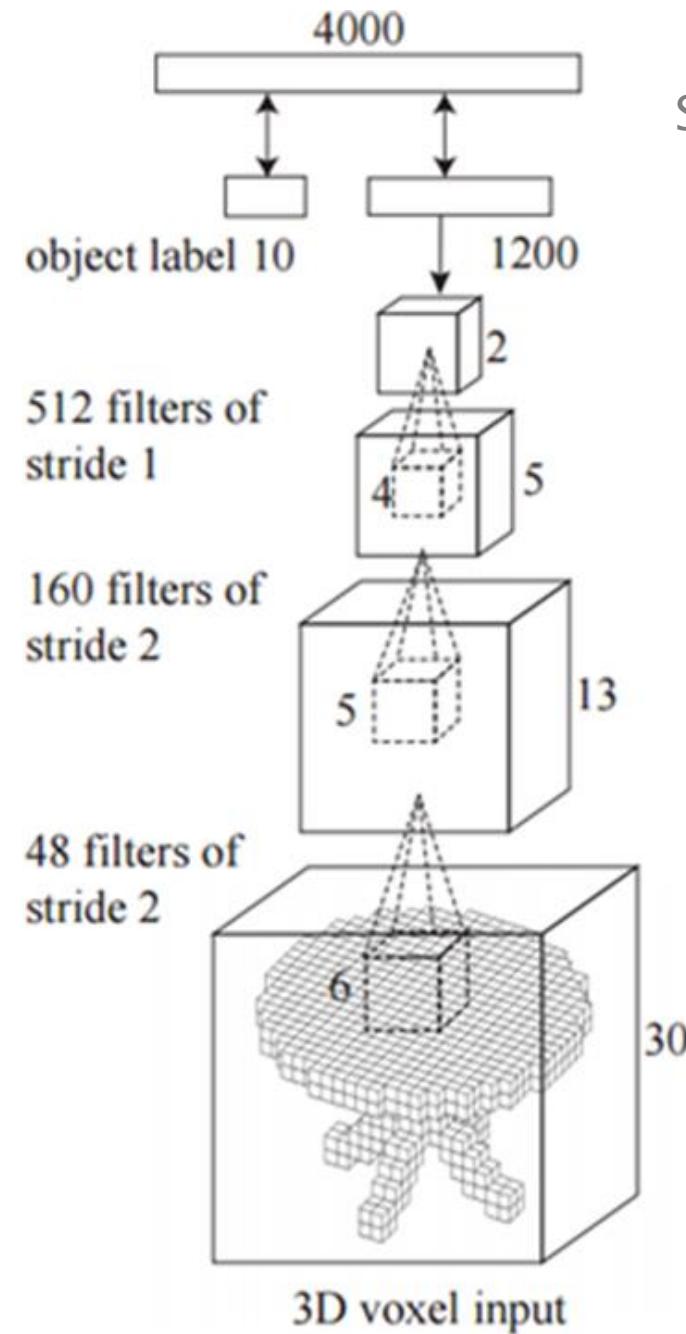
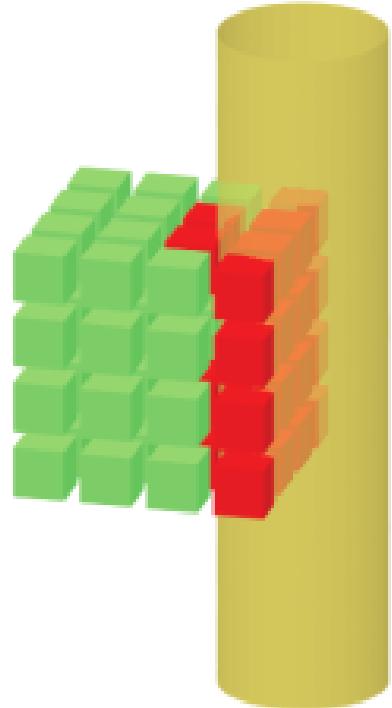


[“3D ShapeNets: A Deep Representation for Volumetric Shapes”, Wu et al., 2015](#)

3D Shape-Net:

Volumetric representation
(binary voxels on 3D grid)

3D convolutional network



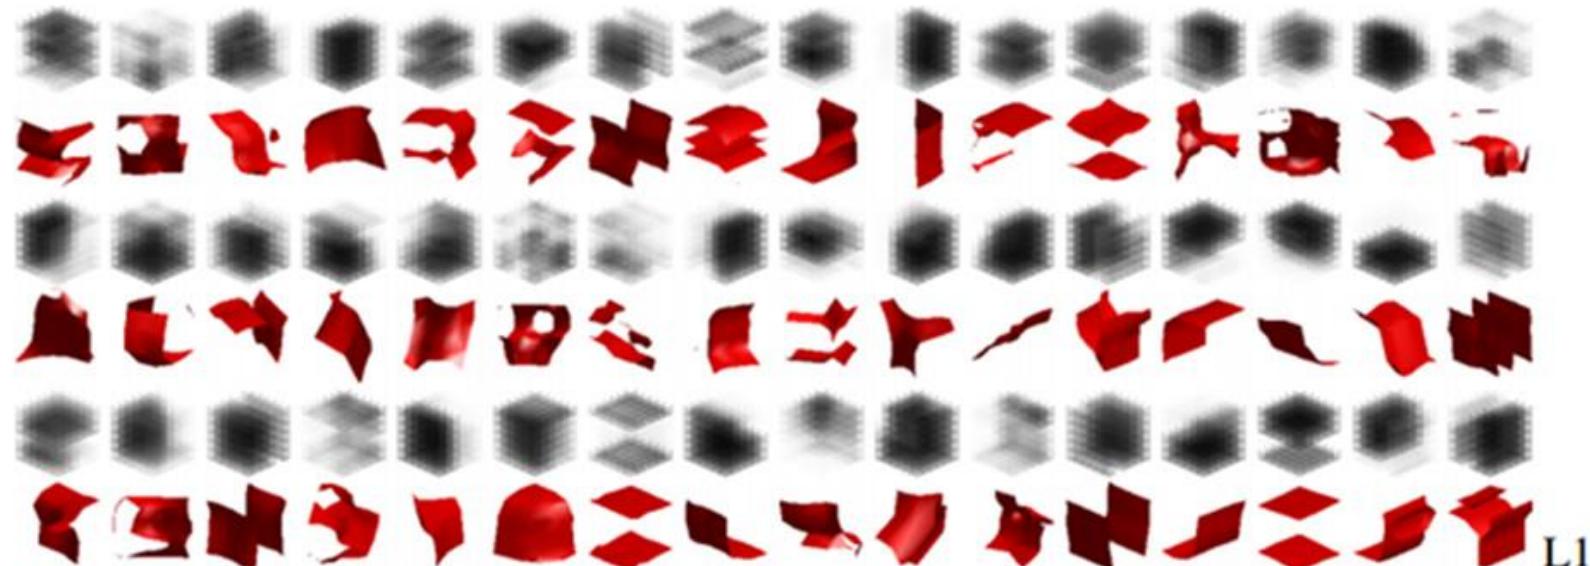
Slide credit
E. Rodolà

[“3D ShapeNets: A Deep Representation for Volumetric Shapes”, Wu et al., 2015](#)

3D Shape-Net learned features:

Slide credit
E. Rodolà

3D primitives

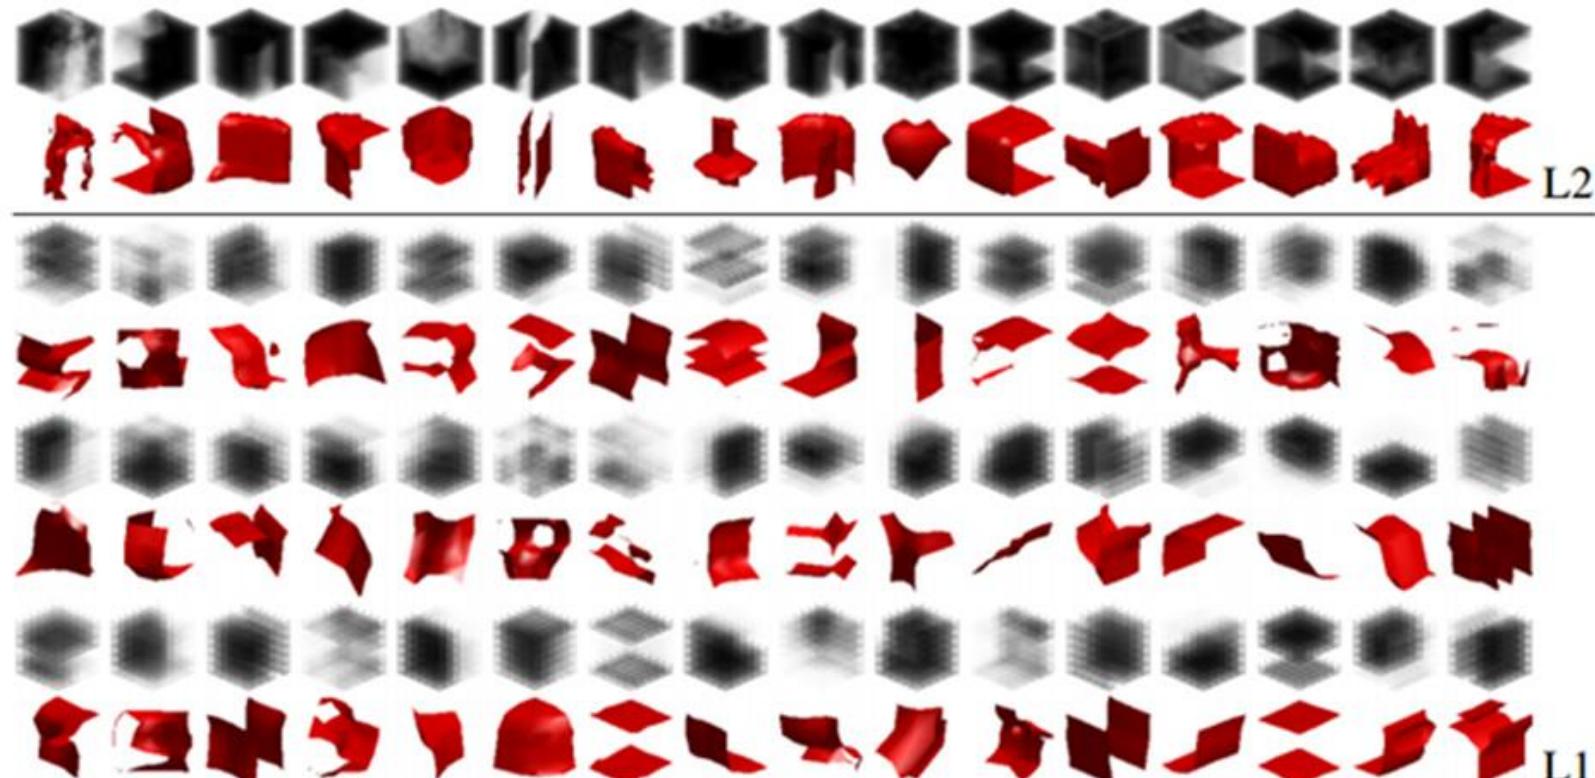


["3D ShapeNets: A Deep Representation for Volumetric Shapes", Wu et al., 2015](#)

3D Shape-Net learned features:

Slide credit
E. Rodolà

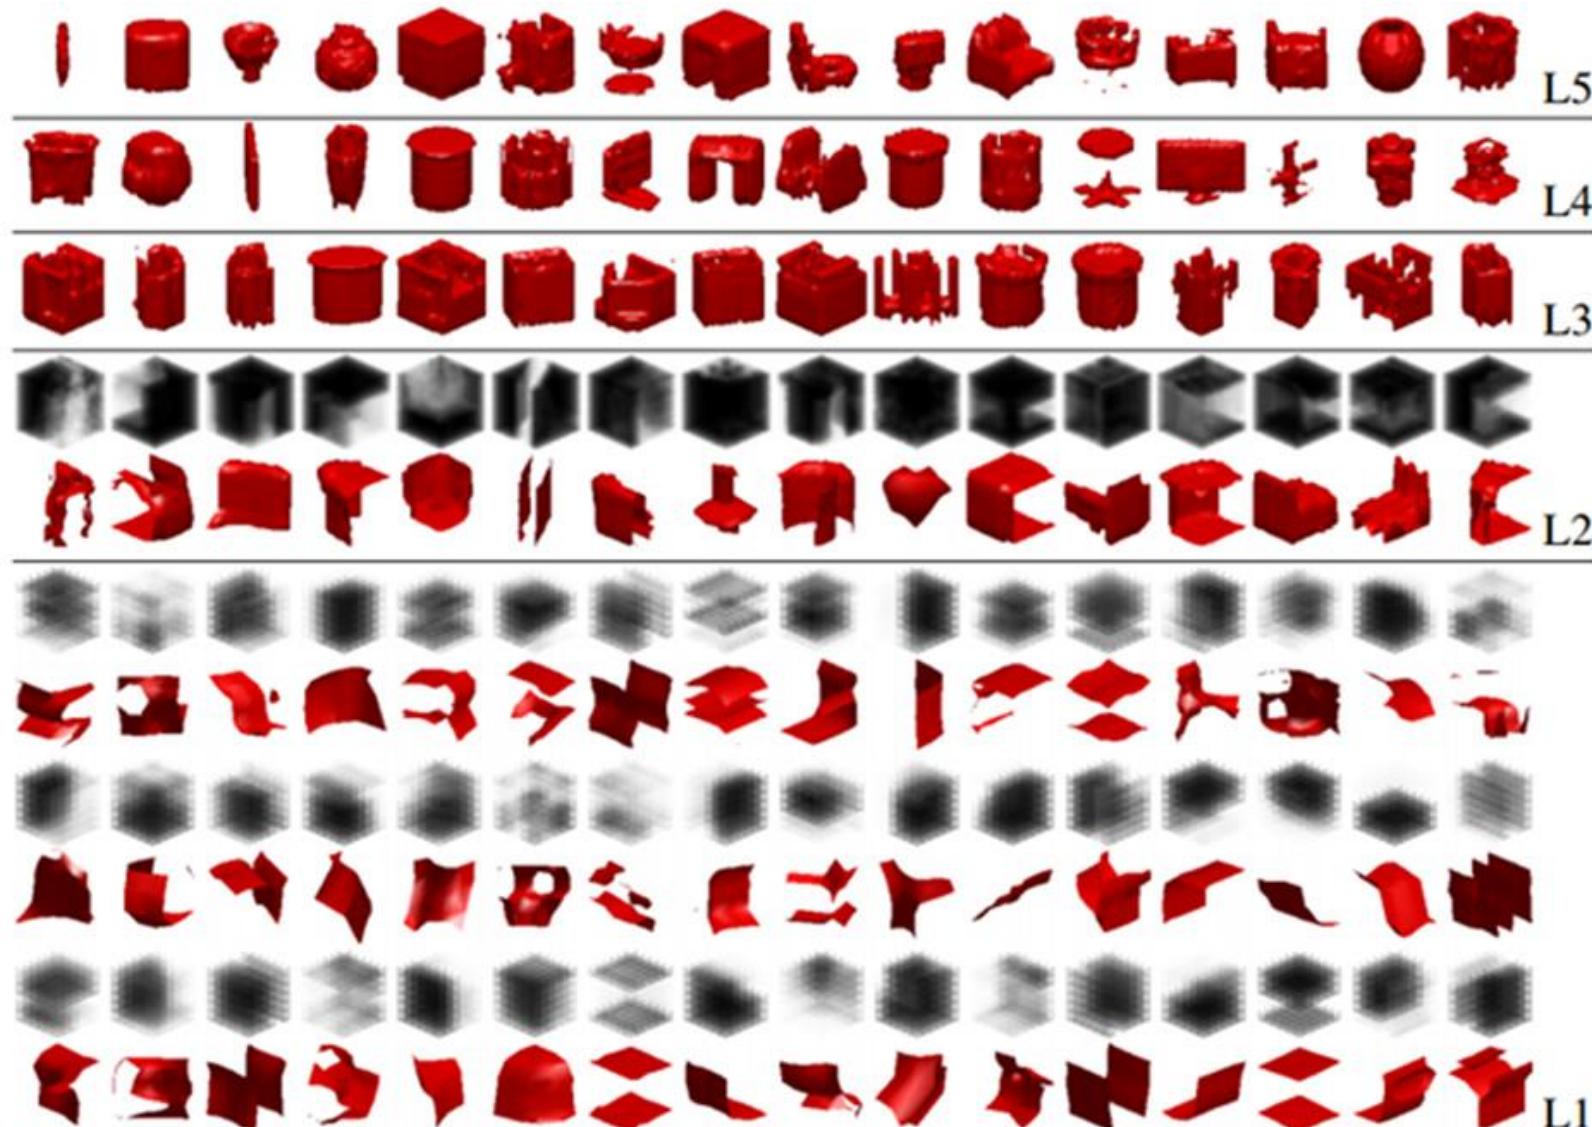
Different scales of features



["3D ShapeNets: A Deep Representation for Volumetric Shapes", Wu et al., 2015](#)

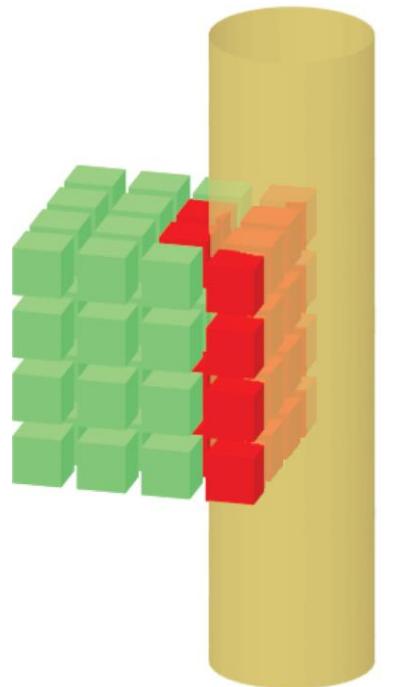
3D Shape-Net learned features:

Slide credit
E. Rodolà

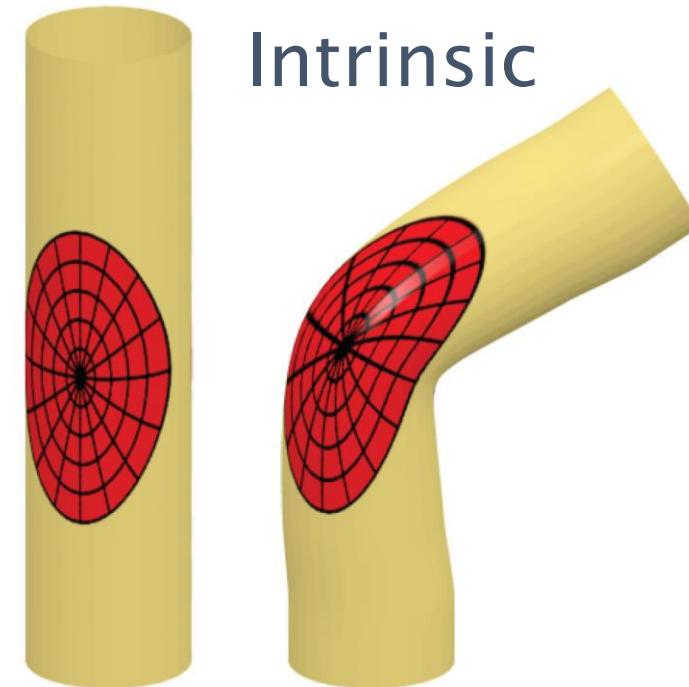
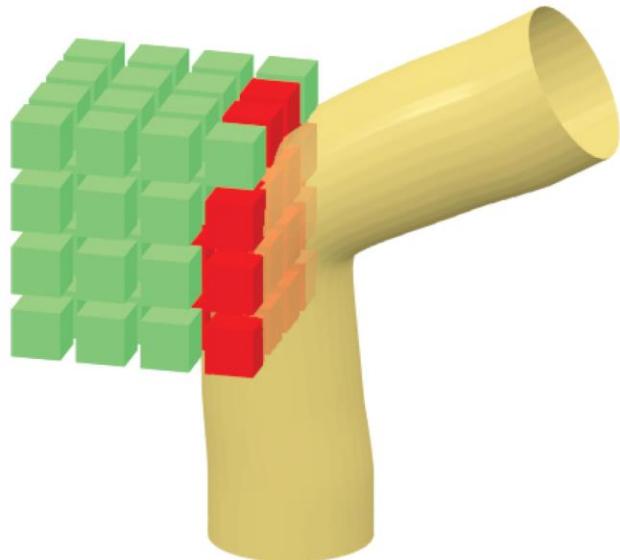


[“3D ShapeNets: A Deep Representation for Volumetric Shapes”, Wu et al., 2015](#)

Challenges of Geometric DL:



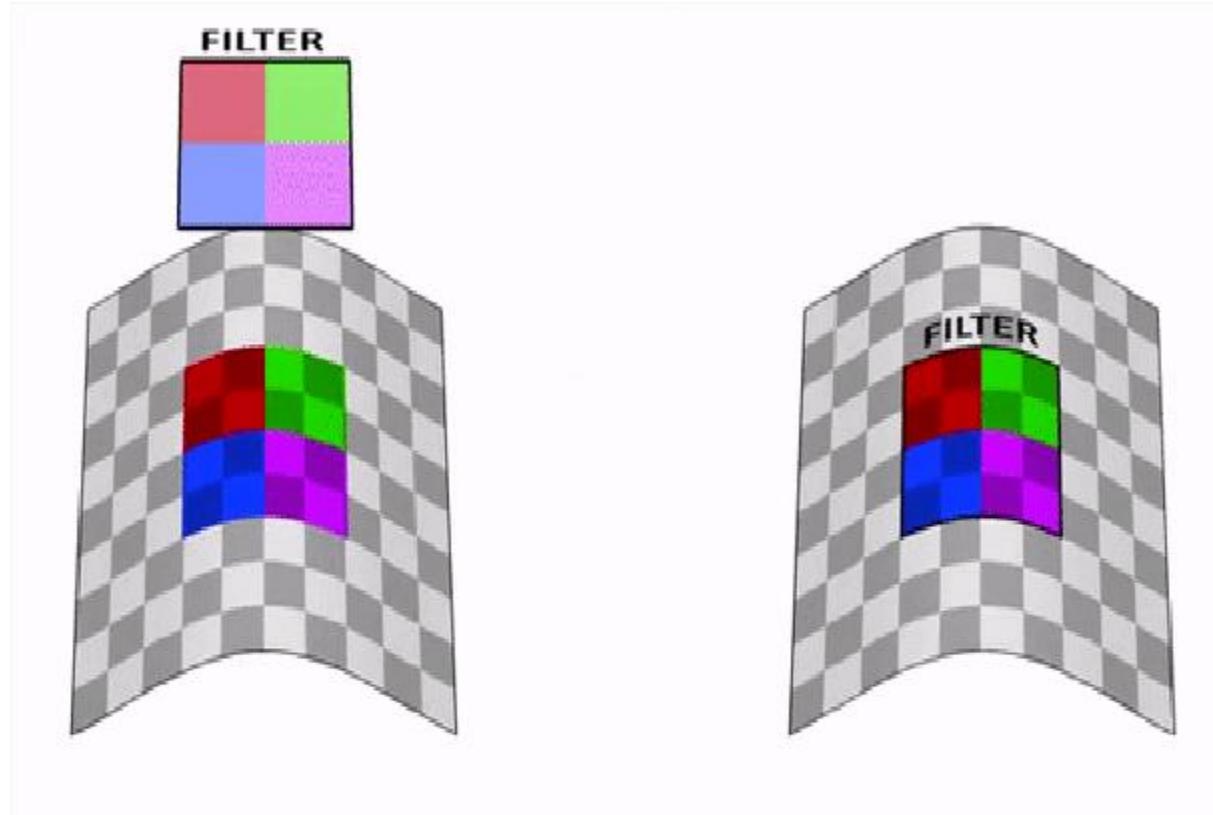
Extrinsic



Intrinsic

- How to define **convolution**?
- How to do **pooling**?
- How to be **efficient**?

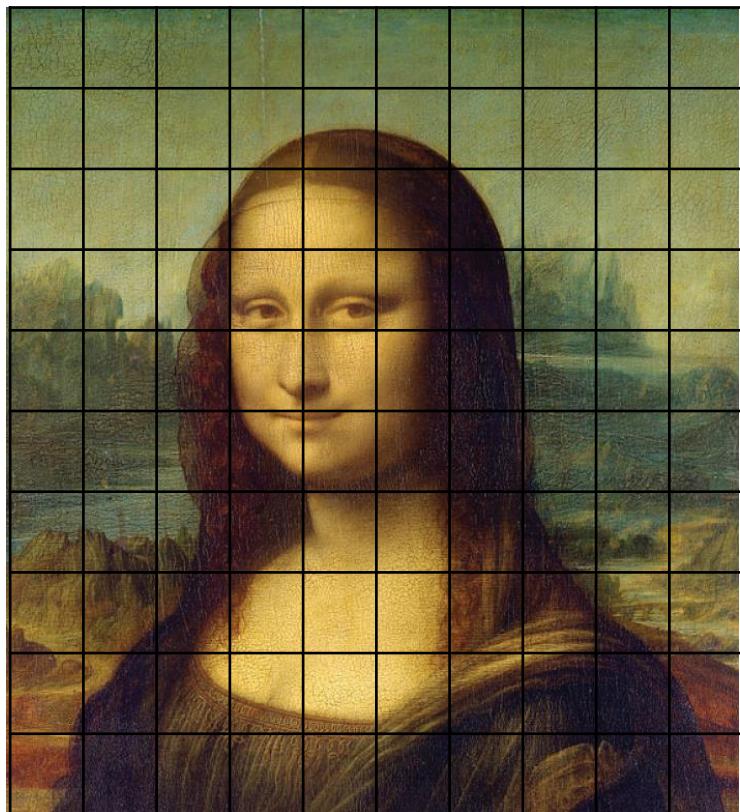
Intrinsic vs extrinsic:



Extrinsic

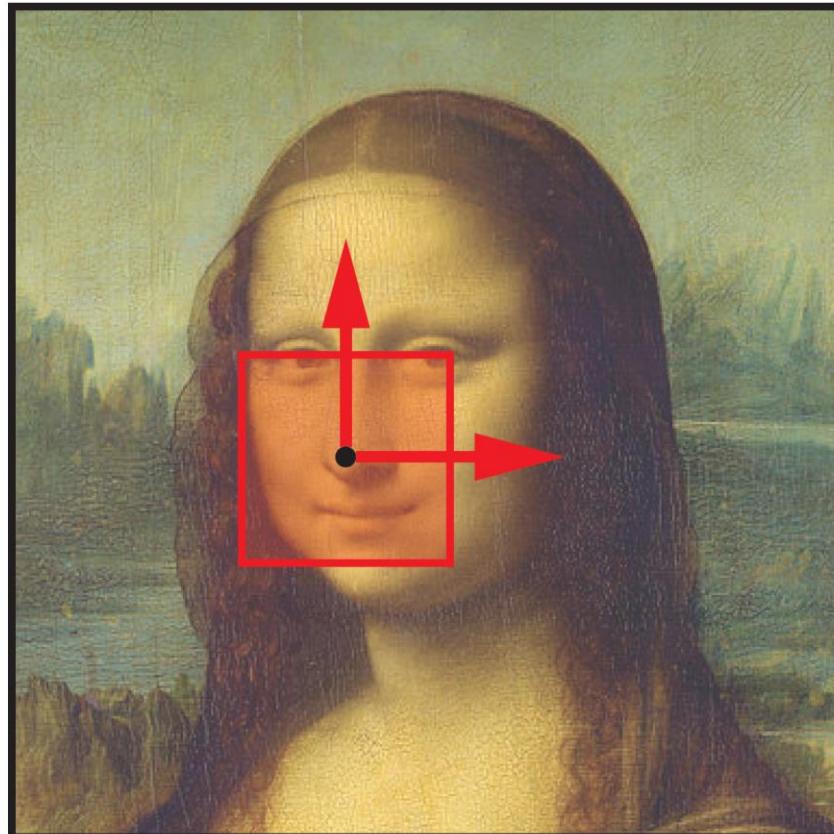
Intrinsic

Fixed pattern vs connectivity:

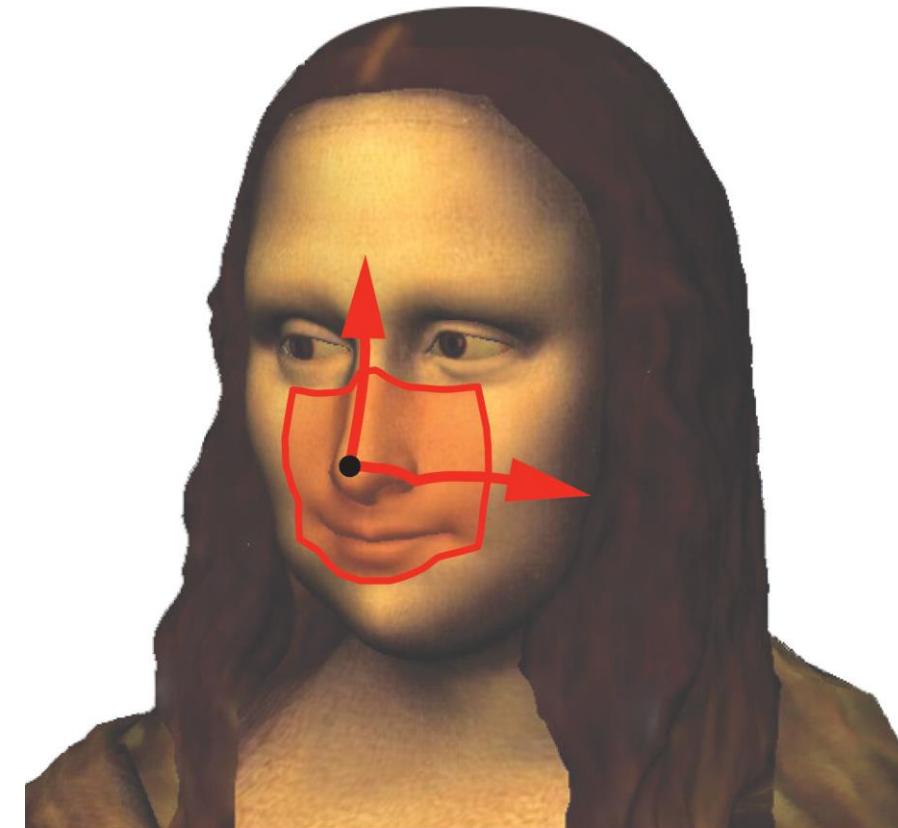


Slide credit E. Rodolà

Non-Euclidean convolution:

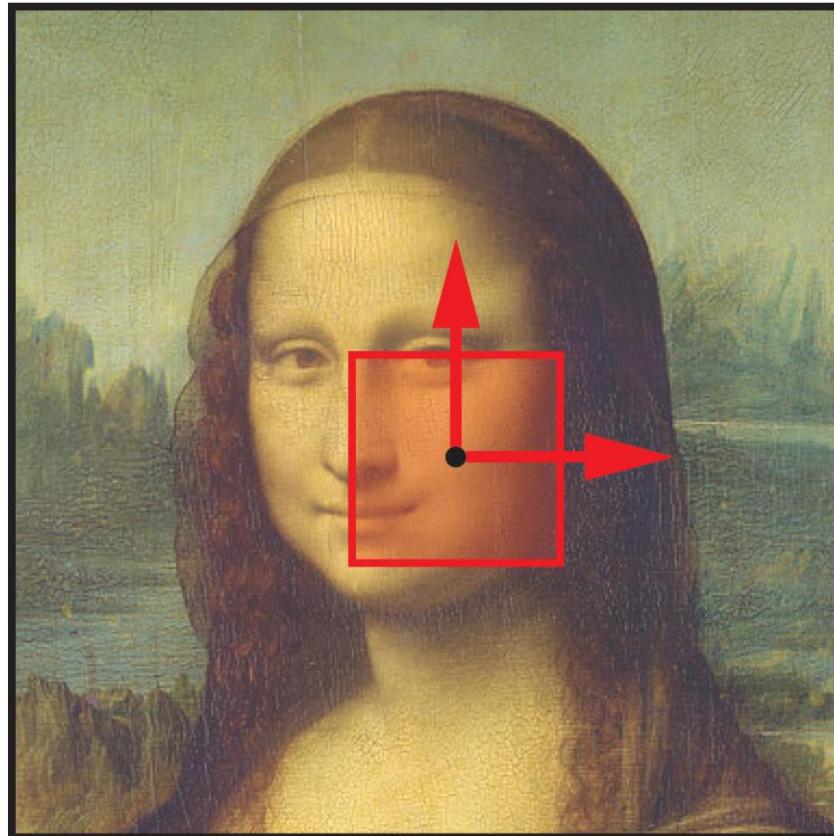


Euclidean

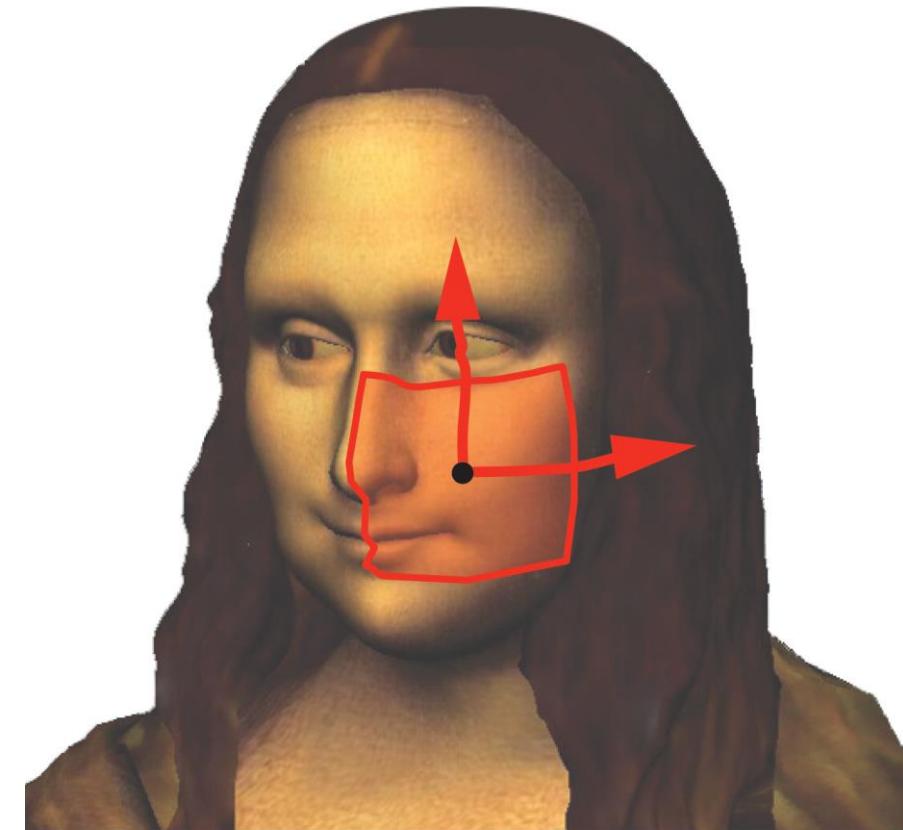


Non- Euclidean

Non-Euclidean convolution:

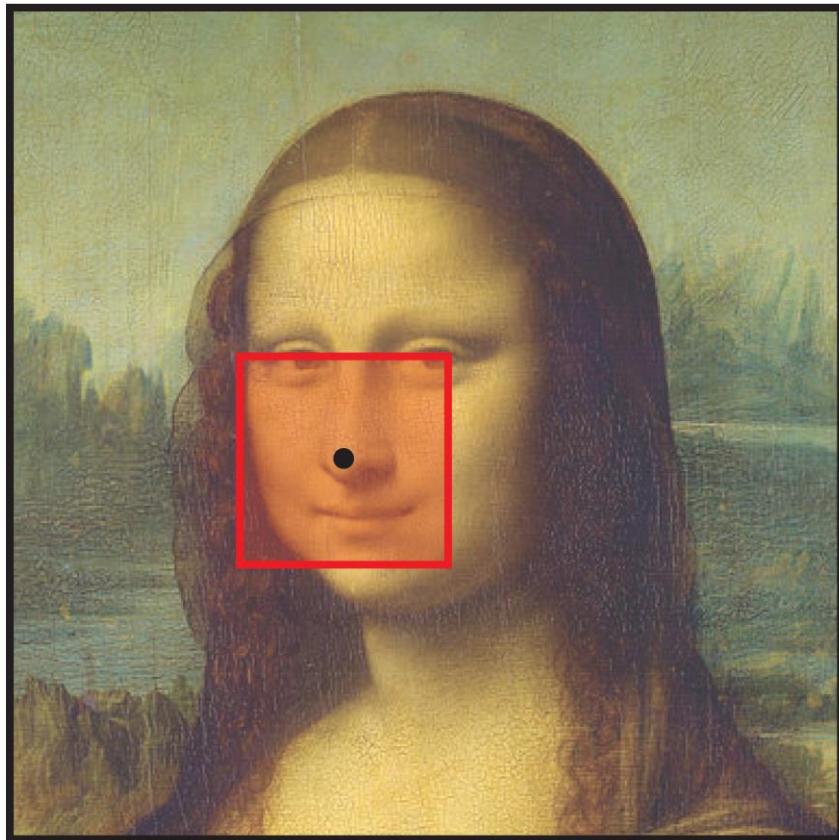


Euclidean

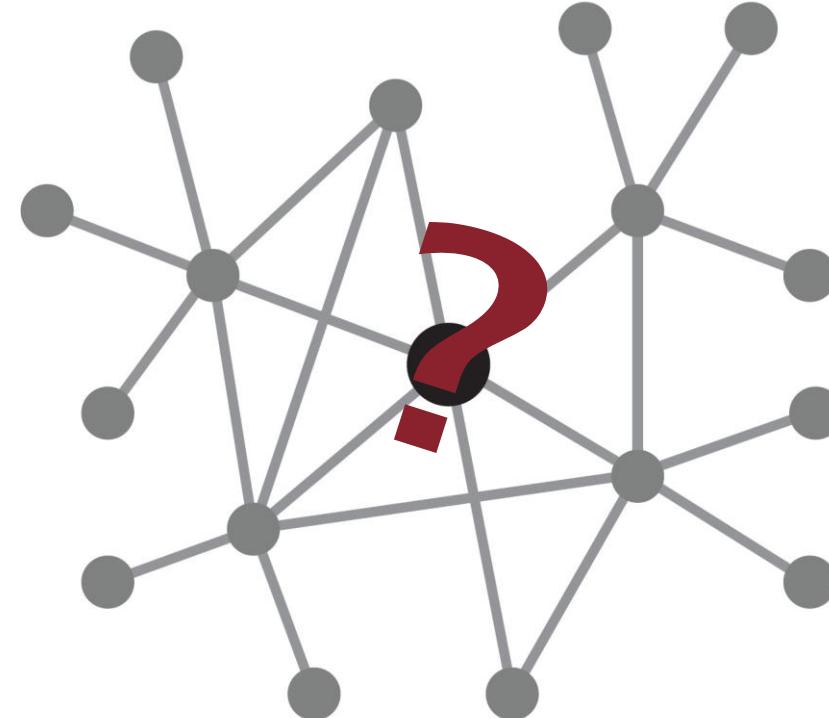


Non- Euclidean

Non-Euclidean convolution:



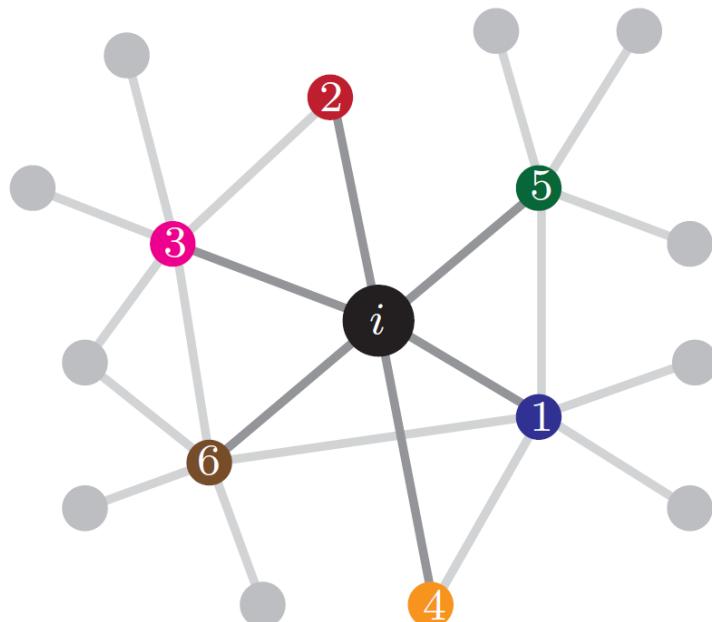
Image



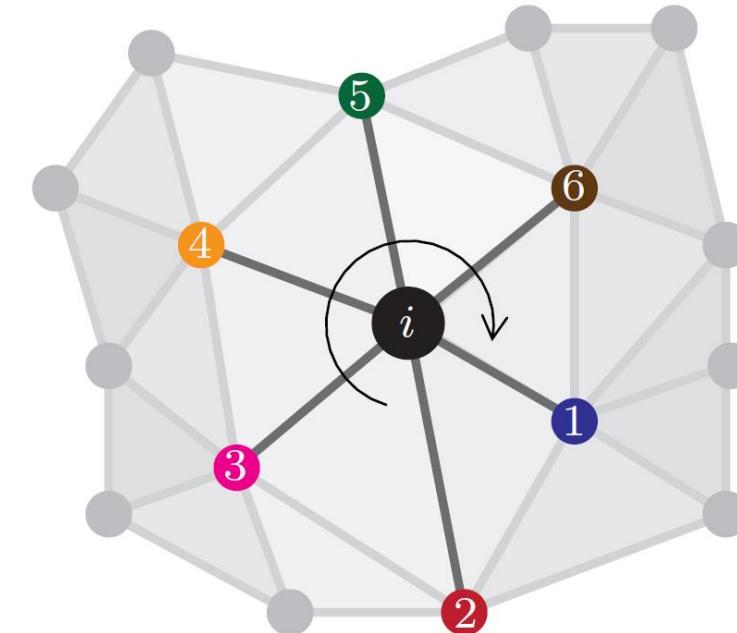
Graph

Local ambiguity

Unlike images, there is not a **canonical ordering** of the points in the domain



Graph rotation

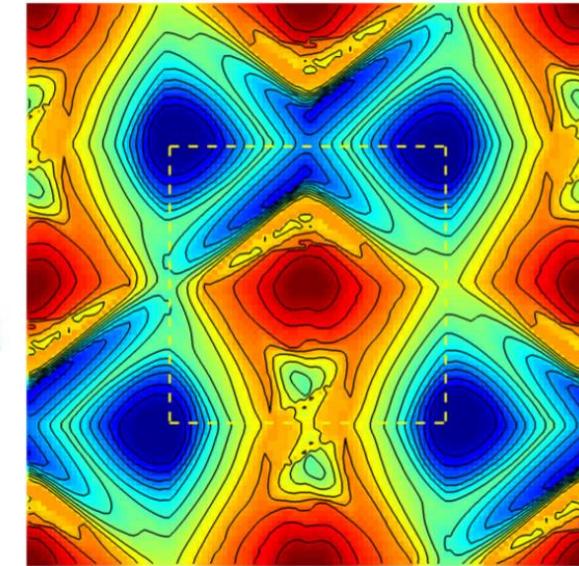
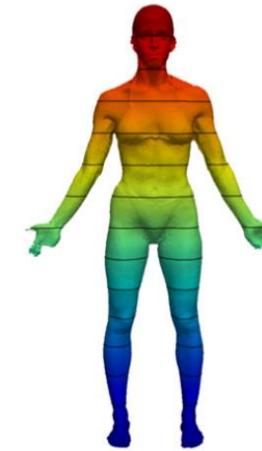
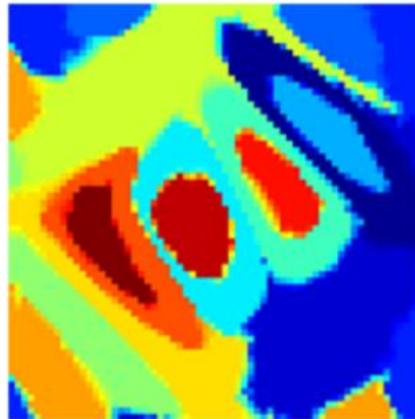


Mesh rotation

Global parametrization:

Slide credit E. Rodolà

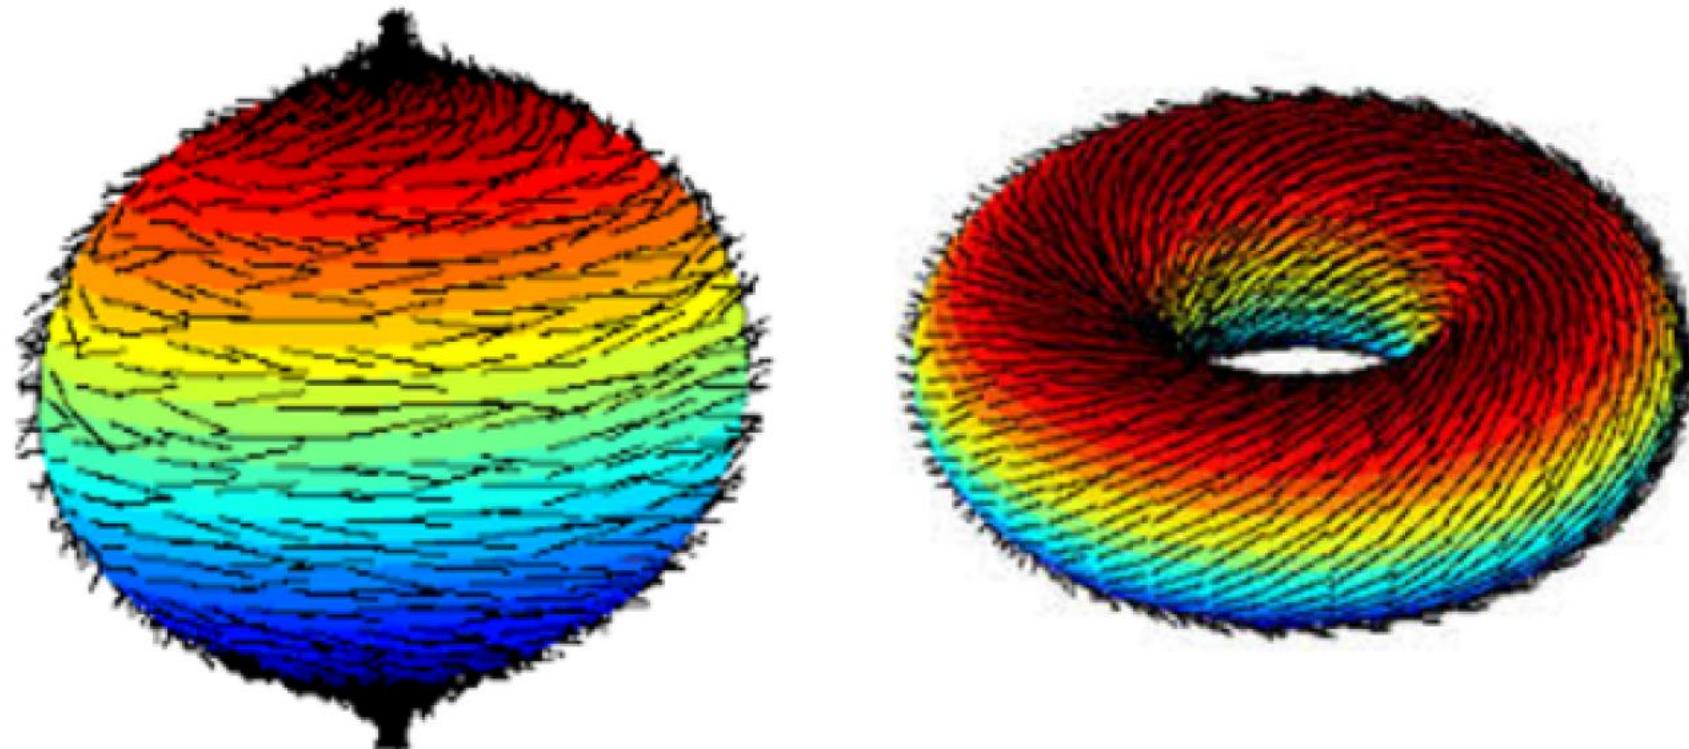
We can map a geometric object in 2D and:



- Apply Euclidean techniques in the embedding space
- Provide invariance to certain transformations
- Not-unique parametrization
- Distortion rised by the embedding map

Directly on manifolds:

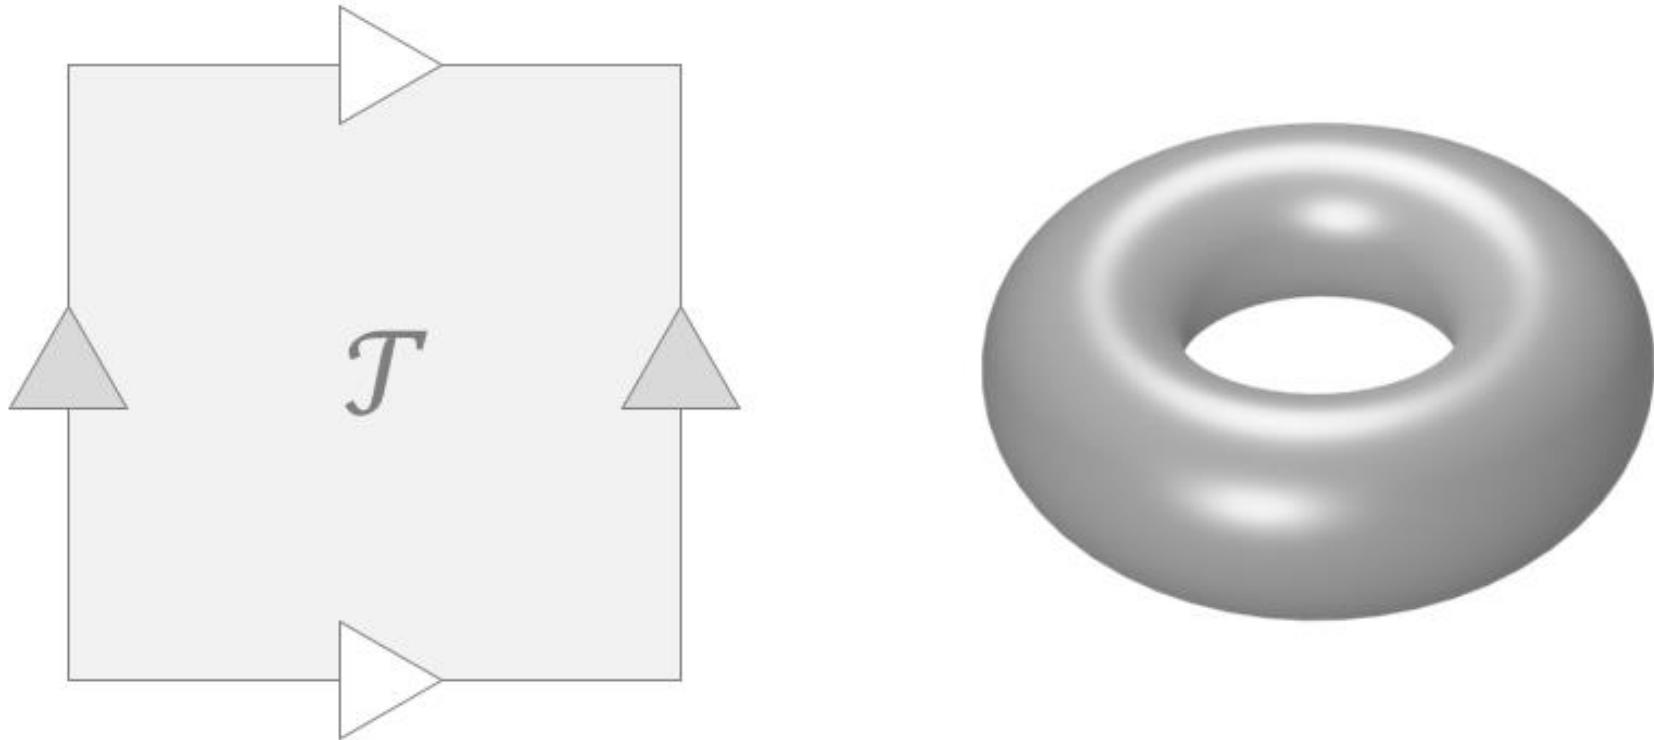
Translation invariant convolution is not possible in general on surfaces due to the singularities in translation fields (proved by the «Hairy ball» theorem from Poincaré-Hopf)



Directly on manifolds:

Slide credit E. Rodolà

The torus is the only closed and orientable surface admitting a translational group



Spatial convolution:

Slide credit E. Rodolà

Local system of coordinates
(geodesic polar around i)

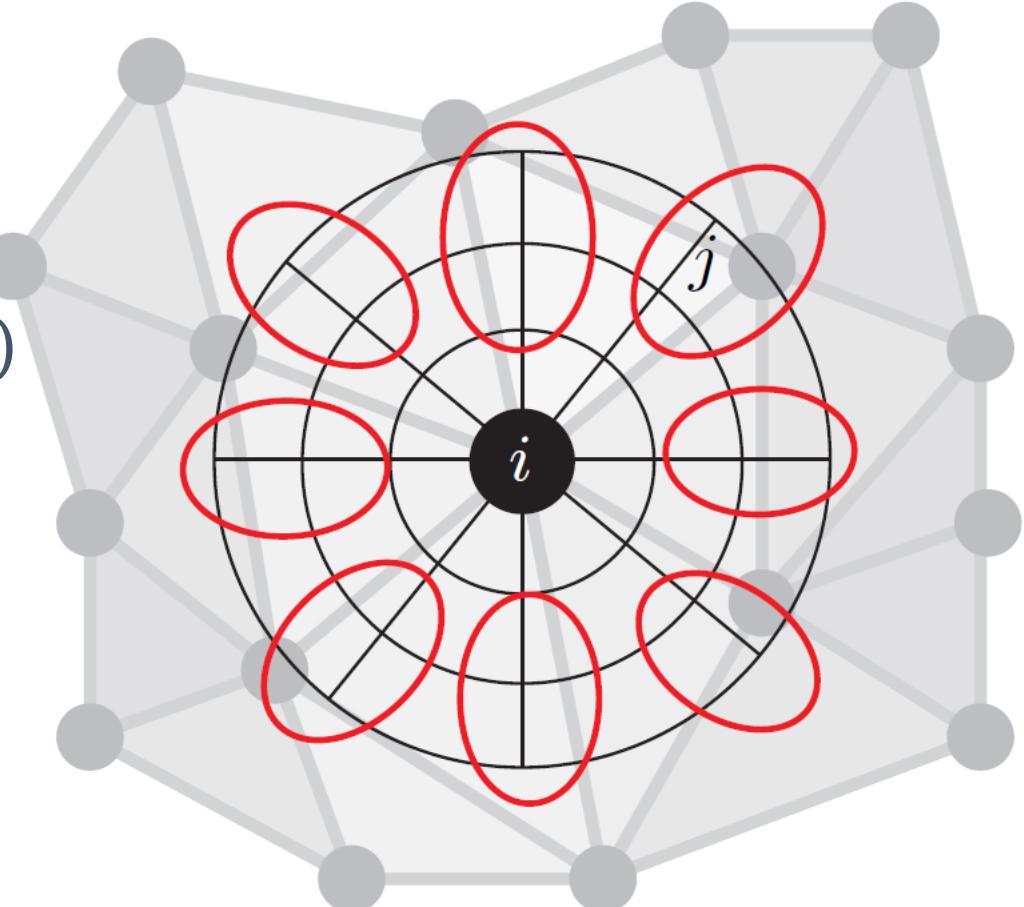
Local Weights $w(u_{ij})$ (Gaussians)

$$w(u_{ij}) = \exp(-(u_{ij} - \mu)^\top \Sigma^{-1} (u_{ij} - \mu))$$

Input signal: $\mathbf{f} \in \mathbb{R}^n$

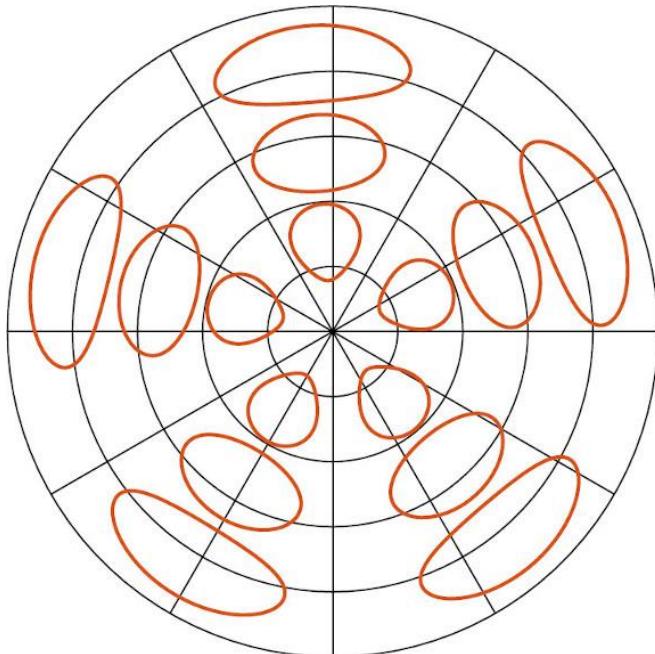
Learnable filter: $\mathbf{g} \in \mathbb{R}^n$ (local)

Apply the filter: $\mathbf{f}^\top \mathbf{g} \in \mathbb{R}^n$ (local)

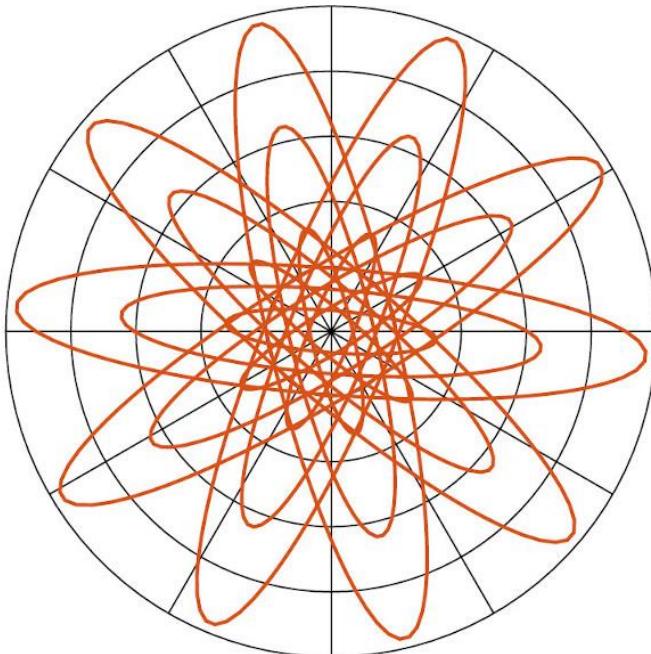


Spatial convolution:

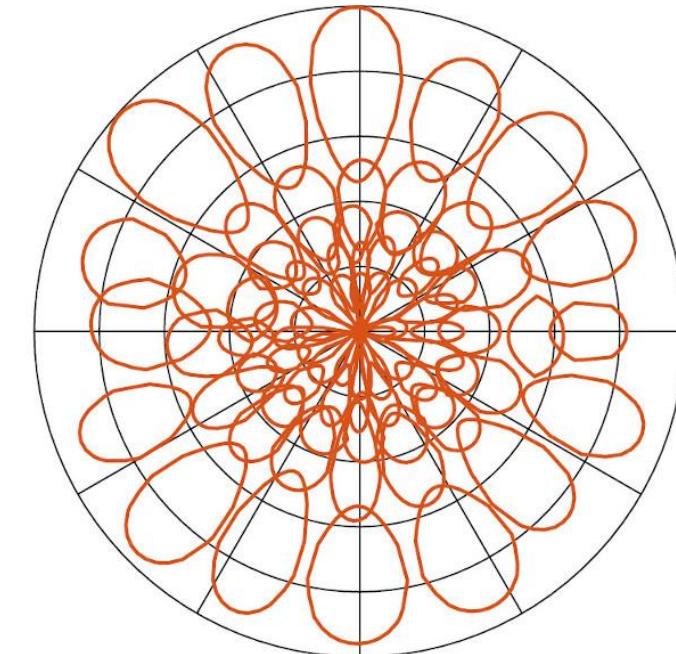
Slide credit E. Rodolà



GCNN



ACNN



MoNet

Fourier analysis

The Fourier coefficients depend on the Global geometry of the surface

We want to compute pointwise descriptors



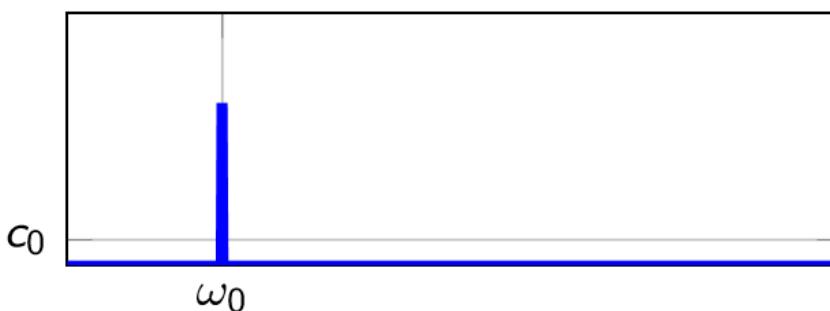
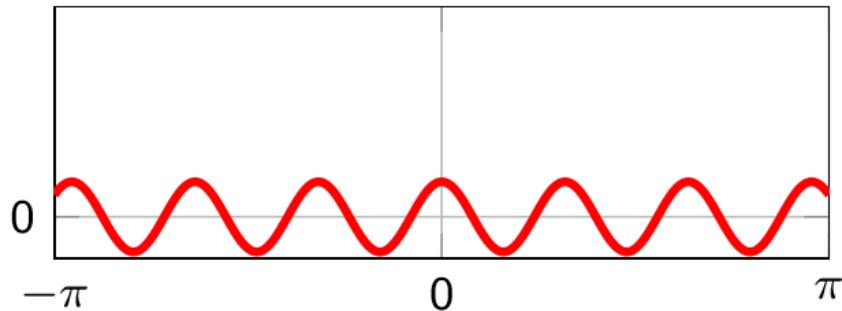
We would like to enforce LOCALIZATION of the Fourier analysis

$$f(x) = \sum_{k \geq 1} \int_X f(\xi) \phi_k(\xi) d\xi \quad \phi_k(x)$$

$\mathcal{F}(f)_k = \langle f, \phi_k \rangle_{L^2(X)}$

The only characterization of these coefficients is the frequency that each of them is representing

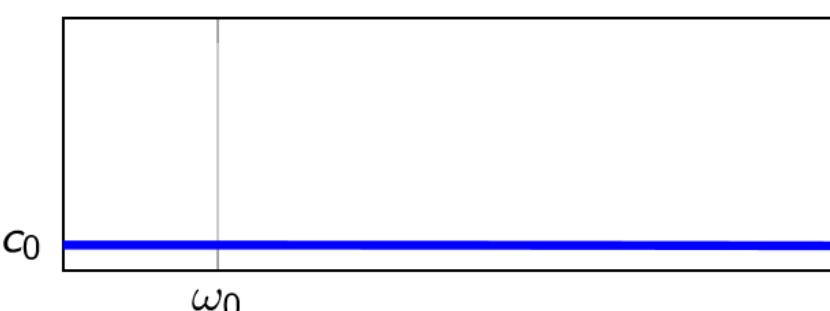
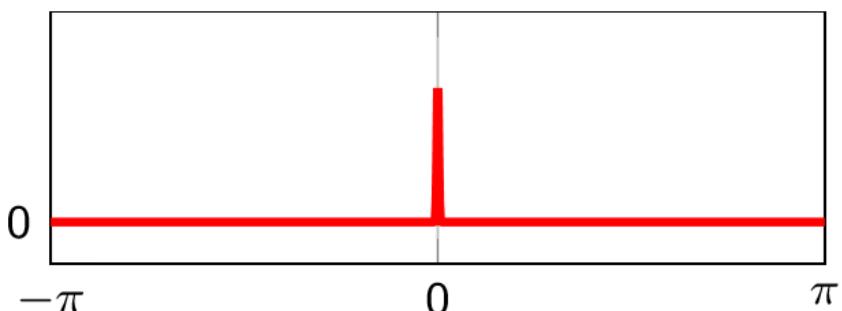
Fourier and the need for localization



Poor spatial localization



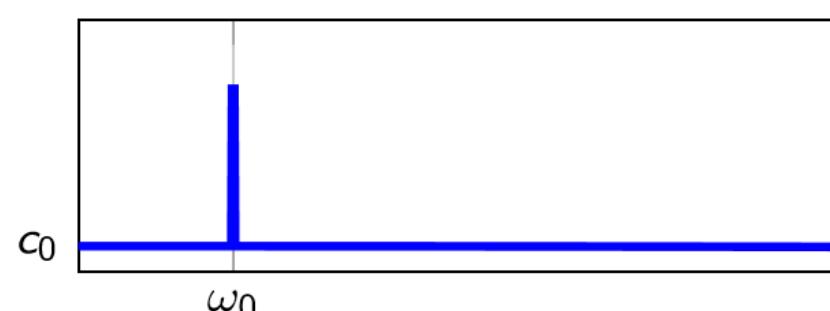
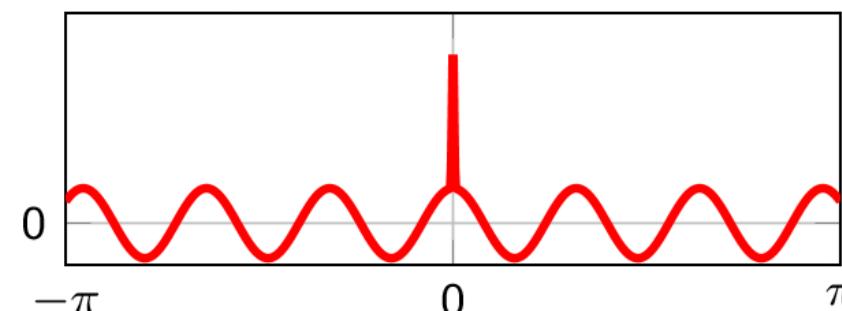
Good spectral localization



Good spatial localization



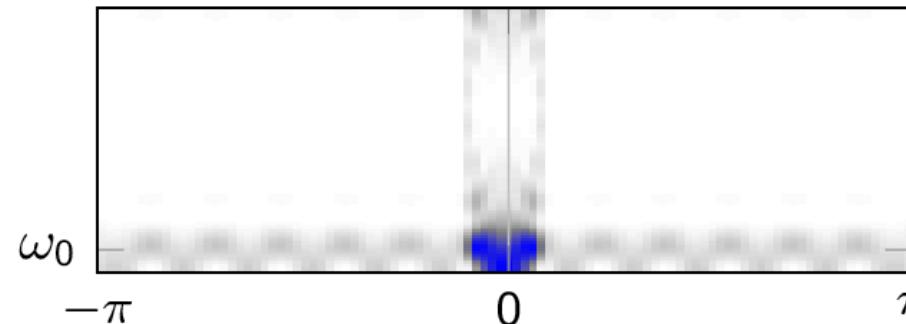
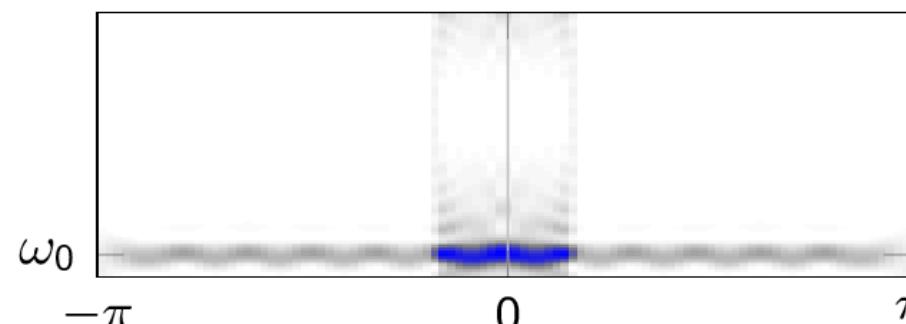
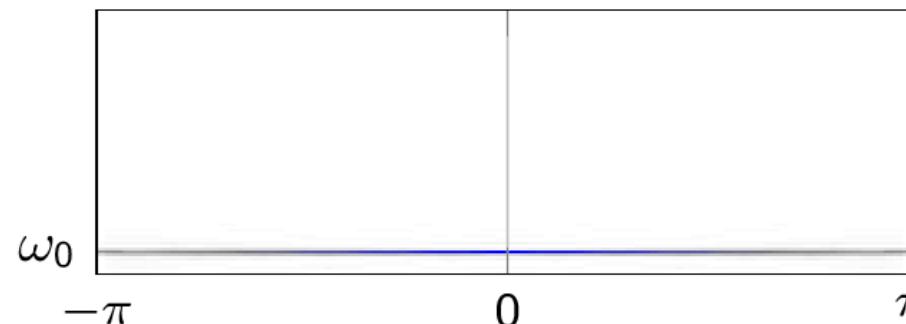
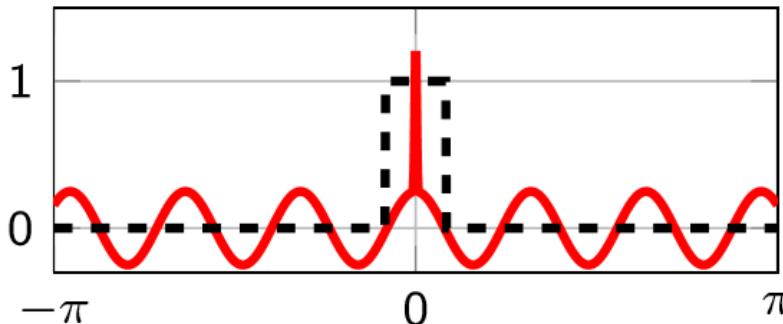
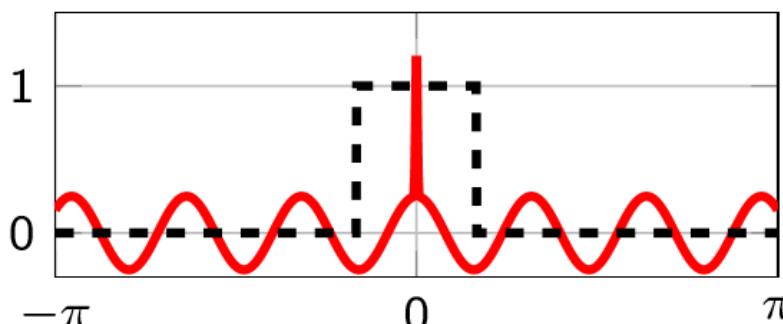
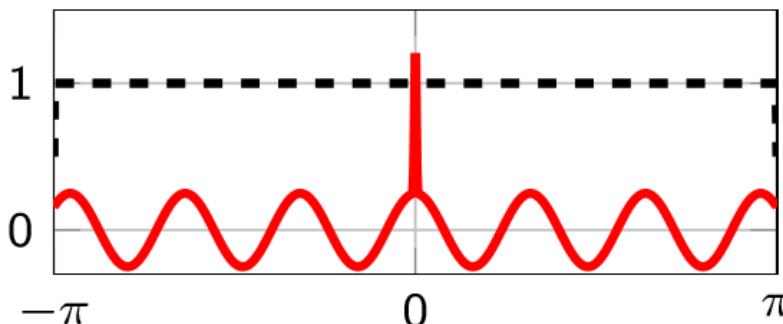
Poor spectral localization



["Learning class-specific descriptors for deformable shapes using localized spectral convolutional networks", Boscaini et al., 2015](#)

QUESTION: Do you know how this has been solved in signal processing?

WFT is the standard solution



Windowed Fourier Transform of a signal enveloped by a **window** g

$$WFT(f)(\xi, \omega) =$$

$$\underbrace{\int_D f(x)g(x-\xi)e^{-i\omega x}}_{\langle f(x), g_{\xi,\omega}(x) \rangle_{L^2(D)}}$$

WFT

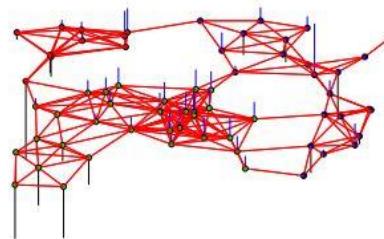
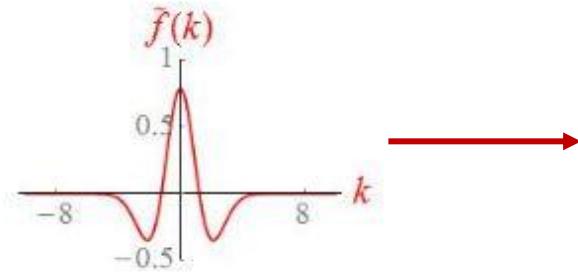
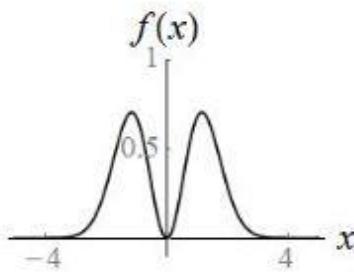
In some case the signal drastically changes in the space-time domain.



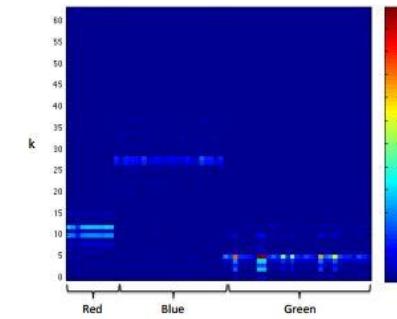
It is desirable to localize signal and analyze it locally.

The **Windowed Fourier Transform (WFT)** is the solution for this problem

Euclidean domains



Graphs



And
on
manifolds
?

Thorem of convolution

Convolution on Euclidean domain : $[-\pi, \pi]$ of two functions $f, g: [-\pi, \pi] \rightarrow \mathbb{R}$ is defined as:

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution Theorem: Fourier transform diagonalizes the convolution operator.

\Rightarrow *Convolution can be computed in the spectral domain*

$$\widehat{(f \star g)} = \hat{f} \cdot \hat{g}$$

Theorem of convolution

Convolution on Euclidean domain : $[-\pi, \pi]$ of two functions $f, g: [-\pi, \pi] \rightarrow \mathbb{R}$ is defined as:

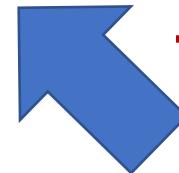
$$(f \star g)[x_i] = \sum_{j=1}^n f[x_j] g[x_i - x_j]$$

$$f \star g = \begin{bmatrix} g_1 & g_2 & \cdots & \cdots & g_n \\ g_n & g_1 & g_2 & \cdots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_2 & g_3 & \cdots & \cdots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix} = \Phi \hat{g} \Phi^\dagger f$$

Theorem of convolution

$$f * g = \begin{bmatrix} g_1 & g_2 & \cdots & \cdots & g_n \\ g_n & g_1 & g_2 & \cdots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_2 & g_3 & \cdots & \cdots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix} = \Phi \hat{g} \Phi^\dagger f$$

Thanks to the
Convolution
Theorem



$$\hat{g} = \begin{bmatrix} \hat{g}_1 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cdots & 0 & 0 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ 0 & 0 & \cdots & 0 & \hat{g}_n \end{bmatrix}$$

$$f * g = \Phi \begin{bmatrix} \hat{f}_1 \cdot \hat{g}_1 \\ \vdots \\ \vdots \\ \hat{f}_n \cdot \hat{g}_n \end{bmatrix} = \Phi \hat{f} \odot \hat{g}$$

WFT on non-Euclidean manifolds

Euclidean domains

Fourier Transform:

$$\hat{f}_\omega = \langle f, e^{2\pi i \omega x} \rangle$$

Convolution: $(f * g)(x) = \int_{-\infty}^{\infty} \hat{f}_\omega \hat{g}_\omega e^{2\pi i \omega x} d\omega$

Translation: $(T_u f)(x) = (f * \delta_u)(x) = f(x - u)$

Modulation: $(M_\omega f)(x) = e^{i\omega x} f(x)$

Manifolds

Fourier Transform: $\hat{f}_k = \langle f, \phi_k \rangle_{L^2(X)}$

Convolution: $(f * g)(x) = \sum_{k \geq 1} \hat{f}_k \hat{g}_k \phi_k(x)$

Translation: $(T_{x'} f)(x) = (f * \delta_{x'})(x) = \sum_{k \geq 1} \hat{f}_k \phi_k(x') \phi_k(x)$

Modulation: $(M_k f)(x) = \phi_k(x) f(x)$

WFT on non-Euclidean manifolds

Euclidean domains

Basic Atom: $g_{u, \omega}(x) =$
 $(M_\omega T_u g)(x) =$
 $g(x - u) e^{2\pi i \omega x}$

Manifolds

Basic Atom: $g_{x', k}(x) =$
 $(M_k T_{x'} g)(x) =$
 $\phi_k(x) \sum_{l \geq 1} \hat{g}_l \phi_l(x')$

Windowed Fourier Transform:

$$Sf(u, \omega) = \langle f, g_{u, \omega} \rangle$$

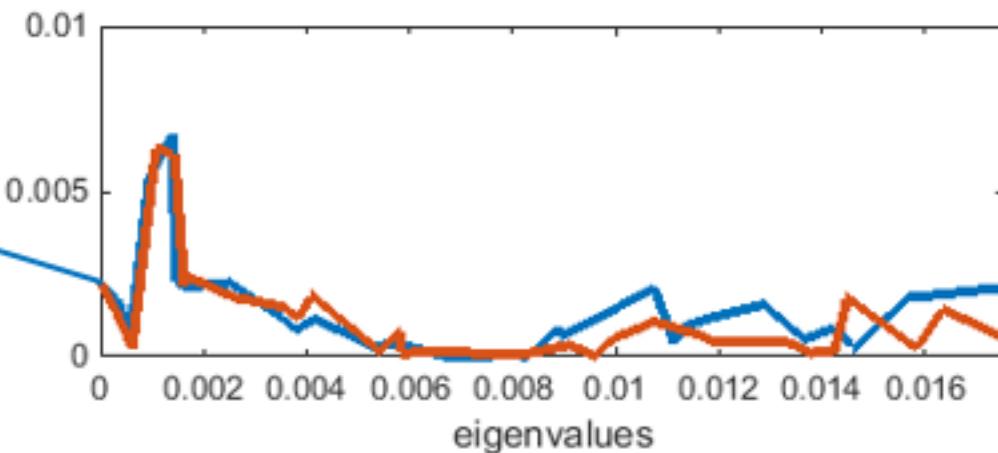
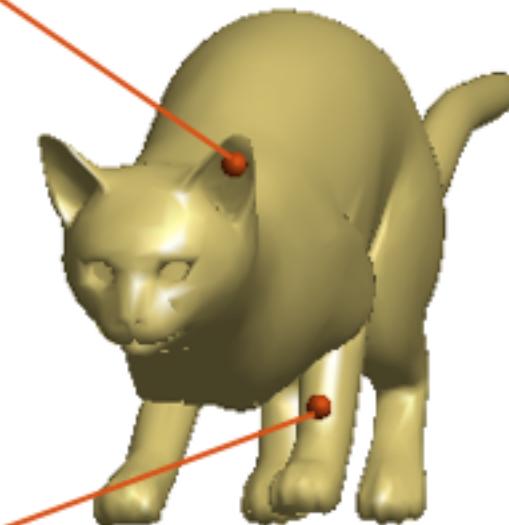
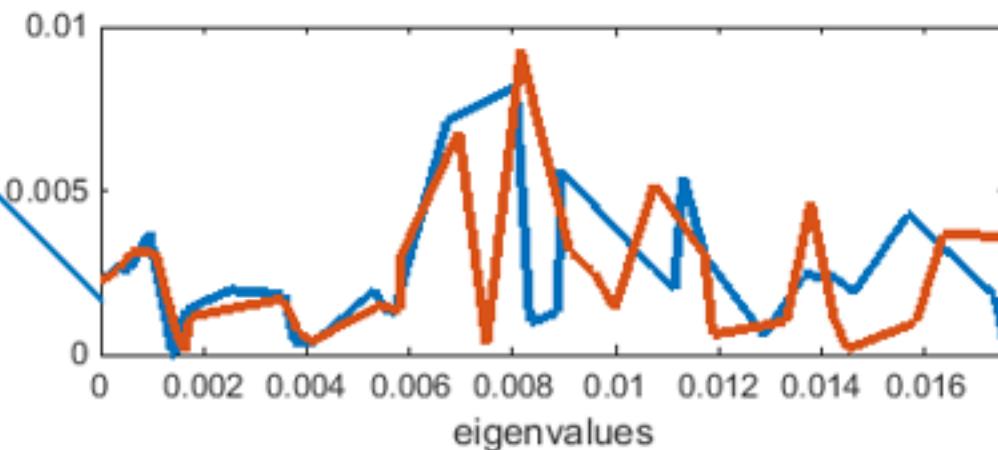
Windowed Fourier Transform:

$$(Sf)_{x, k} = \langle f, g_{x, k} \rangle_{L^2(X)}$$

WFT atoms



WFT problems



Spectral descriptors

A common structure is shared by the spectral descriptors
HKS and **WKS**

$$desc_q(x) = \sum_{l=1}^k g_{t_q}(\lambda_l) \phi_l^2(x), \quad \forall q \in 1, \dots, Q$$

A set of filters on the frequencies
=
functions of the eigenvalues

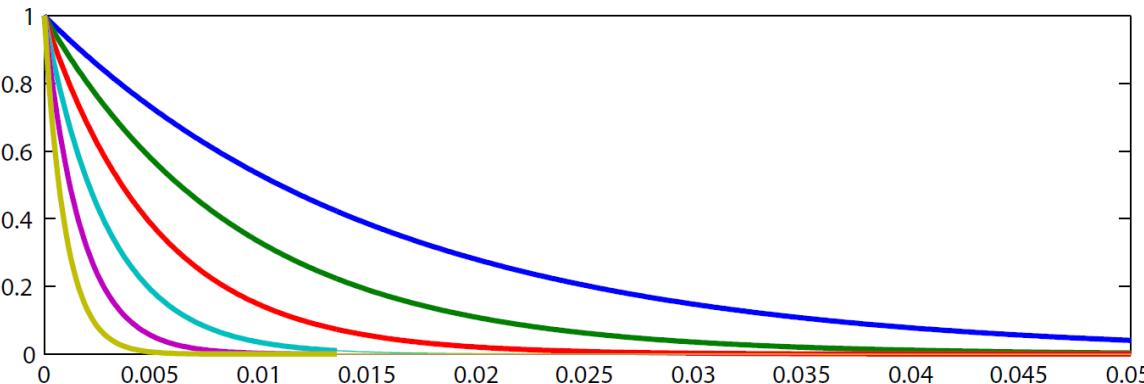
The square of each dimension of the spectral embedding

A signal processing point of view

$$desc_q(x) = \sum_{l=1}^k g_{t_q}(\lambda_l) \phi_l^2(x), \quad \forall q \in 1, \dots, Q$$

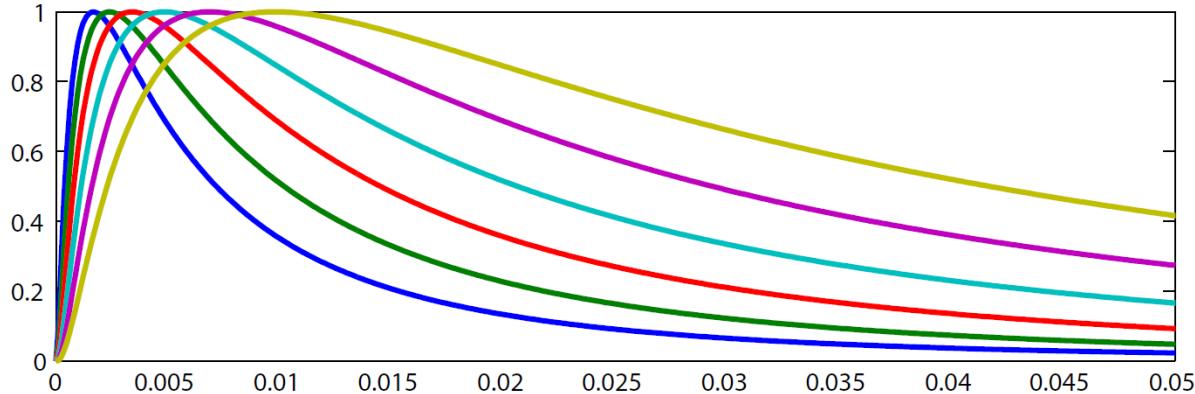
HKS:

$$g_t(\lambda_l) = e^{-t\lambda_l}$$



WKS:

$$g_t(\lambda_l) = e^{-\frac{(\log(E) - \log(\lambda_l))^2}{2\sigma^2}}$$



[“Learning spectral descriptors for deformable shape correspondence”, Litman et al., 2014.](#)

Could we obtain a stronger descriptor

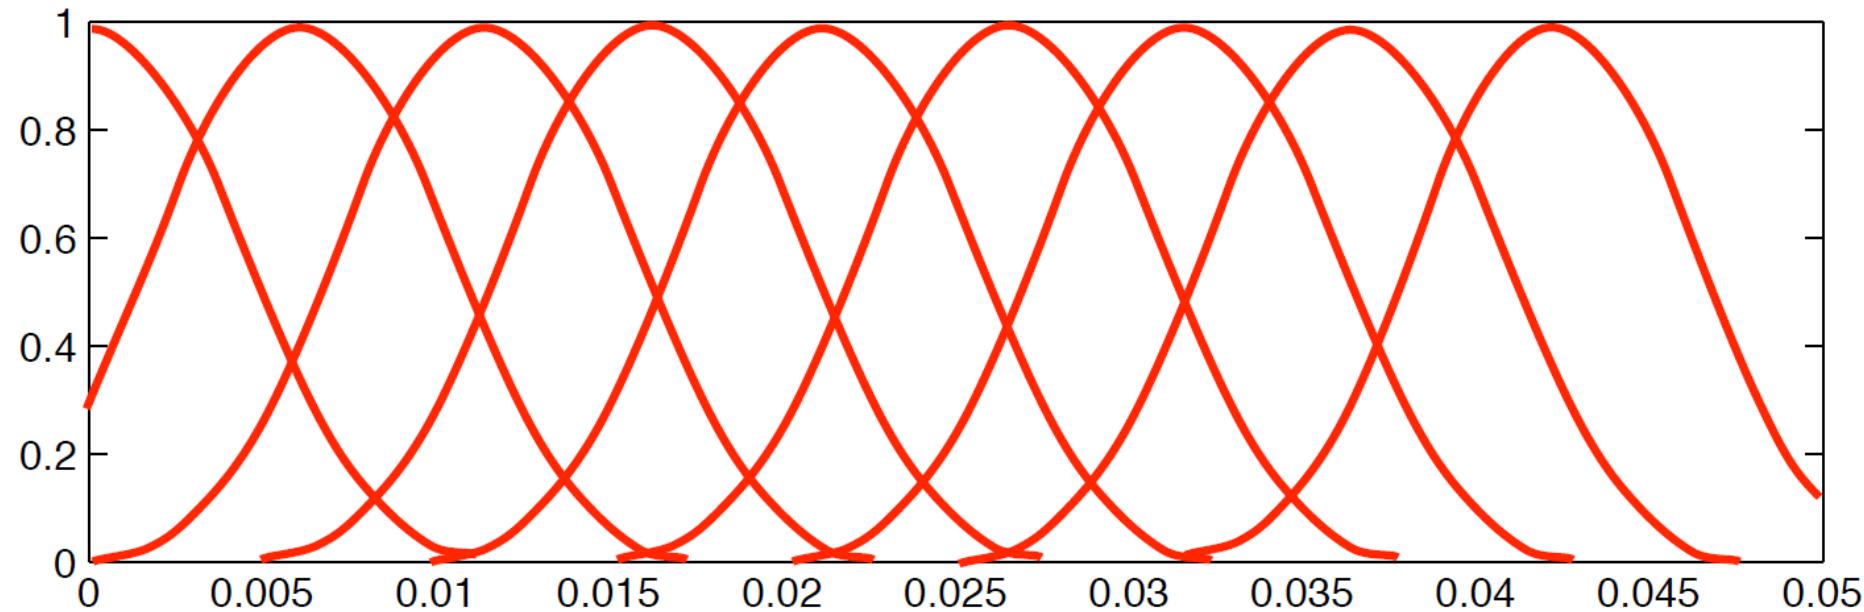
$$desc_q(x) = \sum_{l=1}^k g_{t_q}(\lambda_l) \phi_l^2(x), \quad \forall q \in 1, \dots, Q$$

What are the best filters to apply in this equation to obtain the best descriptors?

Can we learn them?

Learn filters for spectral descriptors

Given a set of basis functions: $\beta_1(\lambda), \beta_2(\lambda), \dots, \beta_Z(\lambda)$



We can learn the best coefficients to linearly combine them to obtain the best filters.

What do we need?

The q -th filter is obtained as

a linear combination of the

basis functions $\{\beta_z(\lambda_l)\}_{z=1}^Z$

$$g_q(\lambda_l) = \left(\sum_{z=1}^Z a_z^q \beta_z(\lambda_l) \right)$$

$$desc_q(x) = \sum_{l=1}^k g_{t_q}(\lambda_l) \phi_l^2(x), \quad \forall q \in 1, \dots, Q$$

$$desc_q(x) = \sum_{l=1}^k \left(\sum_{z=1}^Z a_z^q \beta_z(\lambda_l) \right) \phi_l^2(x)$$

What do we need?

$$desc_q(x) = \sum_{l=1}^k \left(\sum_{z=1}^Z a_z^q \beta_z(\lambda_l) \right) \phi_l^2(x)$$

we should learn the set of coefficients:

$$a_z^q \quad \forall q = 1, \dots, Q \text{ and } z = 1, \dots, Z$$

that is equivalent to learn a matrix:

$$\mathbf{A} \in \mathbb{R}^{Q \times Z} \text{ s.t. } A_{q,z} = a_z^q$$

Learned descriptors

We can compute a learned kernel signature by learning the matrix $A \in \mathbb{R}^{Q \times Z}$

$$LKS(x) = [desc_1^A(x), desc_2^A(x), \dots, desc_q^A(x)]$$

These explicitly depend on the learned matrix

How could we learn this matrix A ?

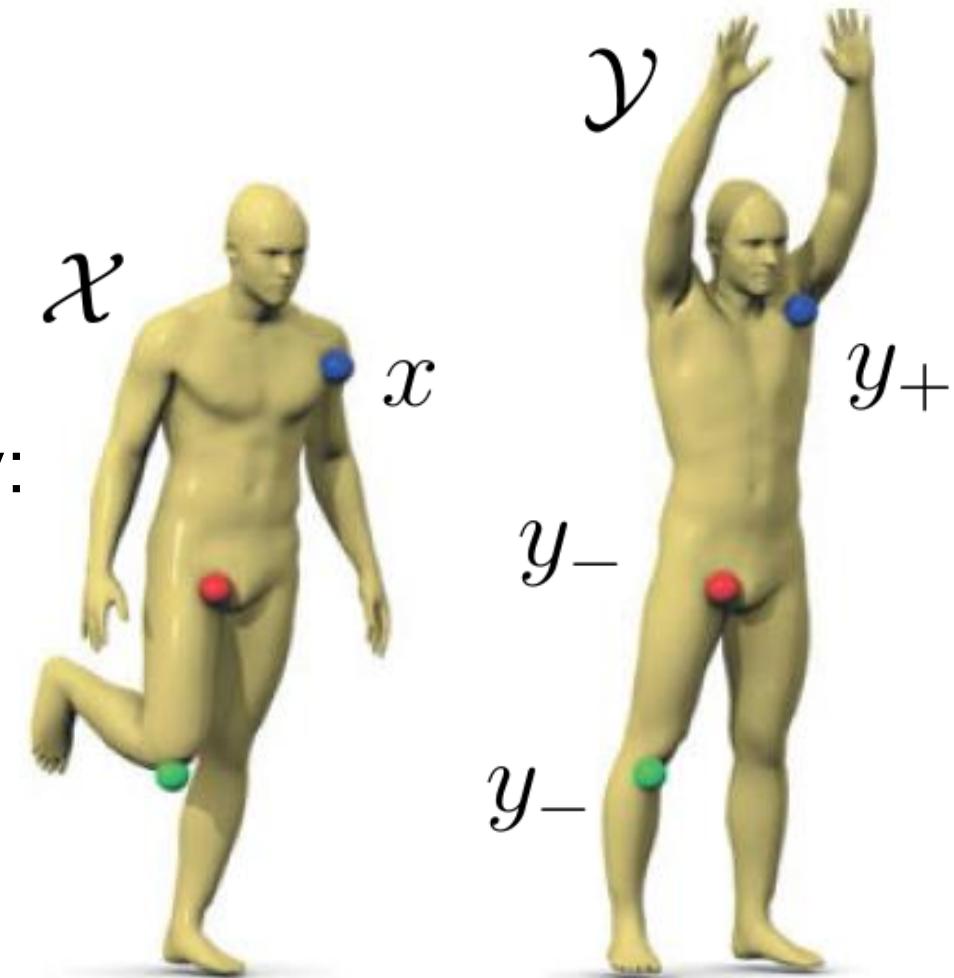
Loss definition

Given a pair of shapes \mathcal{X} and \mathcal{Y}

We consider a set of points X on \mathcal{X}

such that $\forall x \in X$ we can define a set
of points Y on \mathcal{Y} that is composed by:

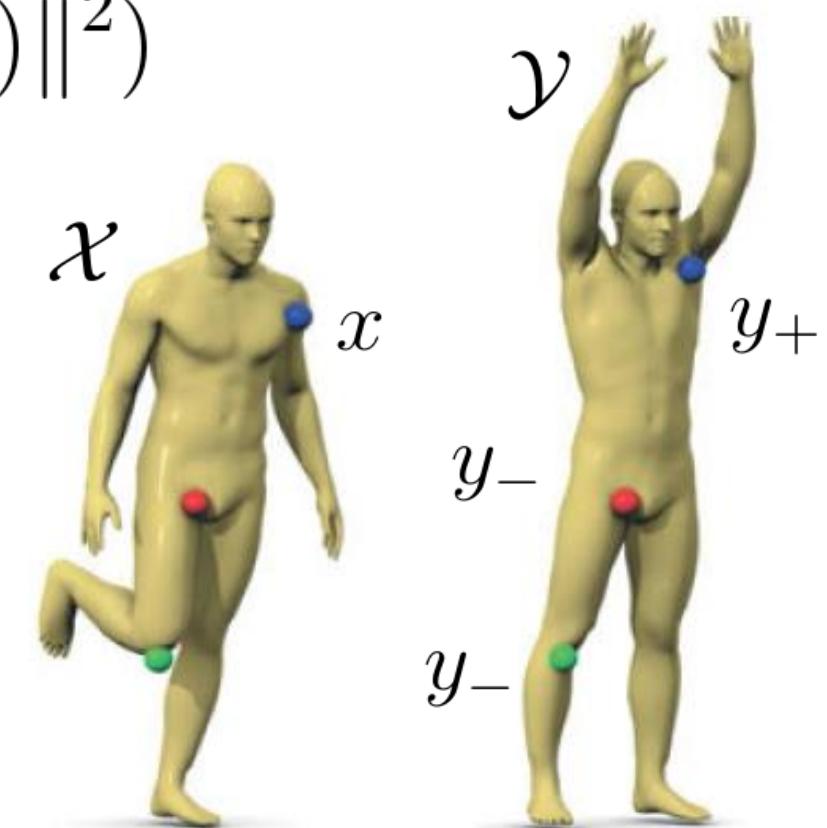
- similar points (**positive**) y_+
- dissimilar points (**negative**) y_-



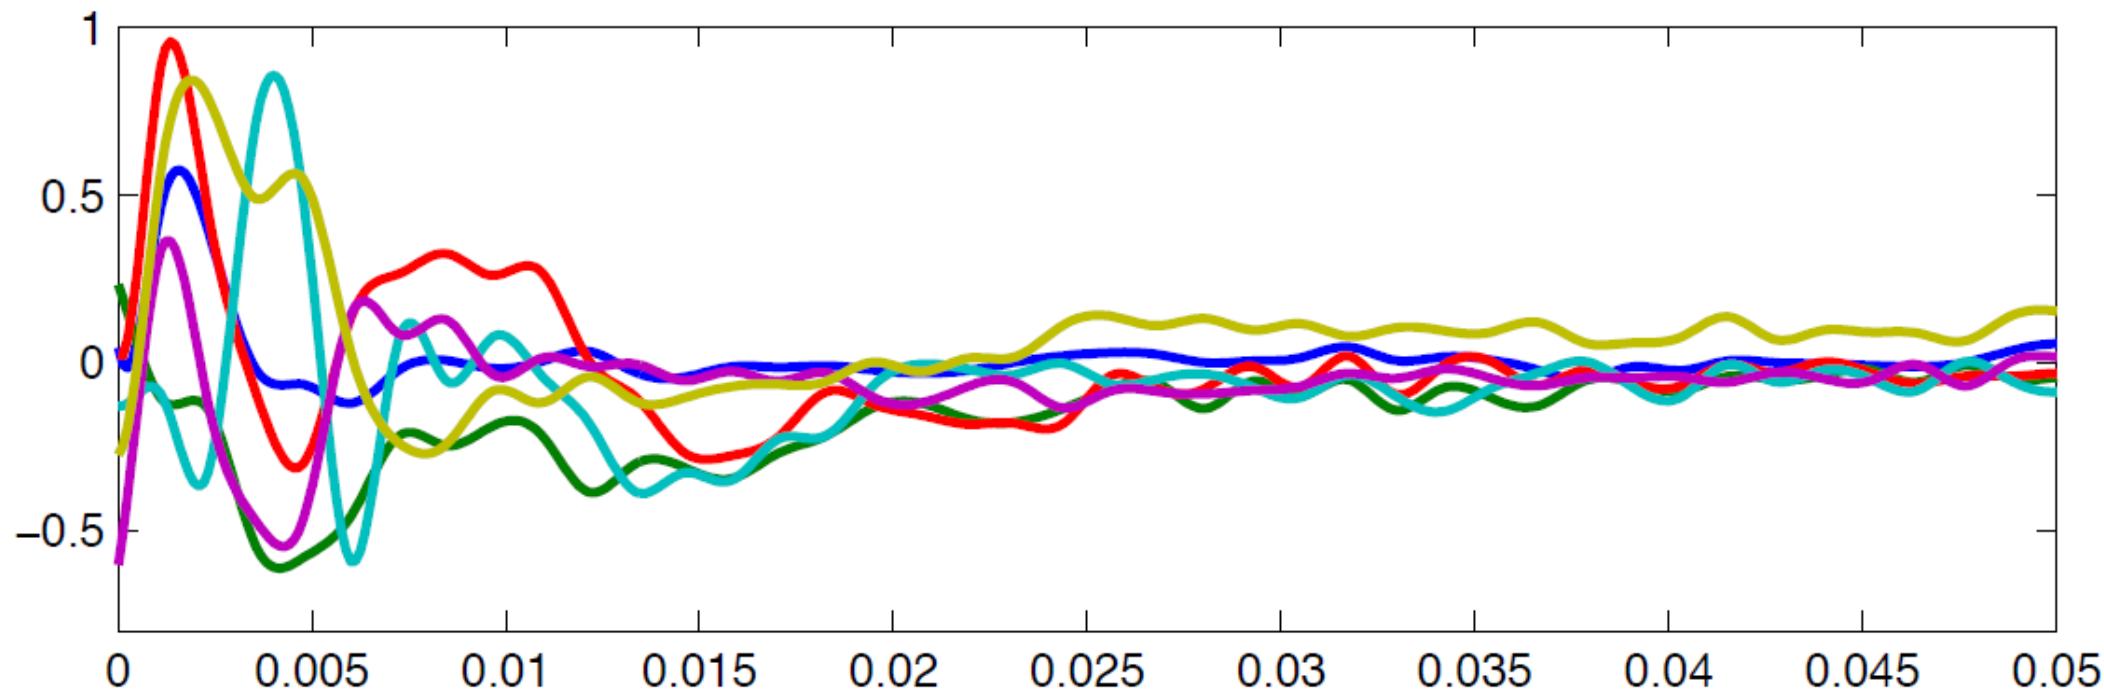
Loss definition

$$LKS(x) = [desc_1^A(x), desc_2^A(x), \dots, desc_q^A(x)]$$

$$\underset{\mathbf{A}}{argmin} \sum_{x \in X} \gamma(\|LKS(x) - LKS(y_+)\|^2) \\ -(1 - \gamma)(\|LKS(x) - LKS(y_-)\|^2)$$



Learned filters



Learn a new spectral descriptor

Given a set of $P \in \mathbb{N}$ functions $f_1, \dots, f_P : \mathcal{X} \rightarrow \mathbb{R}$

$$LSCNN_q(x) = \mathfrak{F}(f_1, \dots, f_P), \quad \forall q = 1, \dots, Q$$

$$LSCNN_q(x) = \sum_{p=1}^P \left(\sum_{k=1}^K a_{q,p,k} (Sf_p)_{x,k} \right)$$

The Windowed Fourier atoms for the function f with translation in x and modulation at frequency k

QUESTION:
what is not defined?

Learn the window for each input function

It is easier to learn it in the spectral domain!

We already did something similar, do you remember where?

In the definition of the optimal shape descriptor!

The windows are obtained
as a linear combination of the
basis functions $\{\beta_z(\lambda_l)\}_{z=1}^Z$

$$g_p(\lambda_k) = \sum_{z=1}^Z b_z^p \beta_z(\lambda_k)$$

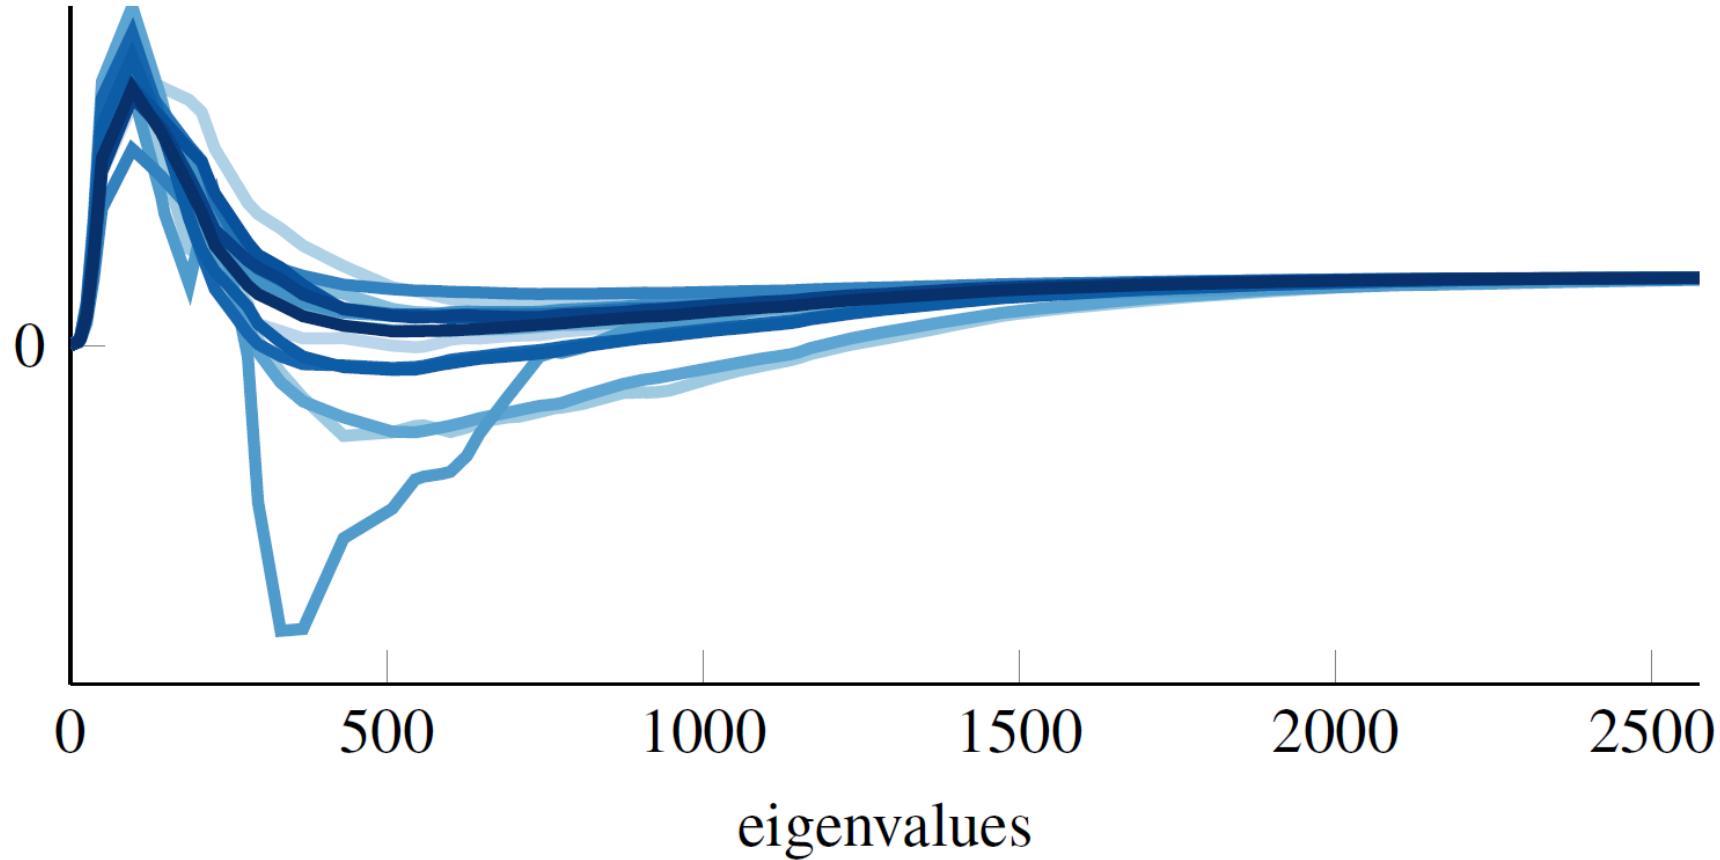
Localized spectral CNN descriptor (LSCNN)

$$desc(x) = [LSCNN_1(x), LSCNN_2(x), \dots, LSCNN_Q(x)]$$

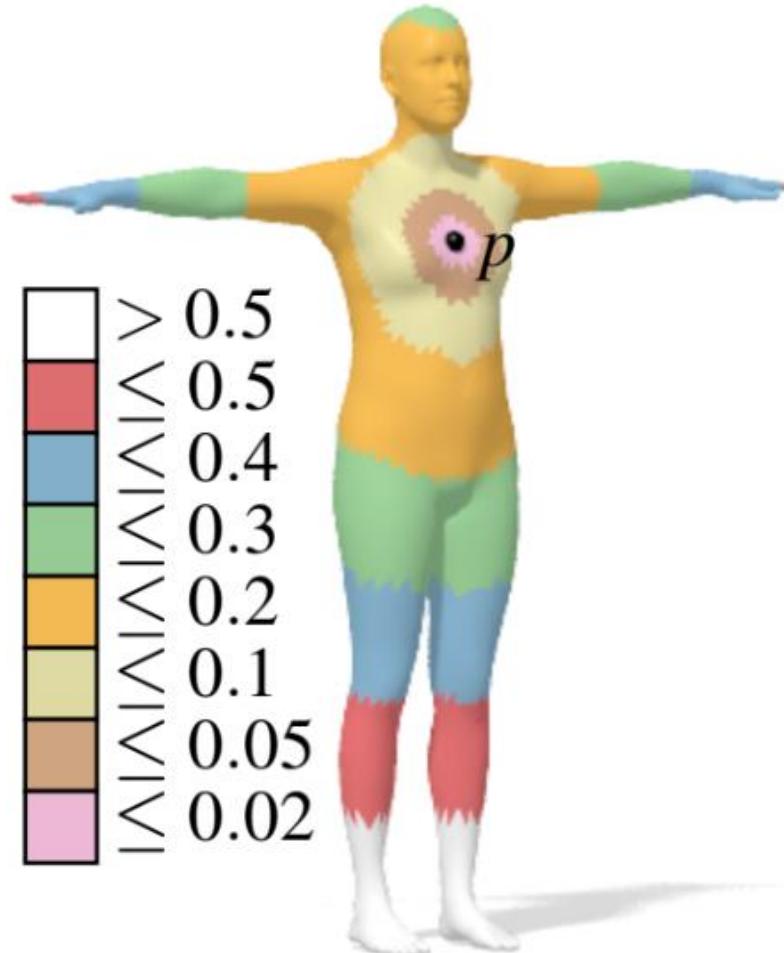
$$LSCNN_q(x) = \sum_{p=1}^P \left(\sum_{k=1}^K a_{q,p,k} (Sf_p)_{x,k} \right)$$

$$(Sf_p)_{x,k} = \langle f_p, \sum_{l=1}^K \left(g_p(\lambda_l) \phi_l(x) \phi_k(x) \right) \phi_l \rangle \chi$$

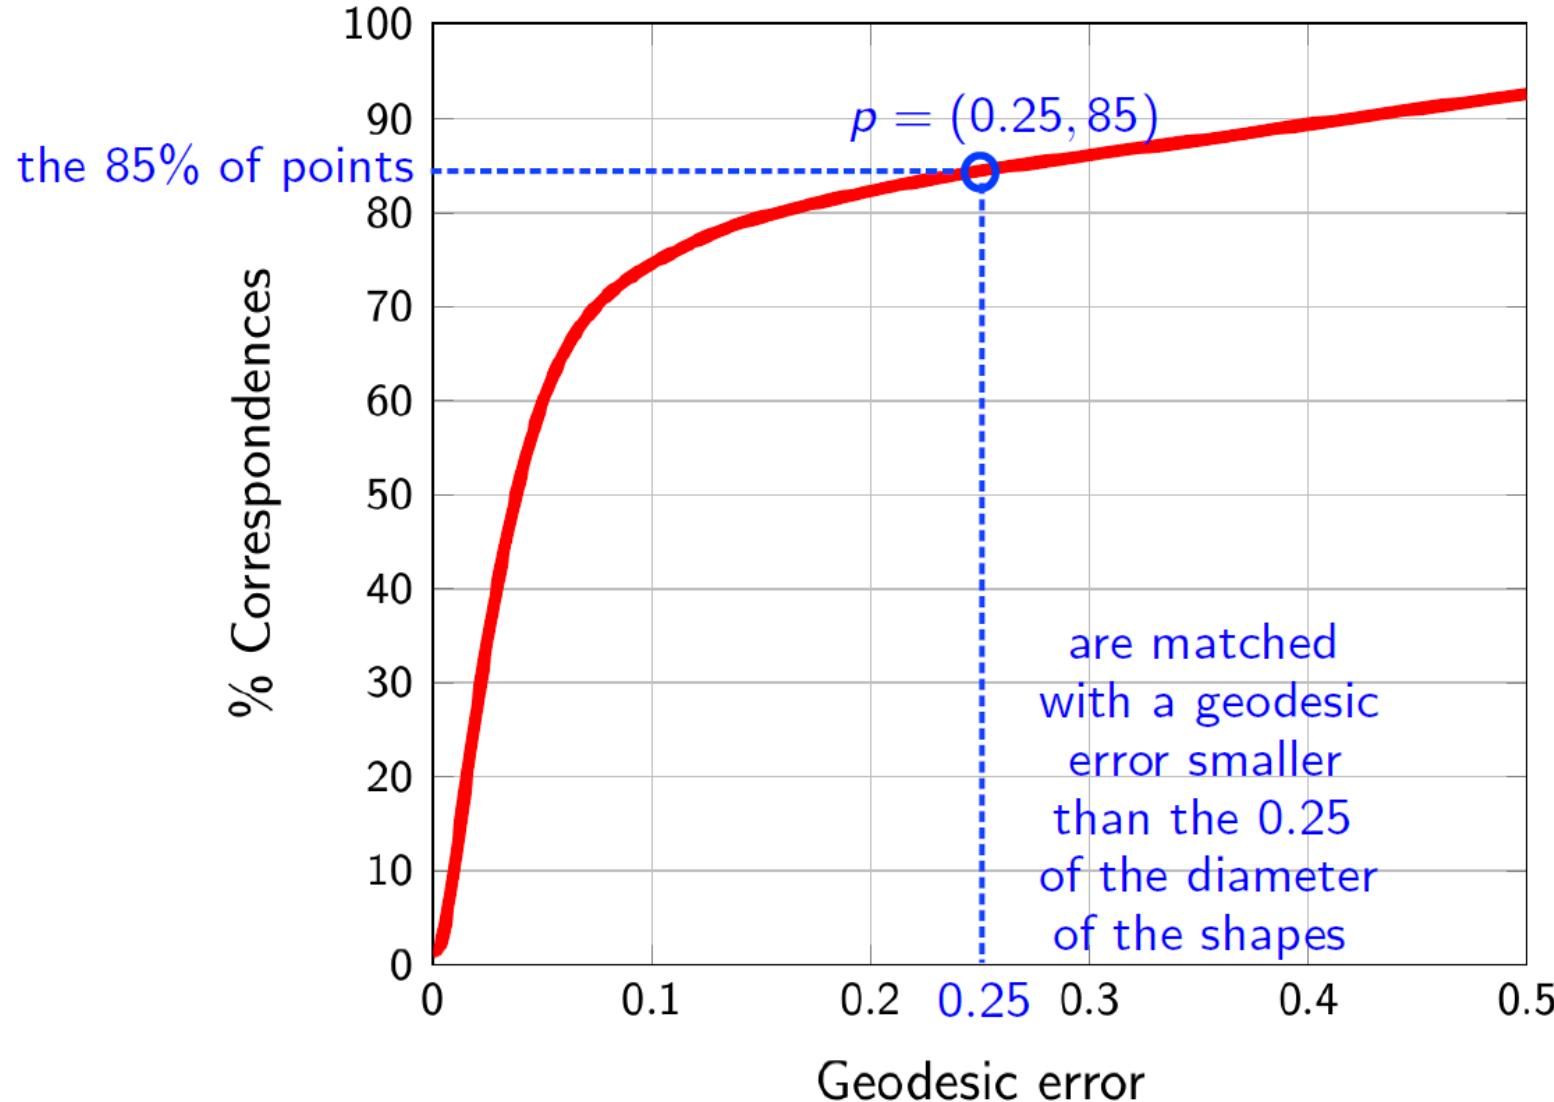
Learned windows



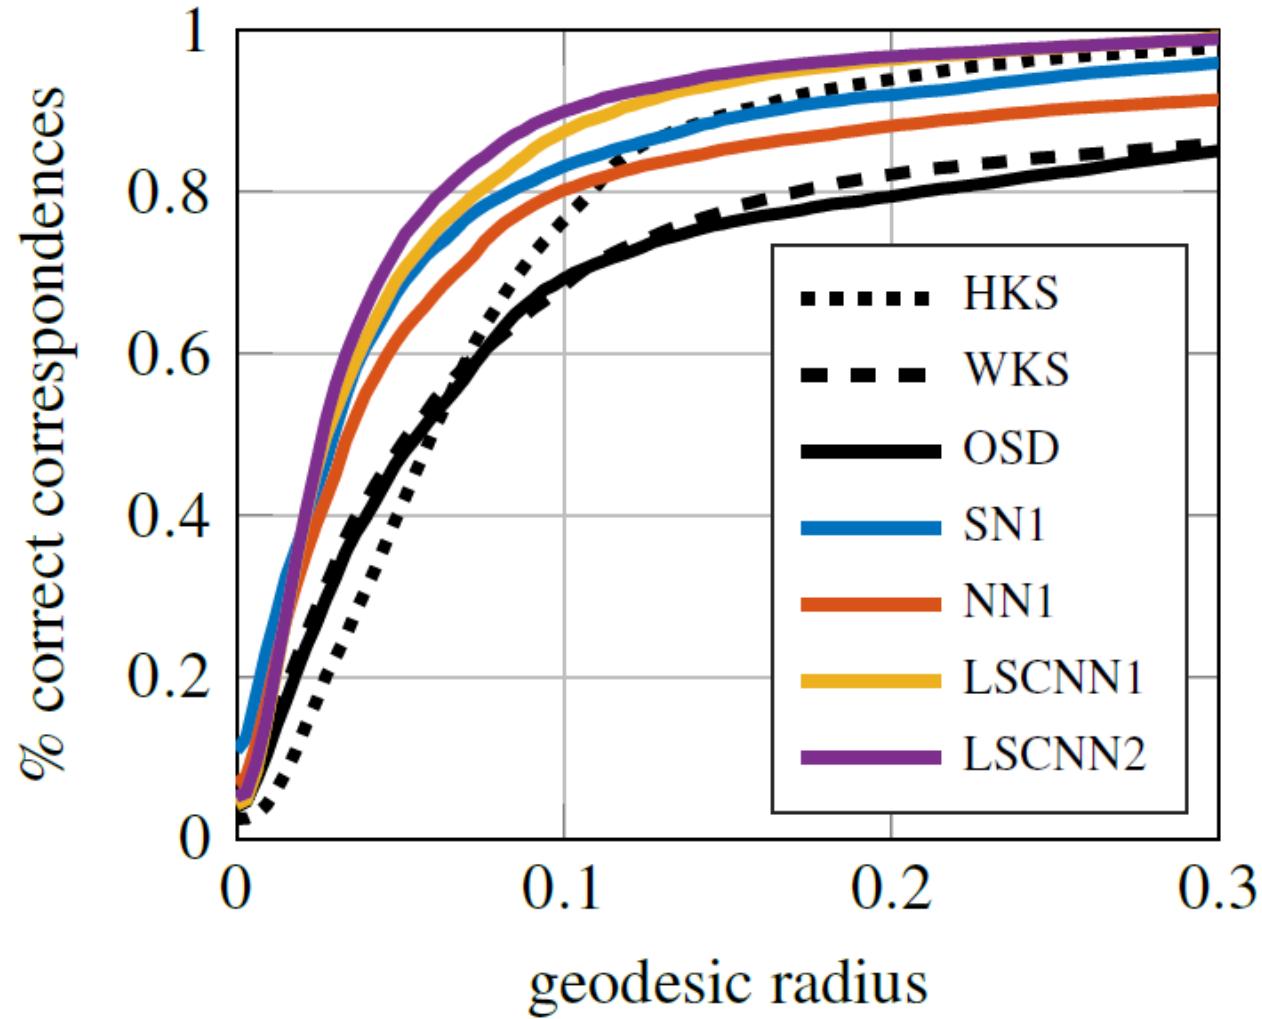
Geodesic error



Cumulative error curve



Quantitative comparison



Qualitative comparison

HKS



WKS



OSD



Qualitative comparison

HKS



WKS



LSCNN



Some conclusions

- Spectral descriptors are invariant to isometric deformations
- Spectral descriptors do not solve the symmetries
- Spectral descriptors can be generalized via data-driven approaches
- WFT can characterize locally the shape
- The data-driven approaches outperform the standard spectral ones
- Other deformations (for from isometries) can not be faced

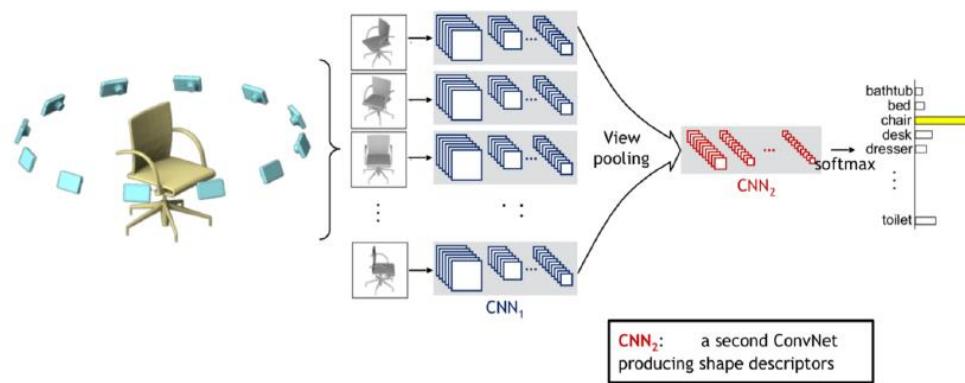
Other data-drive approaches

The data-driven approaches seem well-suited to solve the point-to-point matching problem

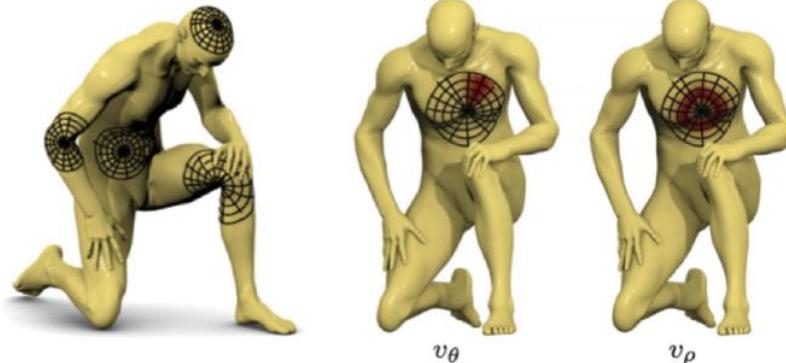
Recently this gives rise to a family of approaches that can be collected under the name of:

GEOMETRIC DEEP LEARNING

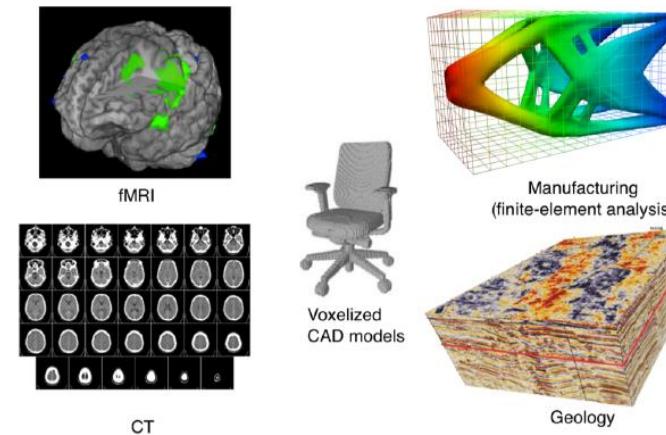
Geometric deep learning



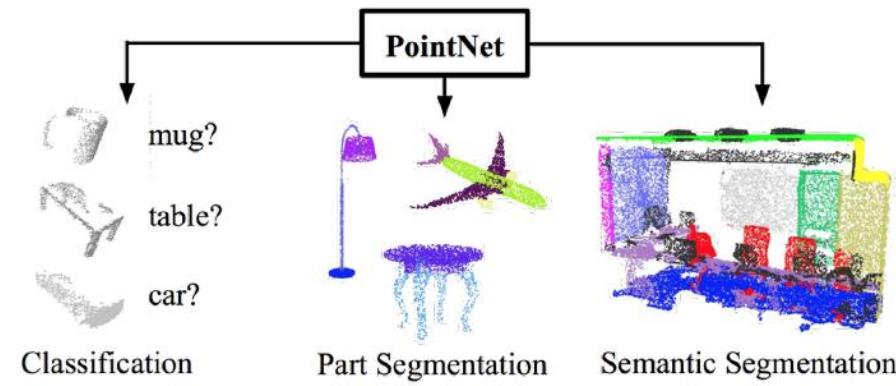
View-based



Intrinsic (surface-based)



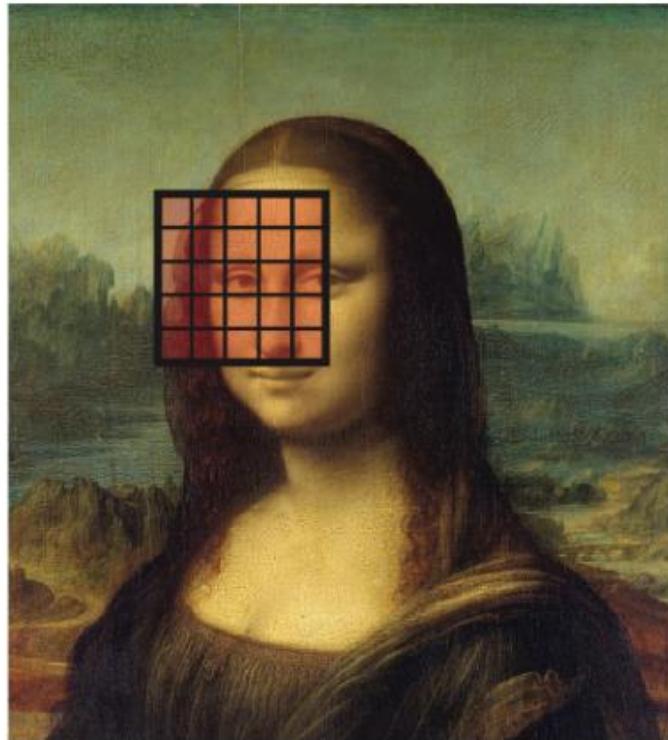
Volumetric



Point-based

Slide credit M. Ovsjanikov

Alternative convolutions



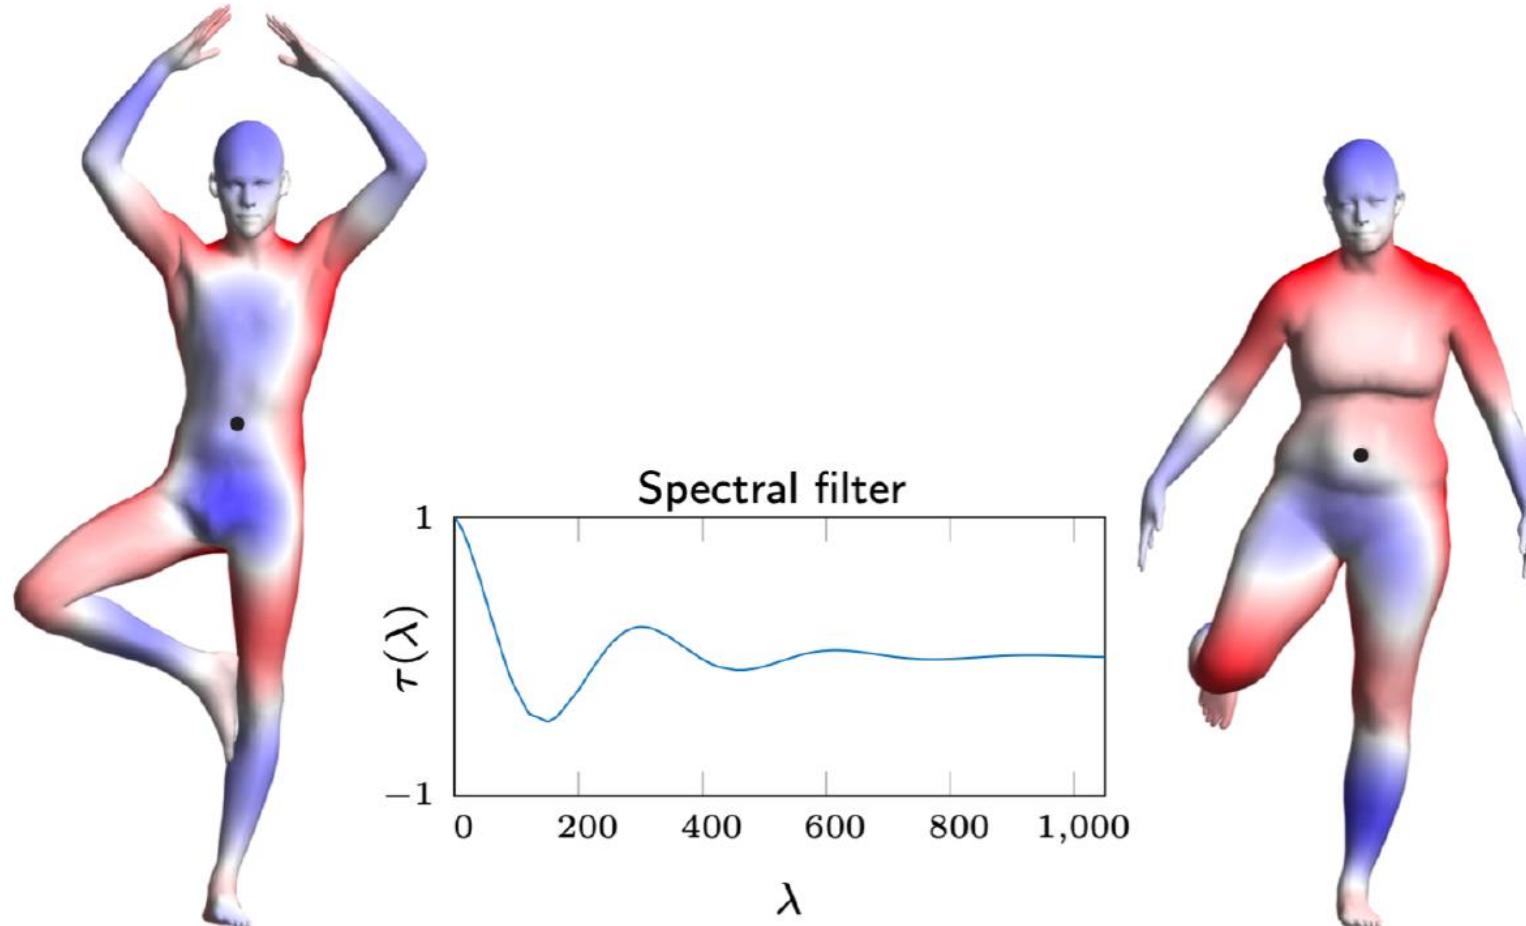
Euclidean



Non-Euclidean

Limits of the spectral convolution

Unfortunately spectral convolution has many limitations (shape, shift invariances).



$$\Phi_1 \tau(\Lambda_1) \Phi_1^\top \delta_0$$

$$\Phi_2 \tau(\Lambda_2) \Phi_2^\top \delta_0$$

Slide credit E. Rodolà

Spectral convolution:

Slide credit E. Rodolà

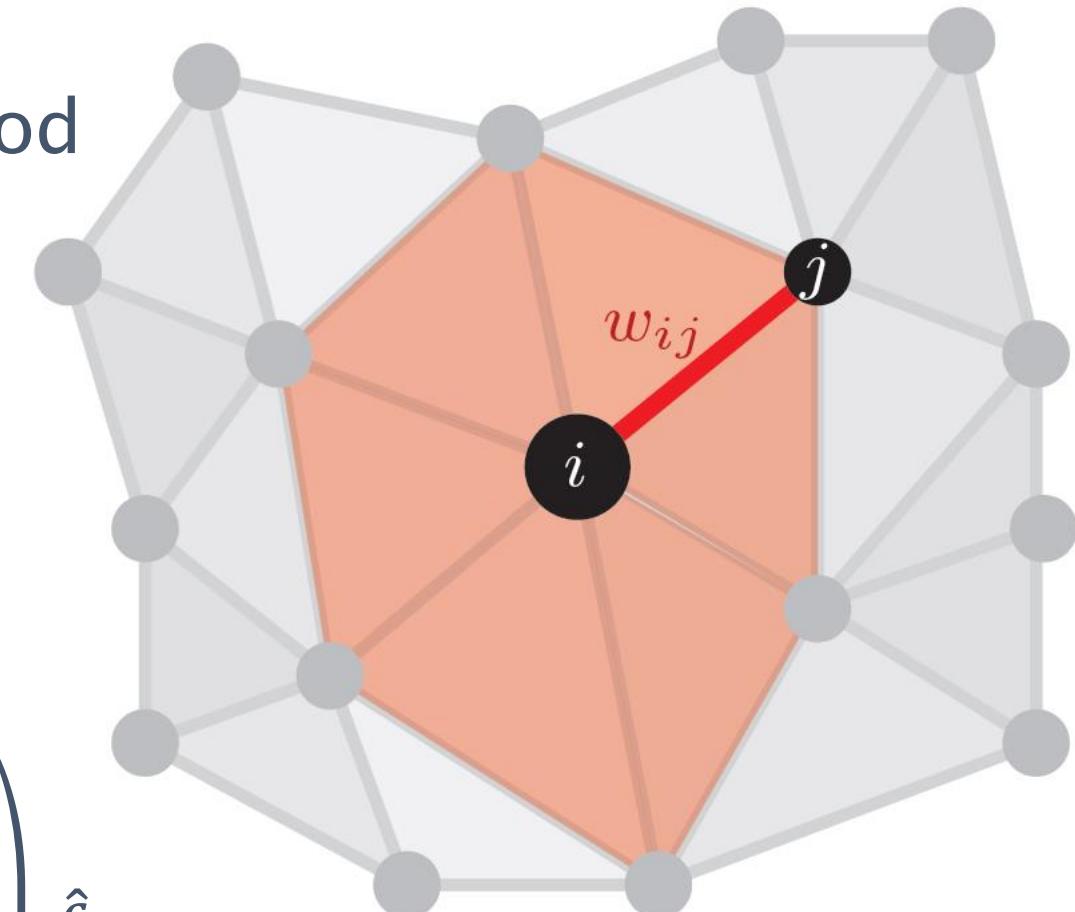
Laplacian operator Δ which
Acts locally on the neighborhood

$$(\Delta f)_i = \sum_{j \in \mathcal{N}(i)} w_{ij} (f_j - f_i)$$

LBO eigenvectors generalize
the Fourier basis $\hat{f} = \Phi^\dagger f$

Spectral
convolution:

$$f * g = \Phi \begin{pmatrix} \hat{g}_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & 0 & \cdots & 0 & 0 \\ \vdots & & & & \ddots & \vdots \\ 0 & 0 & & & & 0 \\ 0 & 0 & & \cdots & & \hat{g}_n \end{pmatrix} \hat{f}$$



Spectral convolution:

Slide credit E. Rodolà

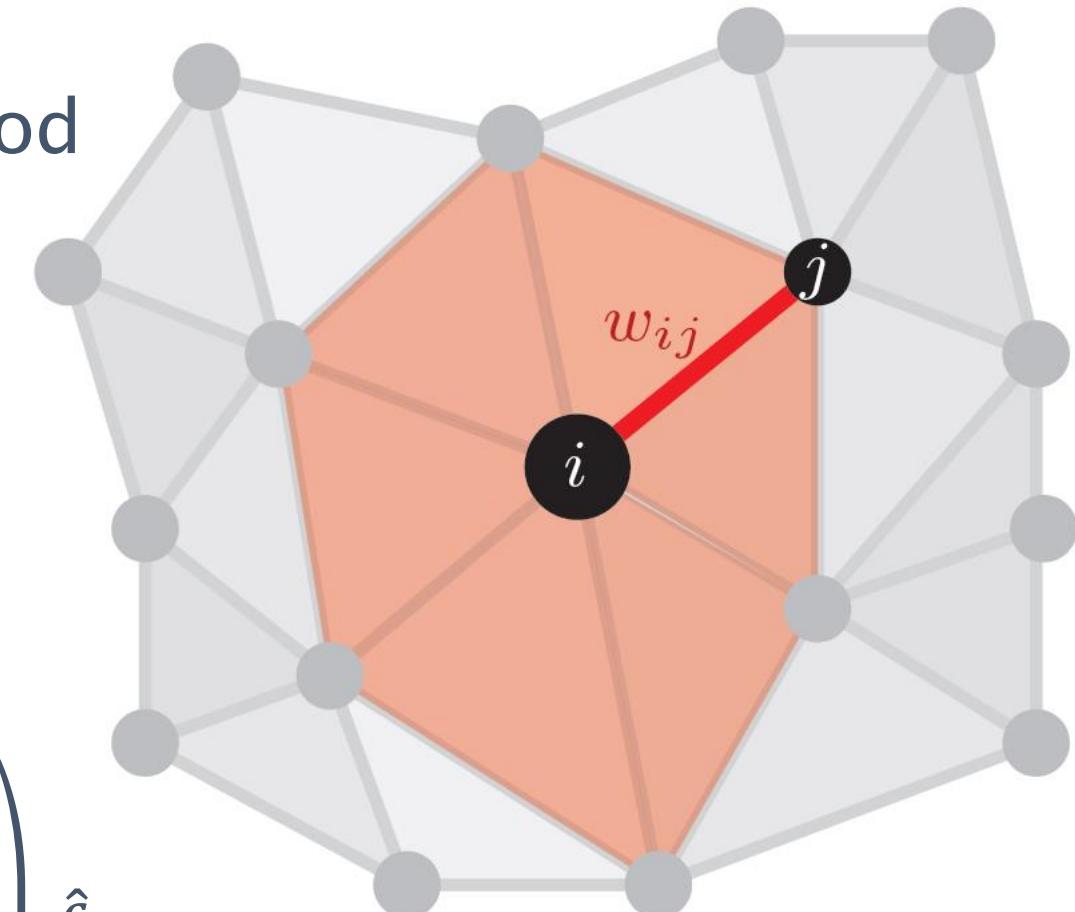
Laplacian operator Δ which
Acts locally on the neighborhood

$$(\Delta f)_i = \sum_{j \in \mathcal{N}(i)} w_{ij} (f_j - f_i)$$

LBO eigenvectors generalize
the Fourier basis $\hat{f} = \Phi^\dagger f$

Spectral
convolution:

$$f * g = \Phi \begin{pmatrix} \hat{g}_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & 0 & \cdots & 0 & 0 \\ \vdots & & & & \ddots & \vdots \\ 0 & 0 & & & & 0 \\ 0 & 0 & & \cdots & & \hat{g}_n \end{pmatrix} \hat{f}$$



Spectral convolution:

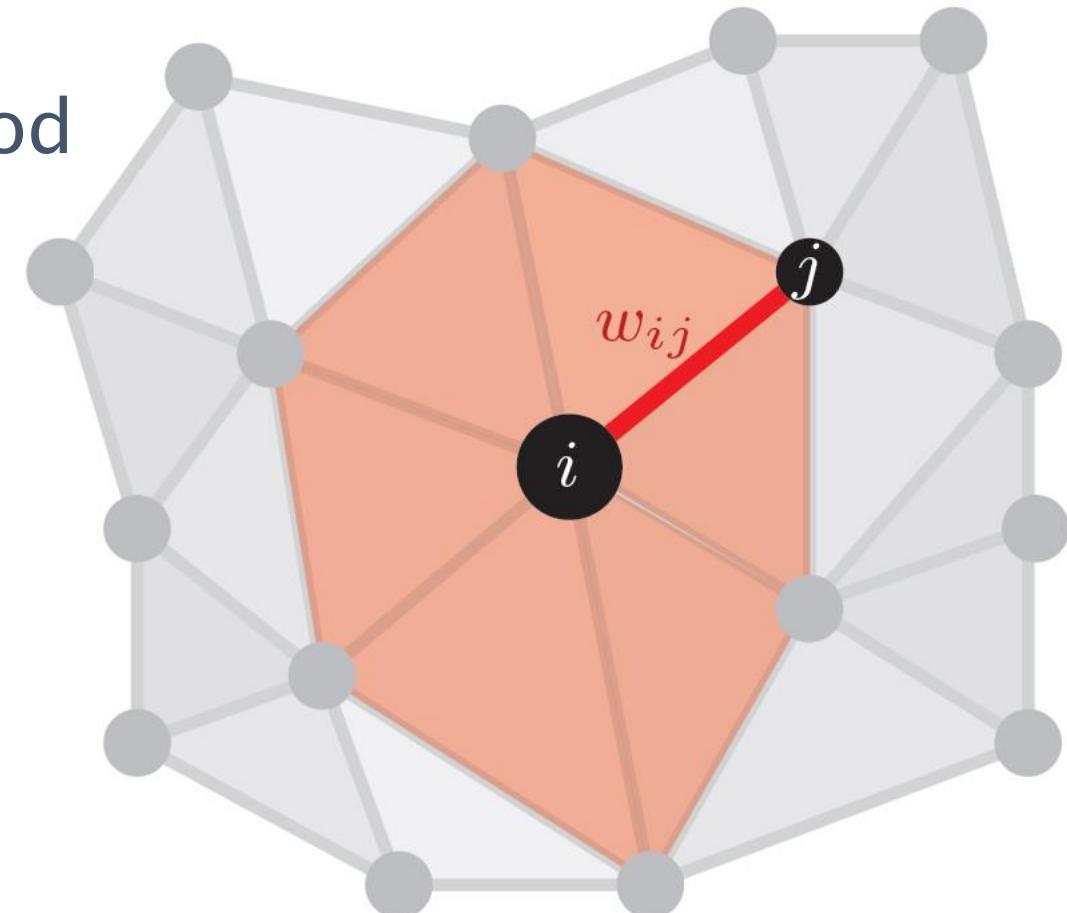
Slide credit E. Rodolà

Laplacian operator Δ which
Acts locally on the neighborhood

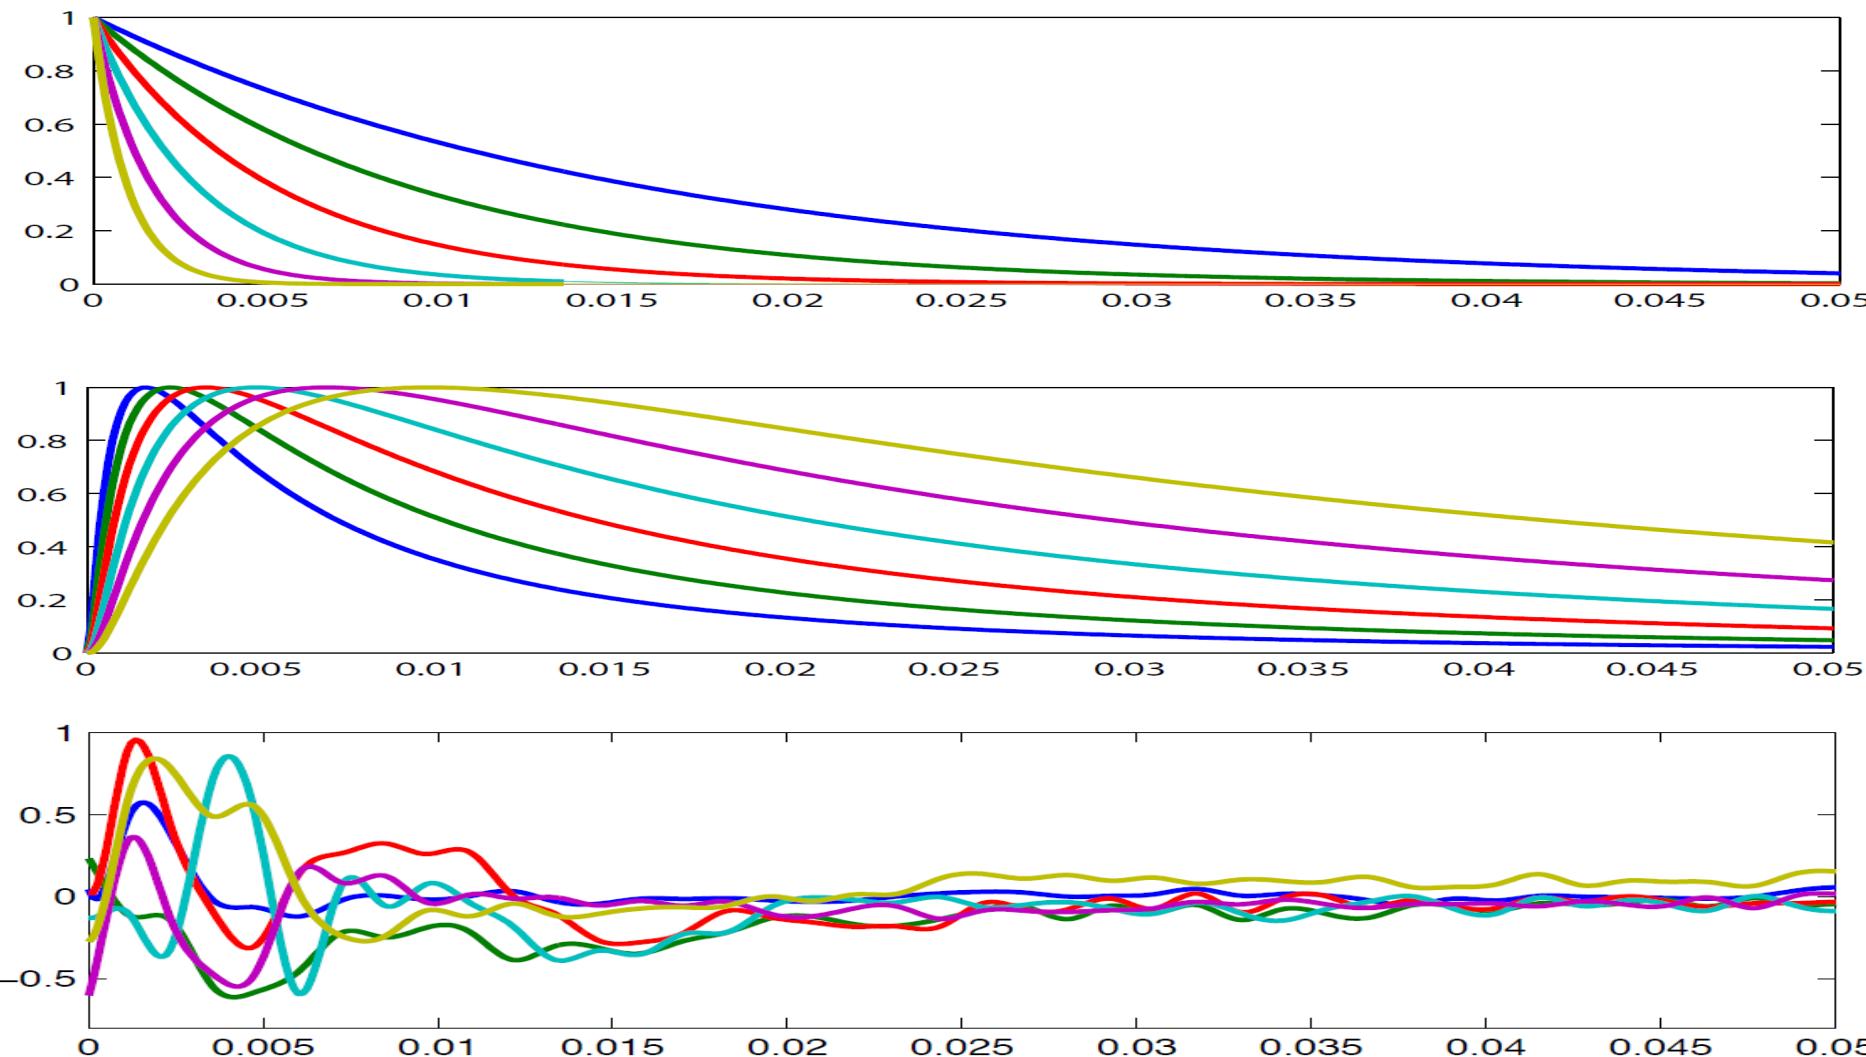
$$(\Delta f)_i = \sum_{j \in \mathcal{N}(i)} w_{ij} (f_j - f_i)$$

LBO eigenvectors generalize
the Fourier basis $\hat{f} = \Phi^\dagger f$

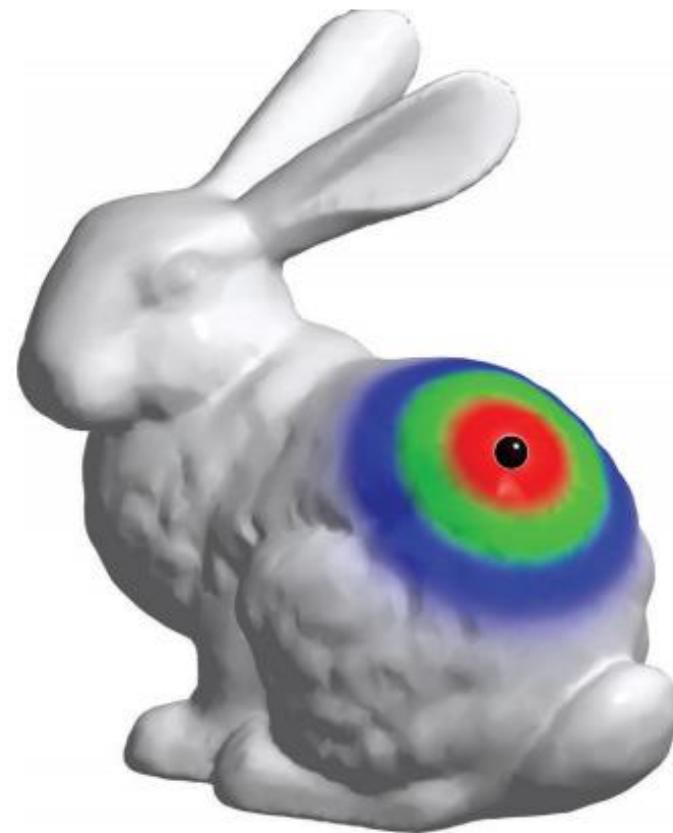
Spectral
convolution: $f' = \Phi \tau(\Lambda) \Phi^\dagger f$



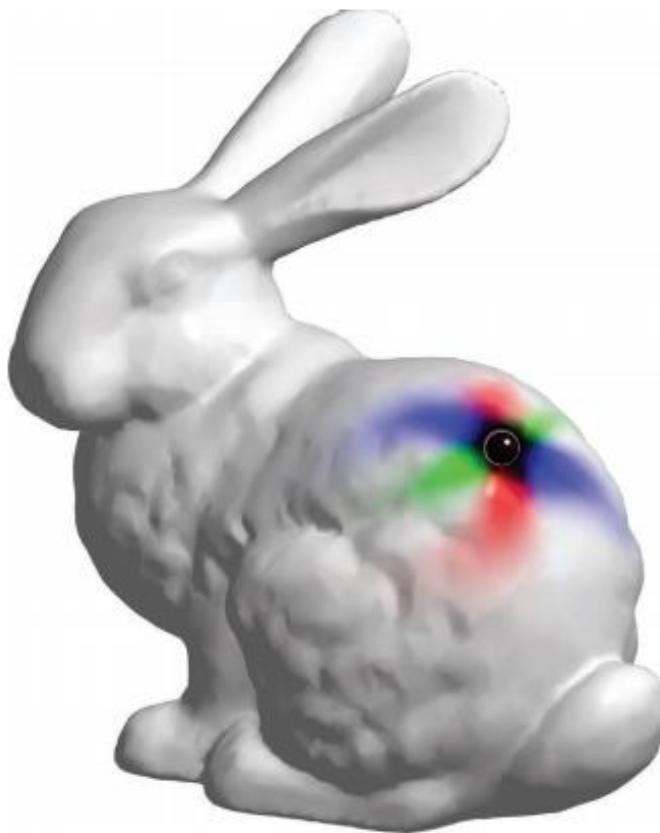
Spectral filtering convolution:



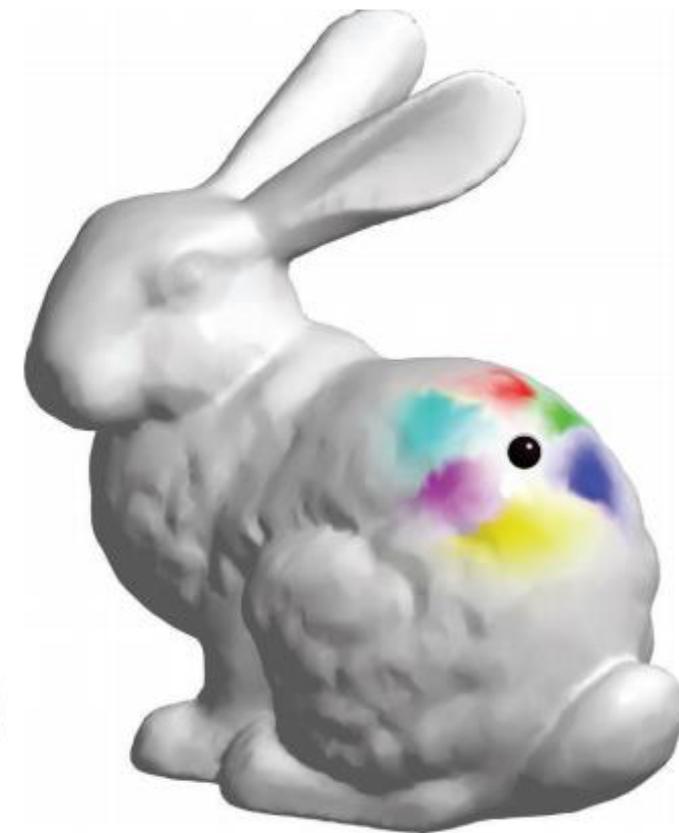
Geometric deep learning



Diffusion distance



Anisotropic
heat kernel



Geodesic polar
coordinates

["Geometric deep learning: going beyond Euclidean data" Bronstein et al., 2016](#)